


# EECE 5550: Mobile Robotics




## Lecture 12: Localization

# Recap

## Last time: mapping

**Given:** Robot poses  $x_{0:t}$ , measurements  $z_{1:t}$


**Estimate:** Belief  $p(m|x_{0:t}, z_{1:t})$  over the map  $M$



## This time: localization

**Given:** Prior  $p(x_0)$ , map  $m$ , controls  $u_{0:t}$ , obs.  $z_{1:t}$

**Estimate:** Belief  $p(x_t | m, z_{1:t})$  over the robot pose




# Plan of the day

- Monte Carlo localization
- Models:
  - Probabilistic motion models
  - Forward beam sensor models
- Practicalities:
  - Particle diversity / depletion
  - Tracking localization performance
  - Adaptive sample sets: KLD-sampling



# References



Chapters 5, 6, and 8 of “Probabilistic Robotics”


# Monte Carlo Localization

**Main idea:** Monte Carlo localization (MCL) is simply localization using the Particle Filter


**Given:** Prior  $p(x_0)$ , map  $m$ , controls  $u_{0:t}$ , observations  $z_{1:t}$

**Estimate:** Belief  $p(x_t | m, z_{1:t})$  over the robot pose


```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):  
2:     $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$   
3:    for  $m = 1$  to  $M$  do  
4:      sample  $x_t^{[m]} \sim p(x_t | u_t, x_{t-1}^{[m]})$   
5:       $w_t^{[m]} = p(z_t | x_t^{[m]})$   
6:       $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$   
7:    endfor  
8:    for  $m = 1$  to  $M$  do  
9:      draw  $i$  with probability  $\propto w_t^{[i]}$   
10:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$   
11:    endfor  
12:    return  $\mathcal{X}_t$ 
```



# Example application





Initial belief




Measurement



Move right



Measurement



Move right


# Monte Carlo Localization

**Main idea:** Monte Carlo localization (MCL) is simply localization using the Particle Filter

**Given:** Prior  $p(x_0)$ , map  $m$ , controls  $u_{0:t}$ , observations  $z_{1:t}$

**Estimate:** Belief  $p(x_t | m, z_{1:t})$  over the robot pose

```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):  
2:     $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$   
3:    for  $m = 1$  to  $M$  do  
4:      sample  $x_t^{[m]} \sim p(x_t | u_t, x_{t-1}^{[m]})$  ← Motion model  
5:       $w_t^{[m]} = p(z_t | x_t^{[m]})$  ← Sensing model  
6:       $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$   
7:    endfor  
8:    for  $m = 1$  to  $M$  do  
9:      draw  $i$  with probability  $\propto w_t^{[i]}$   
10:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$   
11:    endfor  
12:    return  $\mathcal{X}_t$ 
```




## Requires:

- A *sampler* for the motion model
- *Likelihood function* for the sensor model

# Robot motion models

We will consider two primary types of robot motion models:

- Velocity-based
- Odometry-based





# Velocity-based motion models

**Goal:** We need a procedure for **drawing samples**  $x_{t+1} \sim p(x_{t+1}|x_t, u_t)$  from the robot's motion model


**Recall:** In Lecture 7 we derived the kinematic equations for a differential-drive robot. These give the robot's body-centric **velocity**  $\dot{v}_r \triangleq (\dot{x}_r, \dot{y}_r, \dot{\theta}_r)$  as a function of its wheel speeds  $(\dot{\phi}_l, \dot{\phi}_r)$ .

In the **noiseless** case, we can integrate  $\dot{v}_r$  to calculate the next pose  $x_{t+1}$  given the current pose  $x_t$ . (You will do this in Lab #2.)

We can model **noisy motion** by supposing that the robot's velocity  $\dot{v}_r$  is subject to noise. This leads to a very natural **sampling algorithm**:

**sample\_velocity\_motion\_model**( $x_t, \dot{\phi}_l, \dot{\phi}_r$ ):


1. Calculate body-centric velocity  $\dot{v}_r$  from  $(\dot{\phi}_l, \dot{\phi}_r)$  using kinematic model
2. **Add noise:**  $\tilde{v}_r \triangleq \dot{v}_r + \Delta v_r$ , where  $\Delta v_r$  is sampled from a noise distribution (e.g. Gaussian)
3. Calculate next pose  $x_{t+1}$  by integrating **noisy** velocity  $\tilde{v}_r$



$$\begin{pmatrix} \dot{x}_r \\ \dot{y}_r \\ \dot{\theta}_r \end{pmatrix} = \begin{pmatrix} \frac{r}{2}(\dot{\phi}_r + \dot{\phi}_l) \\ 0 \\ \frac{r}{w}(\dot{\phi}_r - \dot{\phi}_l) \end{pmatrix}$$

# The Banana Distribution

This velocity-based motion model induces a “banana”-shaped marginal distribution over position



In robotics, this is often referred to (appropriately enough) as the “banana distribution”

# Odometry-based motion models

Velocity-based motion models provide a simple means of simulating noisy robot motion.

**But:** They are often grossly oversimplified ( $\Rightarrow$  not very accurate)


**Alternative:** Treat an **odometry measurement**  $z_t \approx x_{t-1}^{-1} x_t$  between pose  $x_{t-1}$  and  $x_t$  **as if** it were a **motion command**  $u_{t-1}$ . That is, consider:

$$x_t \sim p(x_t | x_{t-1}, z_t)$$

## Payoffs:


- Often significantly more accurate than velocity models
- Odometry can be gotten from several sources:
  - Wheel encoders
  - Inertial measurement units
  - **Scan matching** or visual odometry

**Con:** Only available **post hoc** – can't use for planning



# Estimating odometry via scan matching

Suppose I have two scans of environment taken from nearby poses (so that they partially overlap)



**Question:** What is the coordinate transformation  $T_{21}$  that relates their overlap?

**Recall:** If  $T_{W1}$  and  $T_{W2}$  are poses of the scanner in the world frame  $W$ , and  $p$  is a point, then:

$$\begin{aligned} T_{W1}p_1 &= p_W = T_{W2}p_2 \\ \Rightarrow p_2 &= \underbrace{T_{W2}^{-1}T_{W1}}_{= T_{21}}p_1 \end{aligned}$$


**Punchline:** Scan matching provides a means of *directly measuring odometry*

**But:** How can we actually *find*  $T_{21}$  given only the two scans?

# Pointcloud registration

**Problem:** Given point sets  $X = \{x_i\}_{i=1}^n$  and  $Y = \{y_i\}_{i=1}^n$  in  $\mathbb{R}^d$ , find the transformation  $T = (t, R) \in SE(d)$  that *optimally aligns* them:

$$(t, R) = \operatorname{argmin} \sum_{i=1}^n \|y_i - (Rx_i + t)\|^2$$



# Pointcloud registration

**Given:** particle sets  $X = \{x_i\}_{i=1}^n$  and  $Y = \{y_i\}_{i=1}^n$  in  $\mathbb{R}^d$

## **Horn's method:**

1. Calculate centroids:

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i, \quad \bar{y} = \frac{1}{n} \sum_{i=1}^n y_i,$$

2. Center the pointclouds:  $x'_i \triangleq x_i - \bar{x}$ ,  $y'_i \triangleq y_i - \bar{y}$
3. Construct outer product matrix  $W$ :

$$W \triangleq \sum_i y'_i x_i'^T$$

4. Recover optimal rotation  $R$ :

$$R = U \cdot \text{diag}(1, \dots, 1, \det(UV^T)) \cdot V^T$$

where  $W = U\Sigma V^T$  is the singular value decomposition of  $W$

5. Recover optimal translation  $t$ :  $t = \bar{y} - R\bar{x}$

# Iterative Closest Point

**Catch:** Horn's method assumes that we know the *point correspondences*  $x_i \leftrightarrow y_i$ . What if we don't have these?

**One approach:** Estimate these together with the registration  $T = (t, R)$ !

**Iterative closest point (ICP):** Given  $X = \{x_i\}_{i=1}^n$  and  $Y = \{y_j\}_{j=1}^m$ , initial guess  $T = (t_0, R_0)$ , maximum association distance  $d > 0$ :


For  $k = 0, \dots$  until convergence:

1. **Estimate matches:** Let  $y_{l_i} = \operatorname{argmin}_j \|y_j - (R_k x_i + t_k)\|$  be the closest point in  $Y$  to the image of  $x_i$  under the current registration estimate  $T_k = (t_k, R_k)$ .

If  $\|y_{l_i} - (R_k x_i + t_k)\| \leq d$ , accept  $x_i \leftrightarrow y_{l_i}$  as a match

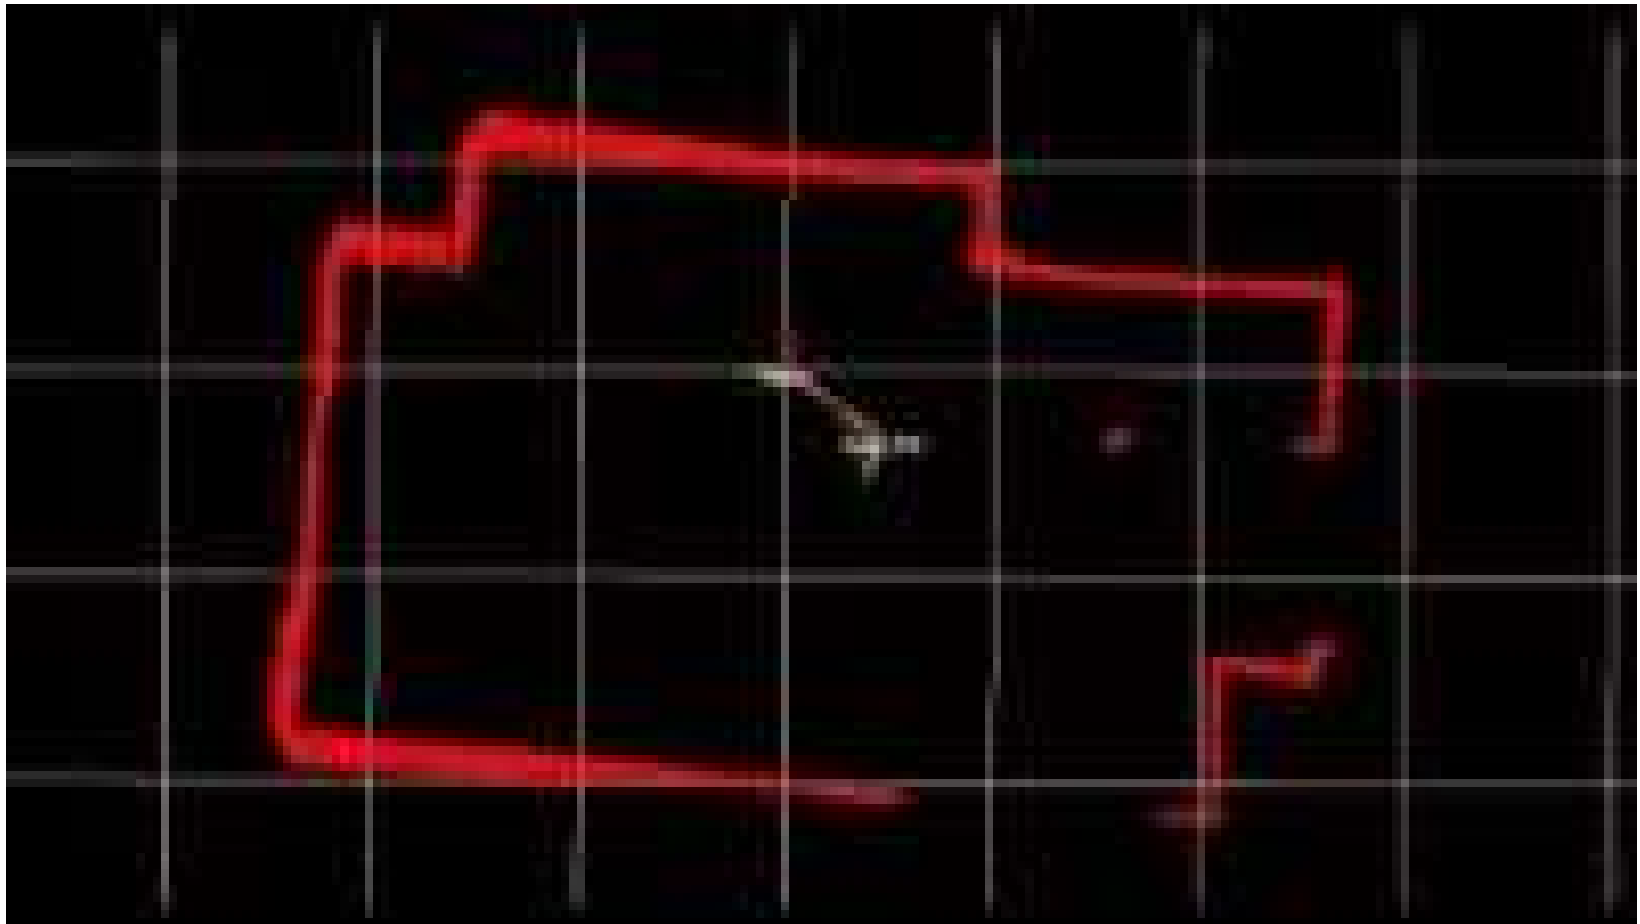
2. **Estimate registration:** Compute next registration estimate  $T_{k+1} = (t_{k+1}, R_{k+1})$  by applying Horn's method to matches  $\{x_i \leftrightarrow y_{l_i}\}$ .

# Iterative closest point





## Odometry estimation using ROS's Canonical Scan Matcher




# Monte Carlo Localization

**Main idea:** Monte Carlo localization (MCL) is simply localization using the Particle Filter

**Given:** Prior  $p(x_0)$ , map  $m$ , controls  $u_{0:t}$ , observations  $z_{1:t}$

**Estimate:** Belief  $p(x_t | m, z_{1:t})$  over the robot pose

```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):  
2:     $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$   
3:    for  $m = 1$  to  $M$  do  
4:      sample  $x_t^{[m]} \sim p(x_t | u_t, x_{t-1}^{[m]})$  ← Motion model  
5:       $w_t^{[m]} = p(z_t | x_t^{[m]})$  ← Sensing model  
6:       $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$   
7:    endfor  
8:    for  $m = 1$  to  $M$  do  
9:      draw  $i$  with probability  $\propto w_t^{[i]}$   
10:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$   
11:    endfor  
12:    return  $\mathcal{X}_t$ 
```



## Requires:

- A *sampler* for the motion model
- *Likelihood function* for the sensor model

# Beam sensor likelihood model

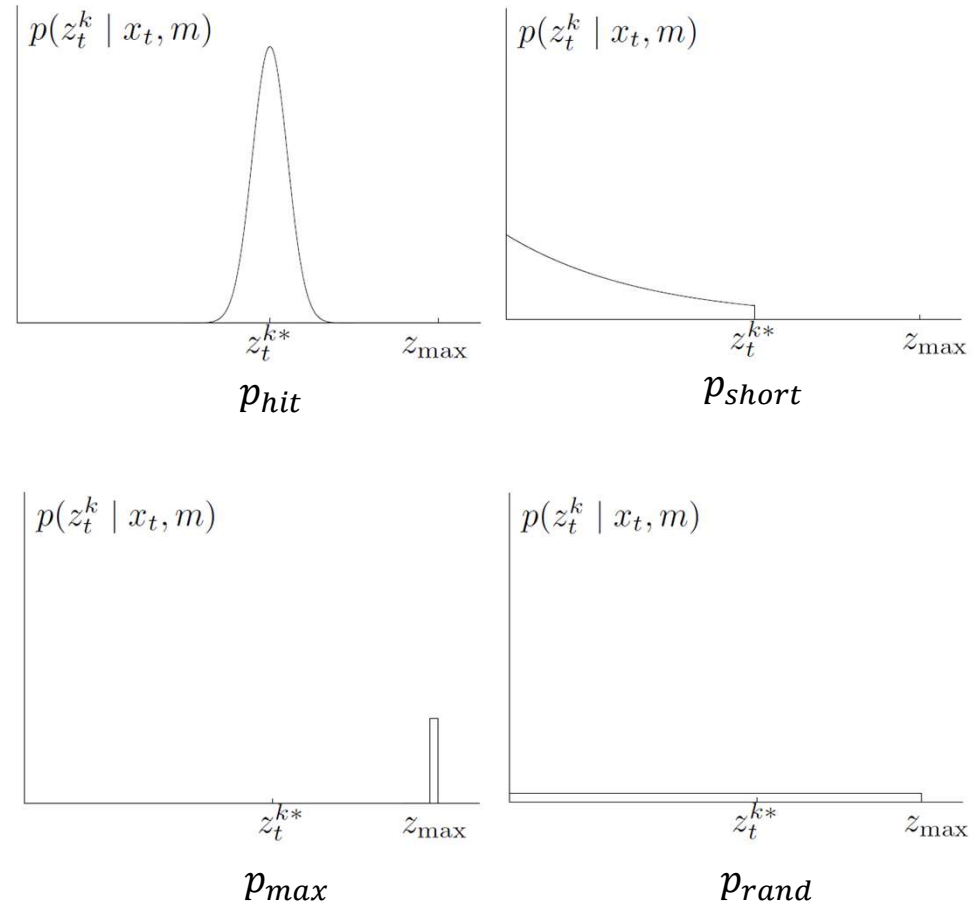
**Last time:** (Approximate) *inverse* sensor model:  $p(m_i|x_t, z_t)$   
⇒ Approximate likelihood of *map* given *pose and measurement*

**Today:** *Forward* sensor model  $p(z|x_t, m)$   
⇒ Likelihood of *measurement* given *pose and map*

# Beam sensor likelihood model

We consider a more physically-faithful sensor model that accounts for **4 distinct sources** of error

- **Correct return w/ small measurement noise**
  - Nominal case for operation
  - Modeled as a **Gaussian**  $p_{hit}$  centered on true range
- **Unexpected / transient objects**
  - Corresponds to e.g. people moving in the scene
  - Causes a **short return**
  - Modeled as an **exponential** distribution  $p_{short}$  (one can justify this mathematically by appealing to e.g. survival analysis).
- **Failed detection**
  - Sensor does not detect a reflected beam – can be caused by e.g. specular reflection or absorption of the beam on target
  - Appears as a **maximum range** return
  - Can be modeled as a **point mass** (or very narrowly peaked distribution)  $p_{max}$  at the sensor's maximum range
- **Otherwise unexplained**
  - Catch-all category for other general weirdness
  - Modeled as a **uniform distribution**  $p_{rand}$  over the sensor's range



# Beam sensor likelihood model


Our overall beam sensor model  $p(z|x_t, m)$  is a *mixture* of these 4 components:

$$p(z|x_t, m) = w_{hit} p_{hit} + w_{short} p_{short} + w_{max} p_{max} + w_{rand} p_{rand}$$

where  $w_{hit}$ ,  $w_{short}$ ,  $w_{max}$ ,  $w_{rand}$  are the *mixture weights* (nonnegative and sum to 1).

We also assume that each beam is sampled *independently*, so the joint probability for an entire ( $n$ -beam) scan is:

$$p(z_t|x_t, m) = \prod_{k=1}^n p(z_t^k|x_t, m)$$




Complete (mixture) beam sensor model  $p(z|x_t, m)$

# Beam sensor measurement model



Laser scan (robot's true position)



Heatmap for scan likelihood as a function of position


# Monte Carlo Localization

**Main idea:** Monte Carlo localization (MCL) is simply localization using the Particle Filter

**Given:** Prior  $p(x_0)$ , map  $m$ , controls  $u_{0:t}$ , observations  $z_{1:t}$

**Estimate:** Belief  $p(x_t | m, z_{1:t})$  over the robot pose


```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):  
2:     $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$   
3:    for  $m = 1$  to  $M$  do  
4:      sample  $x_t^{[m]} \sim p(x_t | u_t, x_{t-1}^{[m]})$  ← Motion model  
5:       $w_t^{[m]} = p(z_t | x_t^{[m]})$  ← Sensing model  
6:       $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$   
7:    endfor  
8:    for  $m = 1$  to  $M$  do  
9:      draw  $i$  with probability  $\propto w_t^{[i]}$   
10:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$   
11:    endfor  
12:    return  $\mathcal{X}_t$ 
```



## Requires:

- A *sampler* for the motion model
- *Likelihood function* for the sensor model

# Monte Carlo localization





## Monte Carlo localization on a RACECAR



# Practicalities

Recall (from Lecture 10) two major issues with particle filters:

- **Particle depletion:** Need to preserve diversity in the particle set  
⇒ **Augmented MCL:** Method for increasing particle diversity
- **Computational cost:** Might need a **lot** of particles to model complex beliefs  
⇒ **Adaptive sampling:** *Dynamically* adjust particle set size

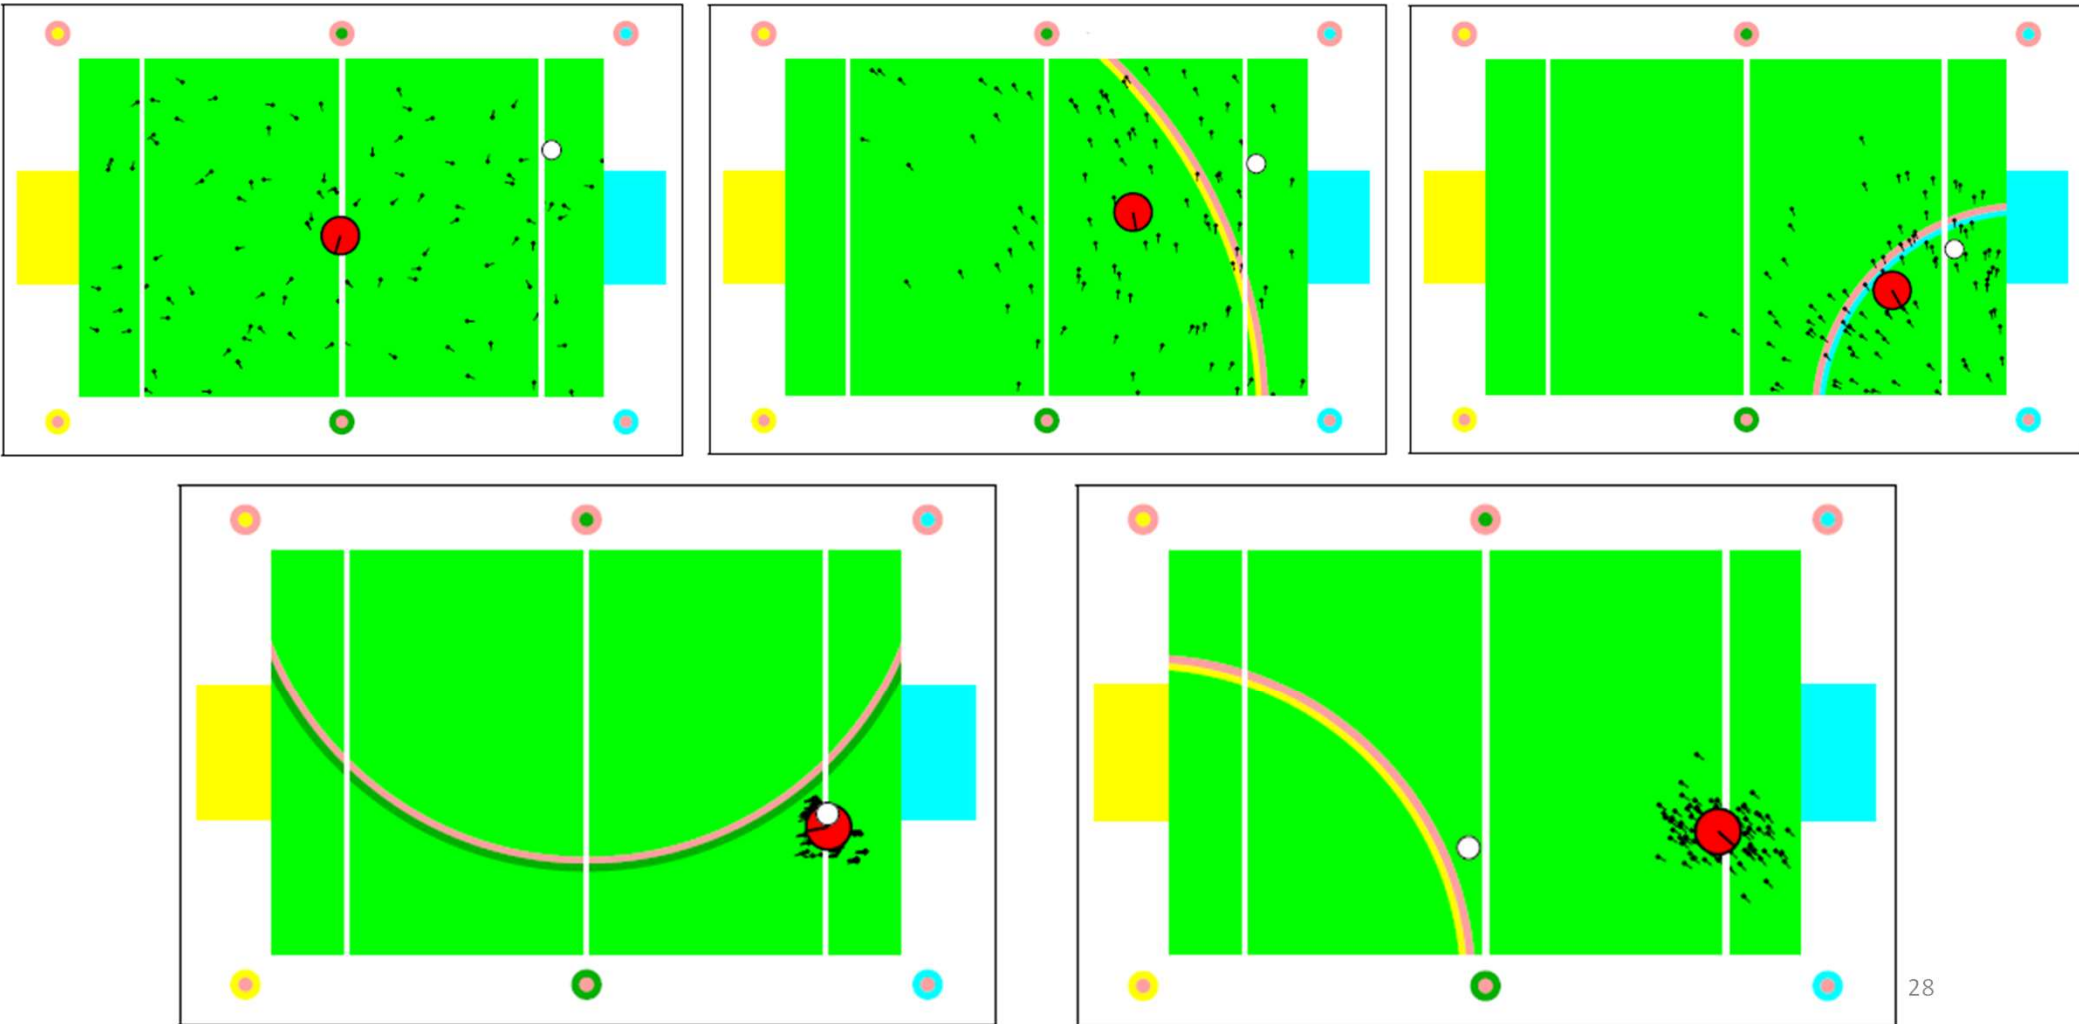
# The “kidnapped robot” problem

**Def:** A scenario in which a robot must relocalize itself after being (instantly) moved to an *arbitrary* location in a map

**Ex:** Repositioning a robot that is powered off (Moving a Roomba up / down a floor)



# Particle depletion and the kidnapped robots



# Preserving particle diversity: Augmented MCL

**Problem:** We need to preserve *diversity* in the filter's particle set

⇒ Need to have particles near the true state to recover from localization failures

**Idea:** We can *increase diversity* by resampling a small fraction of (*purely*) *random* particles

**Approach:** Use *measurement likelihood* (weights) as a measure of localization quality





⇒ Sudden decrease in avg particle likelihood vs. historical avg may indicate localization failure

⇒ In that case, we should increase diversity of particle set

```
1: Algorithm Augmented_MCL( $\mathcal{X}_{t-1}, u_t, z_t, m$ ):
2:   static  $w_{\text{slow}}, w_{\text{fast}}$ 
3:    $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$ 
4:   for  $m = 1$  to  $M$  do
5:      $x_t^{[m]} = \text{sample\_motion\_model}(u_t, x_{t-1}^{[m]})$ 
6:      $w_t^{[m]} = \text{measurement\_model}(z_t, x_t^{[m]}, m)$ 
7:      $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$ 
8:      $w_{\text{avg}} = w_{\text{avg}} + \frac{1}{M} w_t^{[m]}$ 
9:   endfor
10:   $w_{\text{slow}} = w_{\text{slow}} + \alpha_{\text{slow}}(w_{\text{avg}} - w_{\text{slow}})$ 
11:   $w_{\text{fast}} = w_{\text{fast}} + \alpha_{\text{fast}}(w_{\text{avg}} - w_{\text{fast}})$ 
12:  for  $m = 1$  to  $M$  do
13:    with probability  $\max(0.0, 1.0 - w_{\text{fast}}/w_{\text{slow}})$  do
14:      add random pose to  $\mathcal{X}_t$ 
15:    else
16:      draw  $i \in \{1, \dots, N\}$  with probability  $\propto w_t^{[i]}$ 
17:      add  $x_t^{[i]}$  to  $\mathcal{X}_t$ 
18:    endwith
19:  endfor
20:  return  $\mathcal{X}_t$ 
```

**Table 8.3** An adaptive variant of MCL that adds random samples. The number of random samples is determined by comparing the short-term with the long-term likelihood of sensor measurements.

## Example: Kidnapped robot revisited



# KLD-sampling

**KLD-sampling** is a method for *dynamically adjusting* the number of particles that we maintain in the sample set

**Main idea:** The number of particles that we need is related to the *uncertainty* in our current belief:

More uncertainty  $\Rightarrow$  belief is more “spread out”  $\Rightarrow$  need more particles to model

**Approach:** *Dynamically adjust* the number of particles based upon an estimate of the **Kullback-Leibler divergence** between particle set and true posterior



# KLD-Sampling

```

1:  Algorithm KLD_Sampling_MCL( $\mathcal{X}_{t-1}, u_t, z_t, m, \varepsilon, \delta$ ):
2:     $\mathcal{X}_t = \emptyset$ 
3:     $M = 0, M_\chi = \infty, k = 0$ 
4:    for all  $b$  in  $H$  do
5:       $b = \text{empty}$ 
6:    endfor
7:    do
8:      draw  $i$  with probability  $\propto w_{t-1}^{[i]}$ 
9:       $x_t^{[M]} = \text{sample\_motion\_model}(u_t, x_{t-1}^{[i]})$ 
10:      $w_t^{[M]} = \text{measurement\_model}(z_t, x_t^{[M]}, m)$ 
11:      $\mathcal{X}_t = \mathcal{X}_t + \langle x_t^{[M]}, w_t^{[M]} \rangle$ 
12:     if  $(x_t^{[M]}$  falls into empty bin  $b$ ) then
13:        $k = k + 1$ 
14:        $b = \text{non-empty}$ 
15:       if  $(k > 1)$  then
16:          $M_\chi := \frac{k-1}{2\varepsilon} \left\{ 1 - \frac{2}{9(k-1)} + \sqrt{\frac{2}{9(k-1)}} z_{1-\delta} \right\}^3$ 
17:       endif
18:        $M = M + 1$ 
19:     while  $(M < M_\chi)$ 
20:     return  $\mathcal{X}_t$ 

```

