


# EECE 5550: Mobile Robotics




## Lecture 13: Simultaneous Localization and Mapping

# Recap: Two fundamental problems in robot perception

## Localization: Where am I?

**Given:** Prior  $p(x_0)$ , map  $m$ , controls  $u_{0:t-1}$ , measurements  $z_{1:t}$


**Estimate:** Belief  $p(x_t | m, u_{0:t-1}, z_{1:t})$  over robot pose



## Mapping: What's around me?

**Given:** Robot poses  $x_{0:t}$ , measurements  $z_{1:t}$


**Estimate:** Belief  $p(m | x_{0:t}, z_{1:t})$  over the map  $M$




# Simultaneous Localization and Mapping (SLAM): The Big One

**Given:** Prior  $p(x_0)$ , controls  $u_{0:t-1}$ , measurements  $z_{1:t}$

**Estimate:** Joint belief  $p(x_{0:t}, m | u_{0:t-1}, z_{1:t})$  over sequence of robot poses  $x_{0:t}$  and map




# Simultaneous Localization and Mapping (SLAM)




- **Much** harder than localization or mapping alone
- Enables operation in *unknown environments* (exploration)
- An *essential enabling technology* for mobile robots

# Simultaneous Localization and Mapping (SLAM): The Big One



# Plan of the day

- Motivating example
- Factor graphs
- SLAM problem formulation
- Solving the SLAM problem via maximum-likelihood estimation
- Anatomy of a modern SLAM system
- Practicalities



$$p(Z|\Theta) = \prod_i p_i(Z_i | \Theta_i)$$

$$\Theta_i = \{\theta_j \in \Theta \mid (p_i, \theta_j) \in E\}$$

# References

Foundations and Trends® in Robotics  
Vol. 6, No. 1-2 (2017) 1–139  
© 2017 F. Dellaert and M. Kaess  
DOI: 10.1561/23000000043



## Factor Graphs for Robot Perception

Frank Dellaert  
Georgia Institute of Technology  
dellaert@cc.gatech.edu


Michael Kaess  
Carnegie Mellon University  
kaess@cmu.edu

## Papers

- “Factor Graphs for Robot Perception”
- “Factor Graphs and GTSAM: A Hands-On Introduction”
- “ORB-SLAM2: An Open-Source SLAM System for Monocular, Stereo, and RGB-D Cameras”
- “Bags of Binary Words for Fast Place Recognition in Image Sequences”





# A concrete example

Consider a robot exploring some initially unknown environment ...










# A concrete example



# A concrete example



# As the robot explores




We build up a *graph* of *noisy spatial relations* ...

# Factor graphs

A *factor graph*  $G = (\Theta, F, E)$  is a bipartite graph that models the factorization of a function  $f: \Omega \rightarrow \mathbb{R}$ .

$$f(\Theta) = \prod_i f_i(\Theta_i)$$

$$\Theta_i = \{\theta_j \in \Theta \mid (f_i, \theta_j) \in E\}$$



Here:


- $\Theta$  is the set of *variable nodes*
- $F$  is the set of *factor nodes*
- $E$  is the edge set.  $G$  has an edge  $e_{ij} = (f_i, \theta_j)$  if and only if variable  $\theta_j$  is an argument of factor  $f_i$ .

# The SLAM estimation problem

**Given:** A factor graph representation  $G = (\Theta, F, E)$  of the **joint distribution** for the network of noisy spatial relations:

$$p(Z|\Theta) = \prod_i p_i(Z_i | \Theta_i)$$

$$\Theta_i = \{\theta_j \in \Theta \mid (p_i, \theta_j) \in E\}$$



**Find:** The **maximum likelihood estimate**  $\hat{\Theta}_{MLE}$  for the variables  $\Theta$ :

$$\hat{\Theta}_{MLE} = \operatorname{argmin}_{\Theta} \sum -\log p_i(Z_i | \Theta_i)$$


$\Rightarrow$  This is the **point estimate** that **best explains** the available data.

# Key features of the SLAM inference problem

The SLAM problem is:

- Nonlinear
- High-dimensional
- Nonconvex

⇒ This is a challenging optimization problem




$$\hat{\Theta}_{ML} = \operatorname{argmin}_{\Theta} \sum -\log p_i(Z_i | \Theta_i)$$

**However:** It is also *sparse*.

⇒ This enables *efficient maximum likelihood estimation*

# The “standard model” of SLAM

$$\hat{\Theta}_{ML} = \operatorname{argmin}_{\Theta} \sum -\log p_i(Z_i|\Theta_i)$$



Let's **assume** that each factor  $p_i$  models a nonlinear measurement  $z_i = h(\Theta_i) + \epsilon_i$ , where  $\epsilon_i \sim N(0, \Sigma_i)$  is **additive Gaussian noise**.

Then:

$$p_i(Z_i|\Theta_i) \propto \exp\left(-\frac{1}{2} \|z_i - h_i(\Theta_i)\|_{\Sigma_i}^2\right)$$

⇒ Maximum likelihood inference is equivalent to a **sparse nonlinear least-squares** (NLS) problem:

$$\hat{\Theta}_{MLE} = \operatorname{argmin}_{\Theta} \sum_i \|z_i - h_i(\Theta_i)\|_{\Sigma_i}^2$$


**Payoff:** Sparse NLS problems can be processed **very efficiently**. (More on this next time ...)

# Software for solving sparse NLS problems

Current state-of-the-art SLAM approaches are all based upon **sparse nonlinear least-squares estimation over factor graphs**.

Software libraries:

- Ceres
- iSAM / GTSAM
- g2o






# Example: GTSAM

## Constructing the factor graph in GTSAM

A simple (toy) factor graph




```
1 // Create an empty nonlinear factor graph
2 NonlinearFactorGraph graph;
3
4 // Add a Gaussian prior on pose x_1
5 Pose2 priorMean(0.0, 0.0, 0.0);
6 noiseModel::Diagonal::shared_ptr priorNoise =
7     noiseModel::Diagonal::Sigmas(Vector_(3, 0.3, 0.3, 0.1));
8 graph.add(PriorFactor<Pose2>(1, priorMean, priorNoise));
9
10 // Add two odometry factors
11 Pose2 odometry(2.0, 0.0, 0.0);
12 noiseModel::Diagonal::shared_ptr odometryNoise =
13     noiseModel::Diagonal::Sigmas(Vector_(3, 0.2, 0.2, 0.1));
14 graph.add(BetweenFactor<Pose2>(1, 2, odometry, odometryNoise));
15 graph.add(BetweenFactor<Pose2>(2, 3, odometry, odometryNoise));
```

Listing 1: Excerpt from examples/OdometryExample.cpp

# Example: GTSAM

## Optimizing the factor graph in GTSAM



A simple (toy) factor graph

```
1 // create (deliberately inaccurate) initial estimate
2 Values initial;
3 initial.insert(1, Pose2(0.5, 0.0, 0.2));
4 initial.insert(2, Pose2(2.3, 0.1, -0.2));
5 initial.insert(3, Pose2(4.1, 0.1, 0.1));
6
7 // optimize using Levenberg-Marquardt optimization
8 Values result = LevenbergMarquardtOptimizer(graph, initial).optimize();
```

Listing 2: Excerpt from examples/OdometryExample.cpp

Initial Estimate:

Values with 3 values:

Value 1: (0.5, 0, 0.2)

Value 2: (2.3, 0.1, -0.2)

Value 3: (4.1, 0.1, 0.1)

Final Result:

Values with 3 values:

Value 1: (-1.8e-16, 8.7e-18, -9.1e-19)

Value 2: (2, 7.4e-18, -2.5e-18)

Value 3: (4, -1.8e-18, -3.1e-18)

# The SLAM estimation problem

## Main takeaways:


- Factor graphs provide a *simple*, *flexible*, and *elegant* language for modeling machine perception problems
- Under the assumption of **additive Gaussian noise**:

$$z_i = h(\Theta_i) + \epsilon_i, \text{ where } \epsilon_i \sim N(0, \Sigma_i)$$

**maximum likelihood estimation** reduces to **nonlinear least-squares**:

$$\hat{\Theta}_{ML} = \underset{\Theta}{\operatorname{argmin}} \sum_i \|z_i - h_i(\Theta_i)\|_{\Sigma_i}^2$$

- We can process large-scale but *sparse* NLS problems **very efficiently**



$$p(Z|\Theta) = \prod_i p_i(Z_i | \Theta_i)$$


$$\Theta_i = \{\theta_j \in \Theta \mid (p_i, \theta_j) \in E\}$$

# SLAM practicalities


**Recap:** Factor graph models + nonlinear least-squares estimation provide a general and effective means of solving large-scale geometric estimation problems.

$$p(Z|\Theta) = \prod_i p_i(Z_i | \Theta_i)$$



$$\Theta_i = \{\theta_j \in \Theta \mid (p_i, \theta_j) \in E\}$$



**BUT:** How are we supposed to **obtain** these factor graph models??



Raw sensor data



Factor graph model

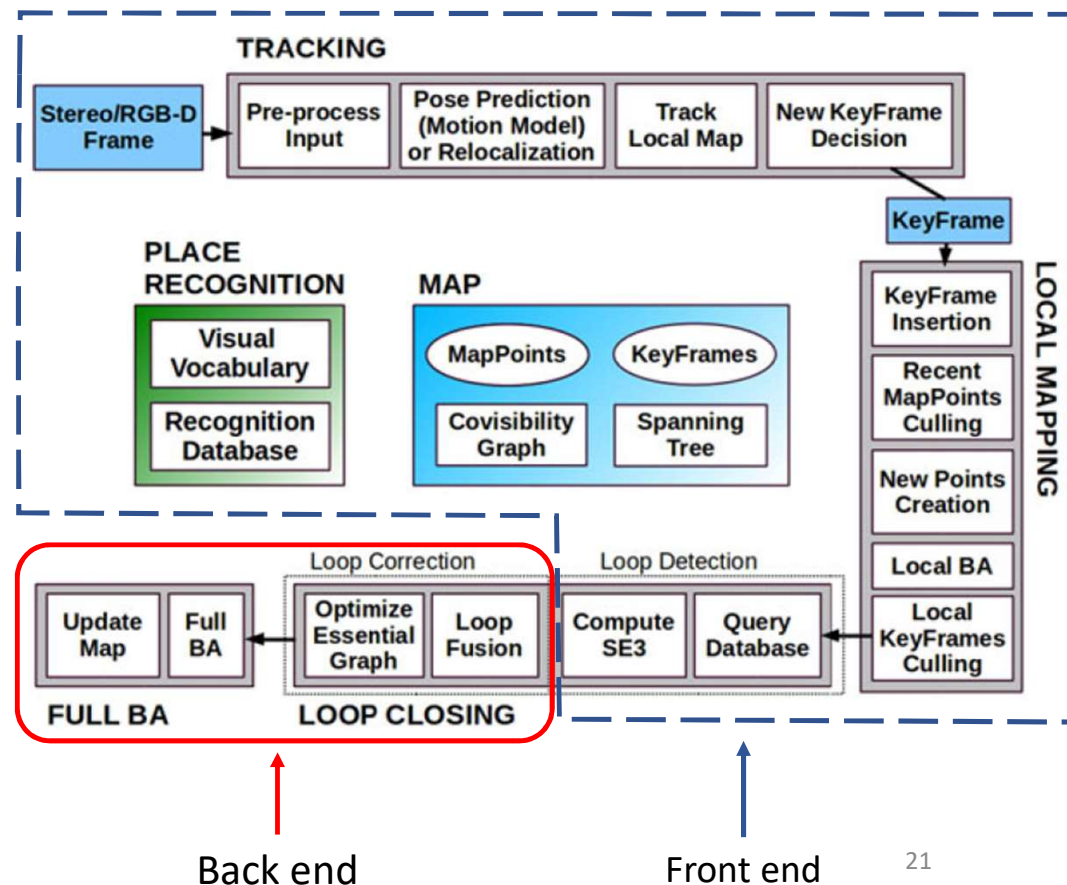
# Anatomy of a modern SLAM system

Modern SLAM systems are typically composed of two components:

- **Front end:** Build factor graph model of the SLAM estimation problem from sensor data
  - Feature extraction
  - Data association
- **Back end:** Perform optimization over the factor graph to recover maximum likelihood estimate  $\hat{\Theta}_{MLE}$  for SLAM solution

**NB:** The majority of a SLAM system's complexity is devoted to constructing the factor graph model!

ORB-SLAM2 system architecture



# Feature extraction

**Main idea:** Process the raw sensor data to extract specific features / measurements / entities that will appear in the back-end factor graph model

## Examples:

- **Feature points** and their descriptors
- **Keyframes** (camera poses + image data)
- **Objects**



ORB feature points

**NB:** ORB-SLAM2 uses **projective** (camera) observations of **3D points**, extracted from images using the **ORB feature detector**

Each detected **image feature**  $z_i \in \mathbb{R}^2$  gives rise to a **factor** of the form  $p(z_i | t_i, x_i)$ , where:

- $t_i \in \mathbb{R}^3$  is the position of the **3D point** that produced the feature  $z_i$
- $x_i \in SE(3)$  is the **pose** of the camera from which the image was taken





# Data association

**Main idea:** Feature extraction can identify *some* object of interest in raw sensor data. However, in order to construct a factor graph, we must also know *which one* it is.

**Data association** is the problem of associating *observations* (in our sensor data) with the *entities* (in the world) that produced them.

**Ex:** In imagery, data association amounts to deciding which *2D features* (in the image) correspond to the same *3D point* (in the world)

**NB:** These (estimated!) *associations* determine the *edge set* in our factor graph



# Data association: Easy and Hard Cases

**Easy case: *Feature tracking*** It's (relatively) easy to track features over a *short sequence* of measurements. Since the sensor is not moving far, we have a good idea of *where to look*





# Data association: Easy and Hard Cases


**Hard case: Loop closures** Conversely, if we **revisit** a location after traveling a long distance (i.e. if we *close a loop*), we may have **high uncertainty** in our position (due to *drift*)

- High pose uncertainty  $\Rightarrow$  **no strong constraint** on *which features to consider*
- **Large changes in view** can make identifying correspondences much harder



**BUT:** Loop closures are **essential for correcting drift!**

# Loop closures correct drift




# Finding loop closures: Place recognition

**Problem:** How can we **recognize** if we have visited a place before?

**Main idea:** Think of this as a **search** problem: try to find a **previous view** that matches the **current view**.

Query view



Database of previously-visited locations

# Bag-of-Words Place Recognition

*Bag-of-words* place recognition is one of the most common approaches to place recognition.


**Main idea:** Treat images as “bags” of “*visual words*” (set of [quantized] descriptors extracted from feature points in the image).

**Then:** we can treat *place recognition* as *text retrieval*!

- Image feature  $\Leftrightarrow$  “word” in a vocabulary
- Image (list of visual words)  $\Leftrightarrow$  Document
- Image similarity search  $\Leftrightarrow$  Similar document retrieval


**Operationally**, given a query image, we:

- **Extract visual features**, and map to quantized vocabulary  
 $\Rightarrow$  This gives a representation of our image as a *sparse feature vector*  $x$ , in which  $x_k$  is the *count* of the  $k$ th word in the image
- **Retrieve similar images** by finding similar feature vectors in a database



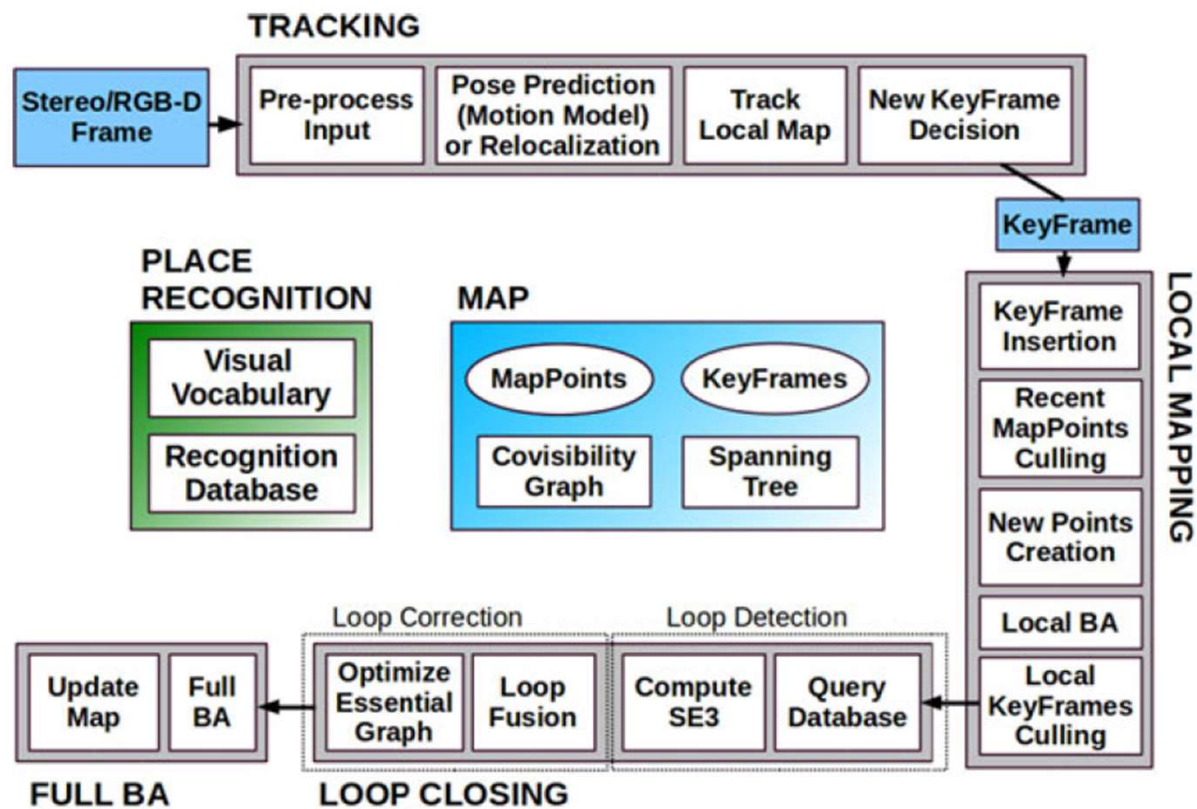
Vocabulary: quantization of visual descriptors

# Bag of words place recognition




# Putting it all together

ORB-SLAM2 system architecture



# Putting it all together





# Summary

- Factor graphs provide a *simple*, *flexible*, and *elegant* language for modeling machine perception problems
- Under the assumption of *additive Gaussian noise*:

$$z_i = h(\Theta_i) + \epsilon_i, \text{ where } \epsilon_i \sim N(0, \Sigma_i)$$

Maximum likelihood estimation reduces to NLS:


$$\hat{\Theta}_{ML} = \underset{\Theta}{\operatorname{argmin}} \sum_i \|z_i - h_i(\Theta_i)\|_{\Sigma_i}^2$$

- We can process *sparse* NLS problems *very efficiently*

**BUT:** A SLAM *system* also needs to address the *front-end*

- Feature extraction
- Data association

⇒ These account for most of the (practical) complexity!



$$p(Z|\Theta) = \prod_i p_i(Z_i | \Theta_i)$$

$$\Theta_i = \{\theta_j \in \Theta \mid (p_i, \theta_j) \in E\}$$

