



# Planning I

## - getting from A to B

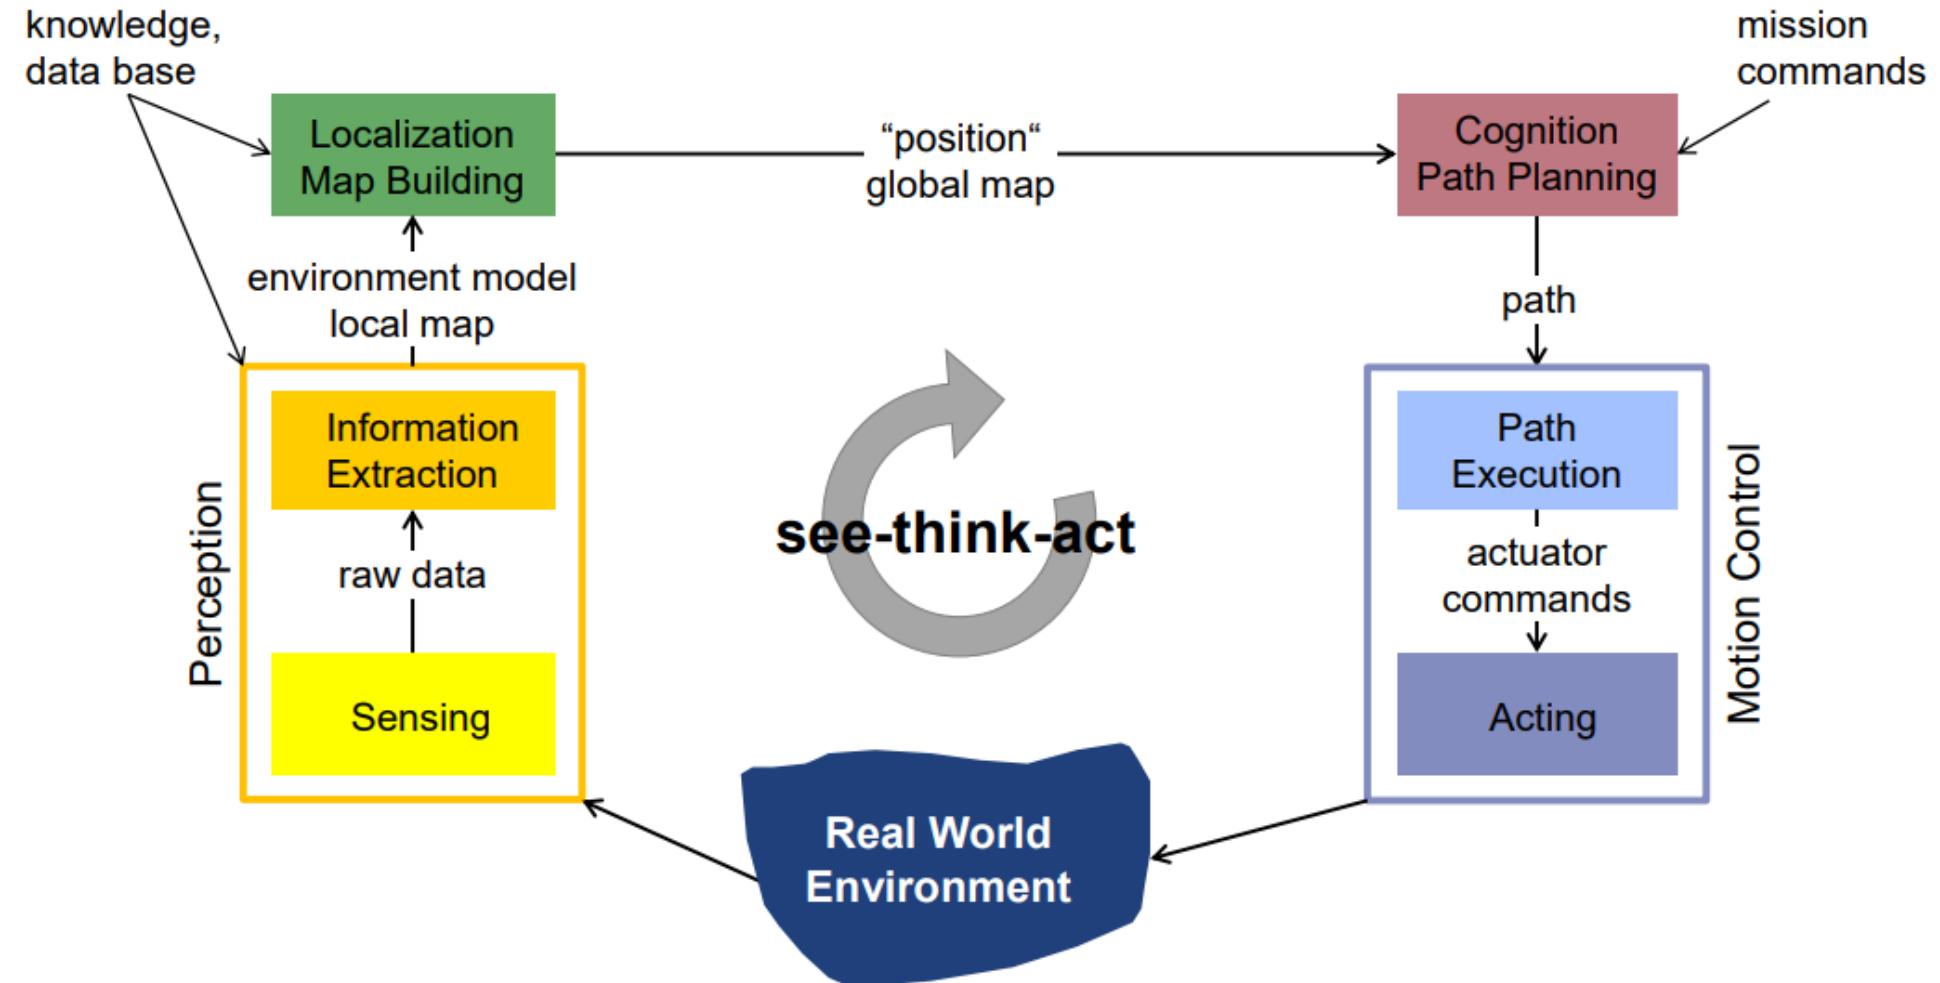
Roland Siegwart, Margarita Chli, Nick Lawrance

Additional thanks to Jen Jen Chung and Juan Nieto for additional slide content

# Learning objectives

- **Goal:** Be able to describe and apply planning algorithms for mobile robots
- **Competencies:**
  - Analyse a robot planning problem and describe suitable planning hierarchies, representations and algorithms for planning
  - Apply fundamental algorithms in planning such as graph search (BFS, Dijkstra, A\*), and collision avoidance (velocity obstacles, dynamic window) and hybrid methods (potential fields)

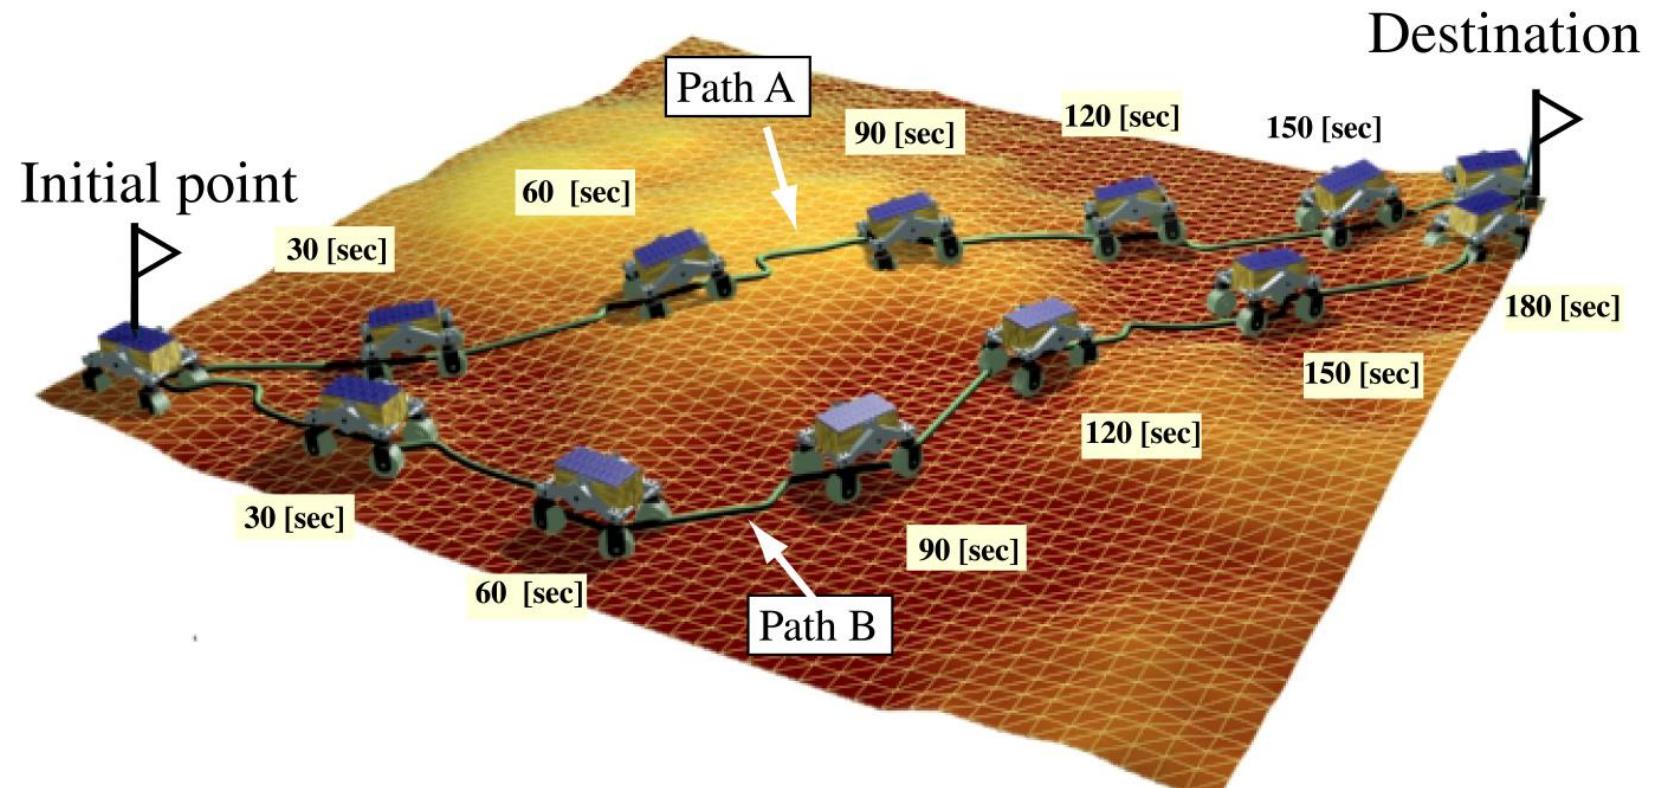
# Introduction – The see-think-act cycle



# What does planning mean?

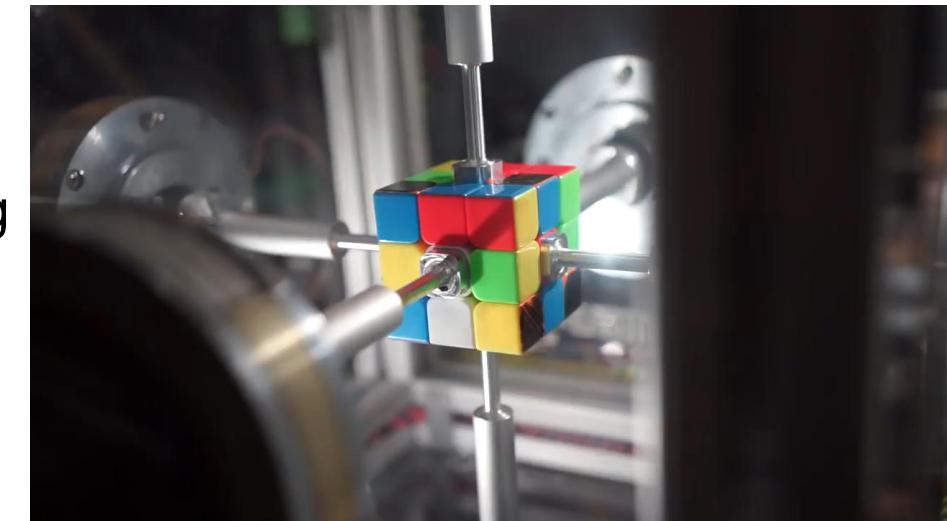
- Different things to different people
- In general, we will focus on ***motion planning*** for robotics, namely determining a set of actions to take a robot from one known state to another known state
- We are also interested in considering some of the ***constraints*** of the platform, ensuring that our generated plans are ***feasible***

# What is (robot) Planning?



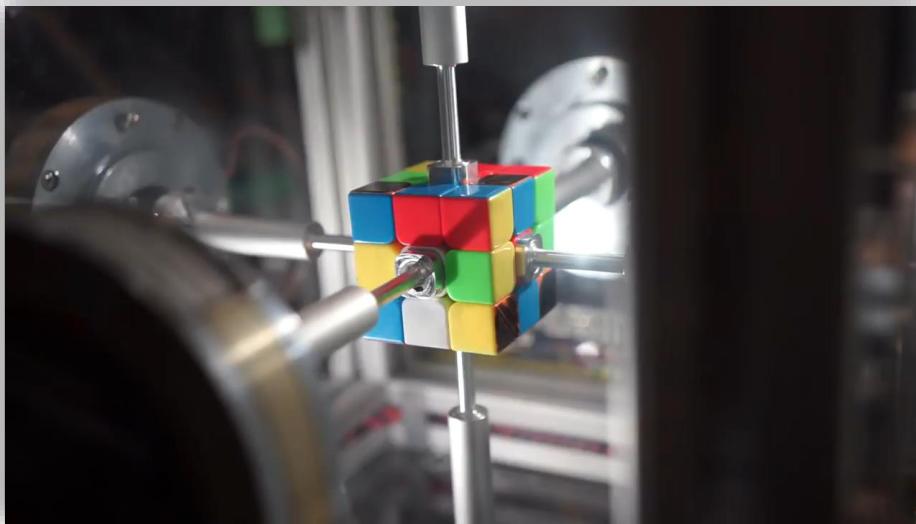
# Related topics

- Control:
  - Generally concerned with reaching and maintaining a desired state in some kind of robust way
  - Often **feedback-based**
  - Success measured in terms of stability, robustness, ability to reject disturbance
- Planning in Computer Science:
  - Generally more focused on **discrete** problems
  - Classic AI planning problems (spanning graphs, travelling salesman, orienteering) often appear in robotic planning approaches



Used with permission of Ben Katz

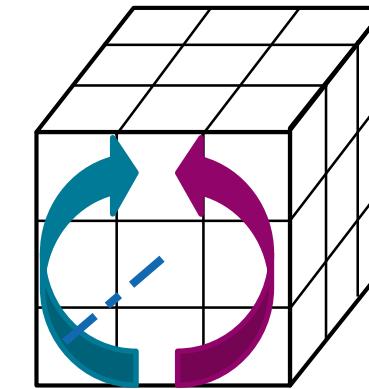
# Discrete planning



Used with permission of Ben Katz

Set of all possible cube configurations  $s \in S$

Set of all possible actions  $a \in A = \{FL, FR, \dots\}$

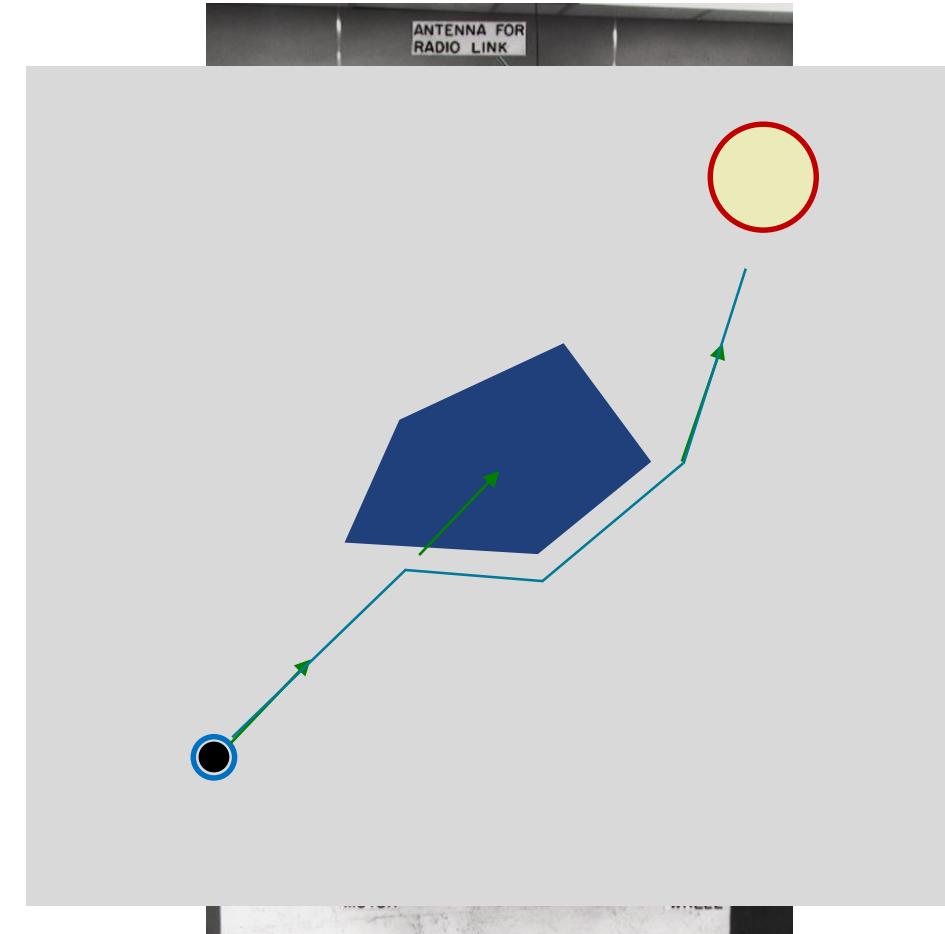


$s_k, a_k \rightarrow s_{k+1}$

**Goal:** Generate a plan  $p_{s_0}$  - a sequence of actions  $[a_0, a_1, \dots, a_n]$  that takes the system from the current state  $s_0$  to the goal state  $s_g$  with the minimum number of actions  $n$

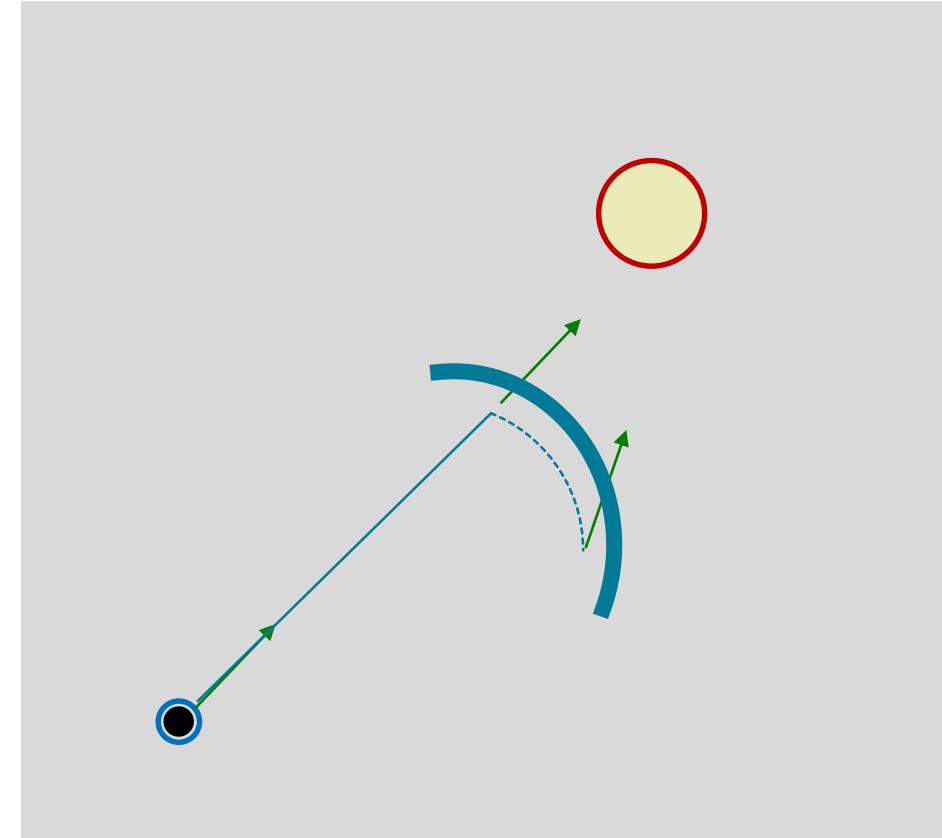
# Historical robot motion planning – Adapted from Howie Choset

- Classical robotics (mid-70s)
  - Exact models, no sensing
- Reactive motion planning (mid-80s)
  - No models, rely on sensors only



# Historical robot motion planning – Adapted from Howie Choset

- Classical robotics (mid-70s)
  - Exact models, no sensing
- Reactive motion planning (mid-80s)
  - No models, rely on sensors only



# Historical robot motion planning – Adapted from Howie Choset

- Classical robotics (mid-70s)
  - Exact models, no sensing
- Reactive motion planning (mid-80s)
  - No models, rely on sensors only
- Hybrid / hierarchical (since early 90s)
  - Use models/planning at high level
  - Reactive (obstacle avoidance) at lower level
- Probabilistic (since mid 90s)
  - Incorporate uncertainty in models and sensors in all stages of planning

# MOTION PLANNING

# Navigation Competence

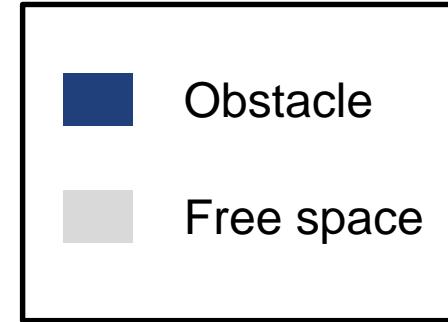
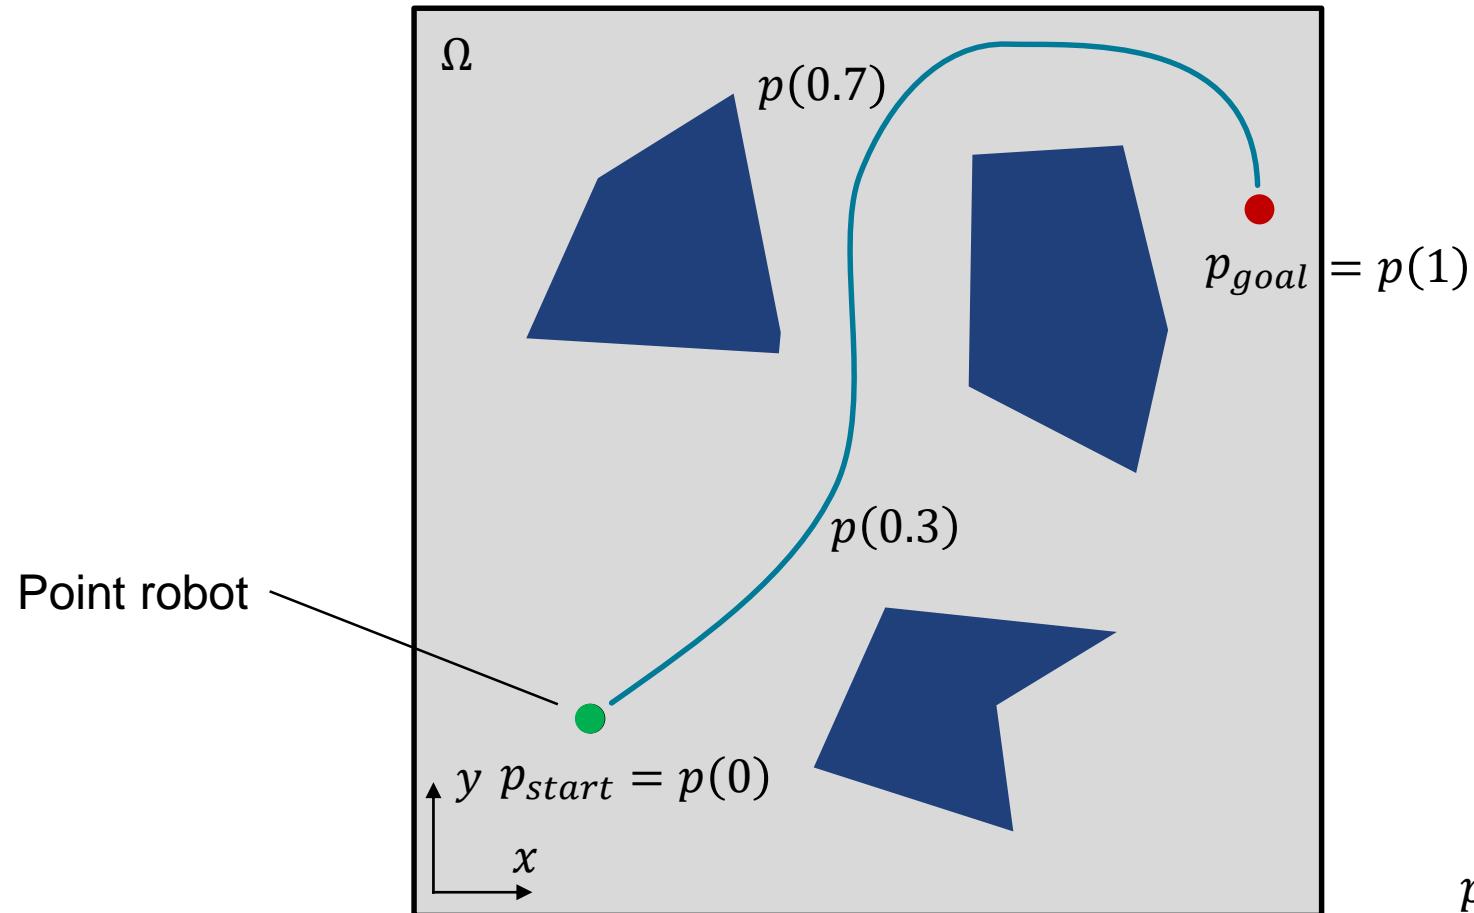
- In a nutshell, work out how the robot could *feasibly* move from one position to another
- Simplifying assumptions (for now...)
  - Our **representation** of the robot and the world is sufficiently expressive
  - We know **where we are** and **where we want to go**
  - We have a **motion model** for our robot
- Typically cast as an optimisation problem – minimise cost (time, distance, energy), within constraints

# REPRESENTATION

# Representation

- How the world is represented and understood by the planner (robot) is important
- Usually some degree of simplification in choosing a representation
- By choosing a suitable representation of the world, we may be able to apply existing algorithms to solve our planning problem

# Representation – workspace and paths



$$\omega_{free} = \Omega - \omega_{obs}$$

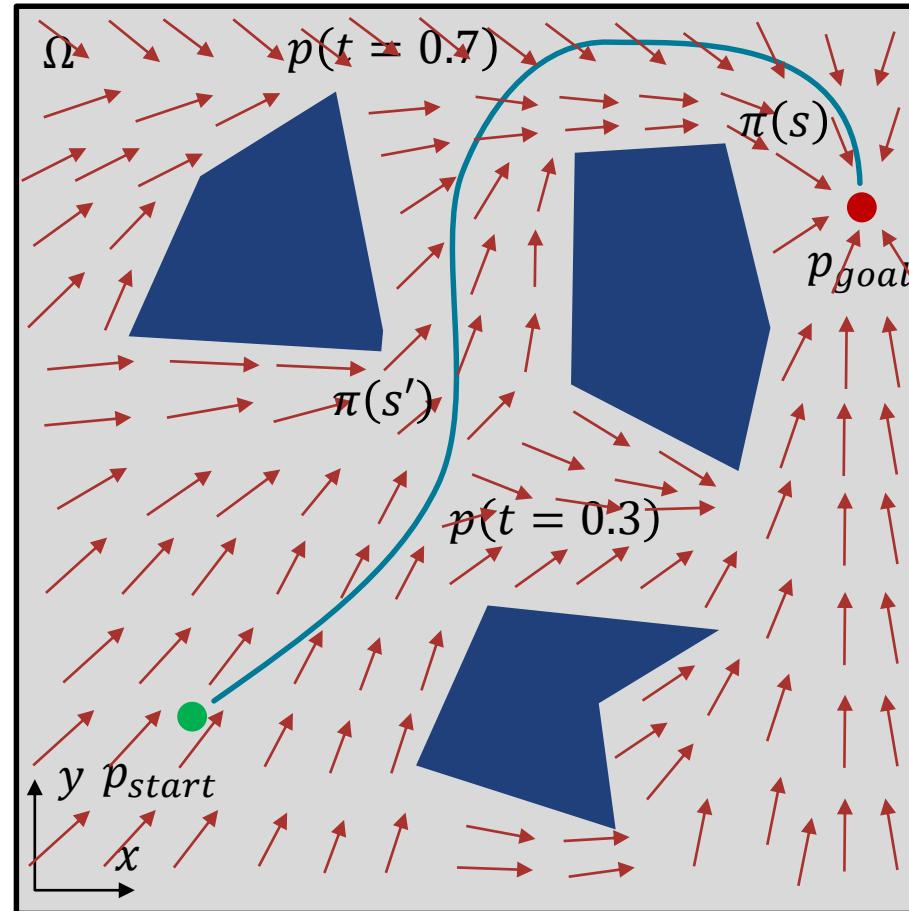
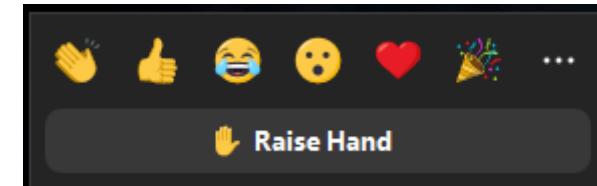
$$p: [0,1] \rightarrow \omega_{free}$$

$$\begin{aligned} p(t): \forall t \in [0,1], \\ p(t) \in \omega_{free} \end{aligned}$$

$$p(0) = p_{start}$$

$$p(1) = p_{goal}$$

# Paths, trajectories and policies

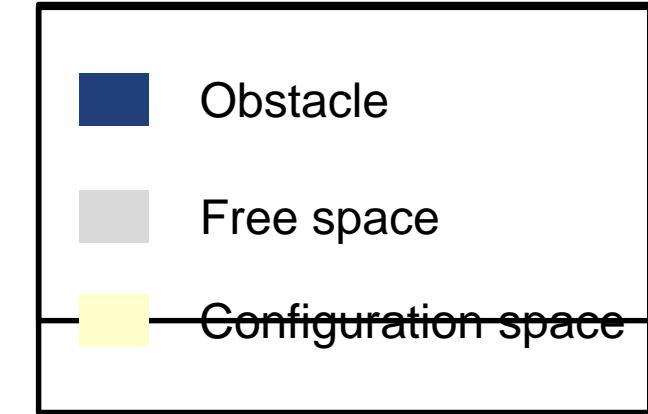
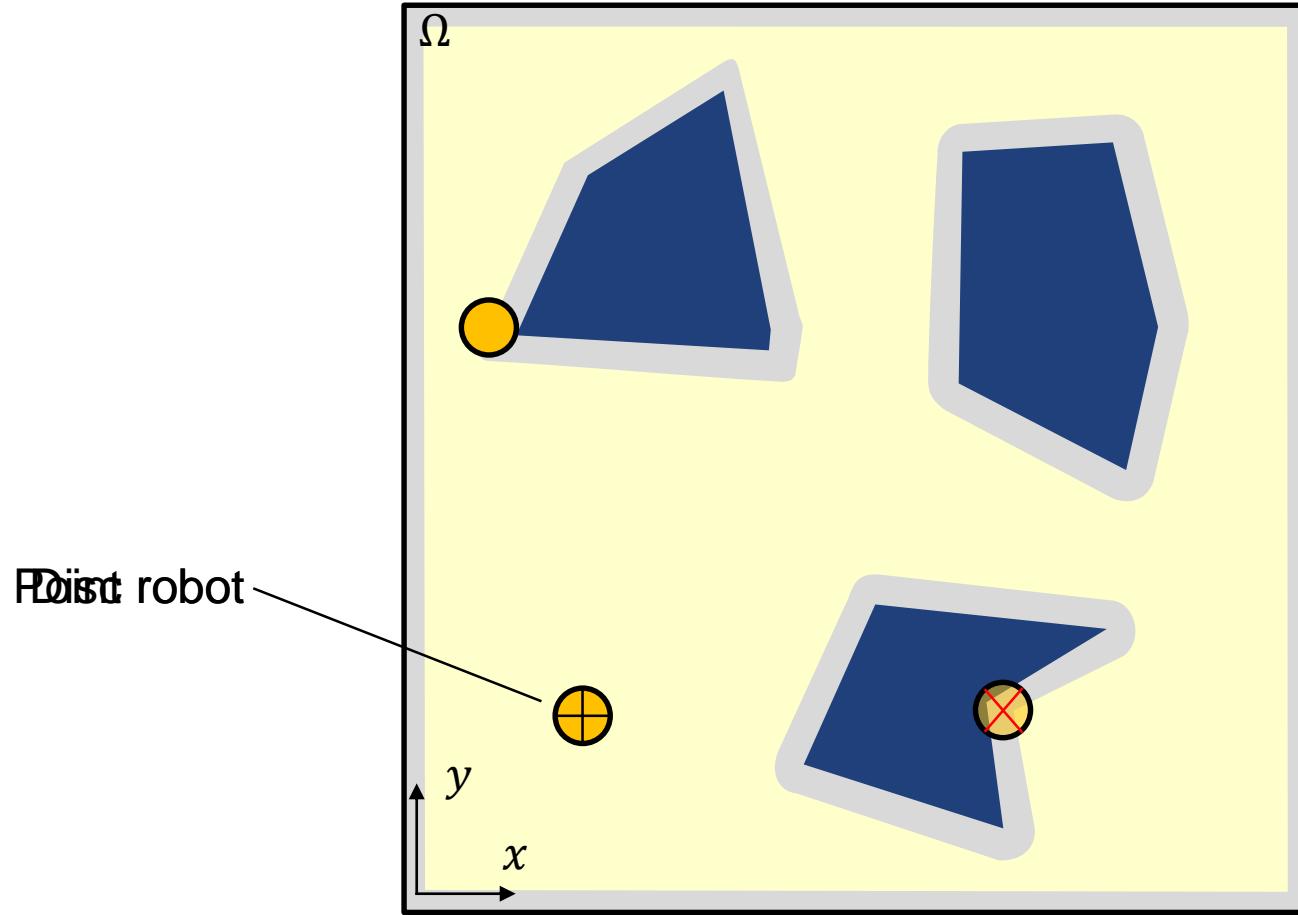
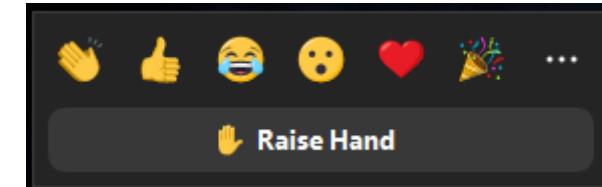


- **Path  $p \subset \Omega$ :** A set of points (geometric path) that the vehicle will travel along (no time consideration)
- **Trajectory  $p(t): T \rightarrow \Omega$ :** Path parameterized by **time** – important when considering dynamics constraints (e.g. turning radius as a function of speed)
- **Policy  $\pi(s): S \rightarrow A$ :** A function that maps (all valid) states to actions

# Representation – Workspace and configuration space

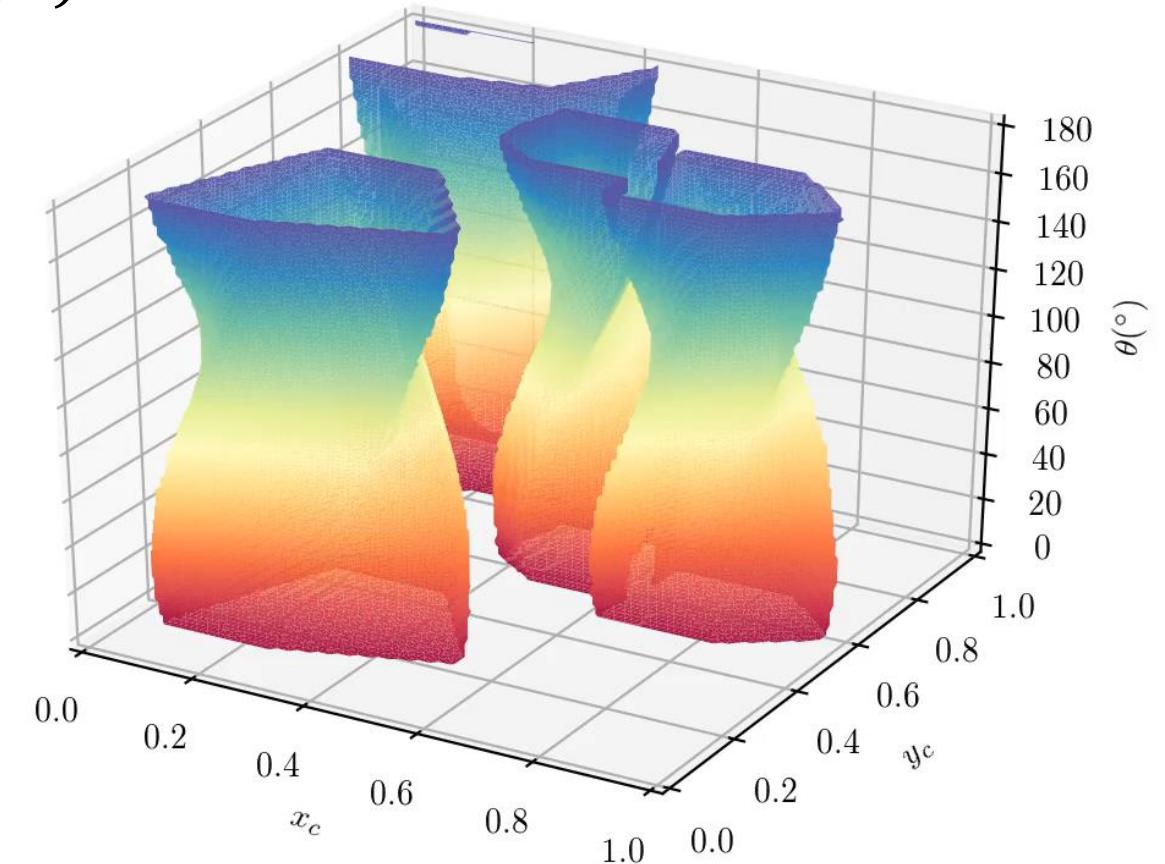
- **Workspace** is often the representation of the world, possibly independent of the robot itself. Often describes some notion of reachability, what space is free or occupied?
- **Configuration space** describes the full state of the robot in the world (actuator positions, orientation, etc.)

# Representation – configuration space

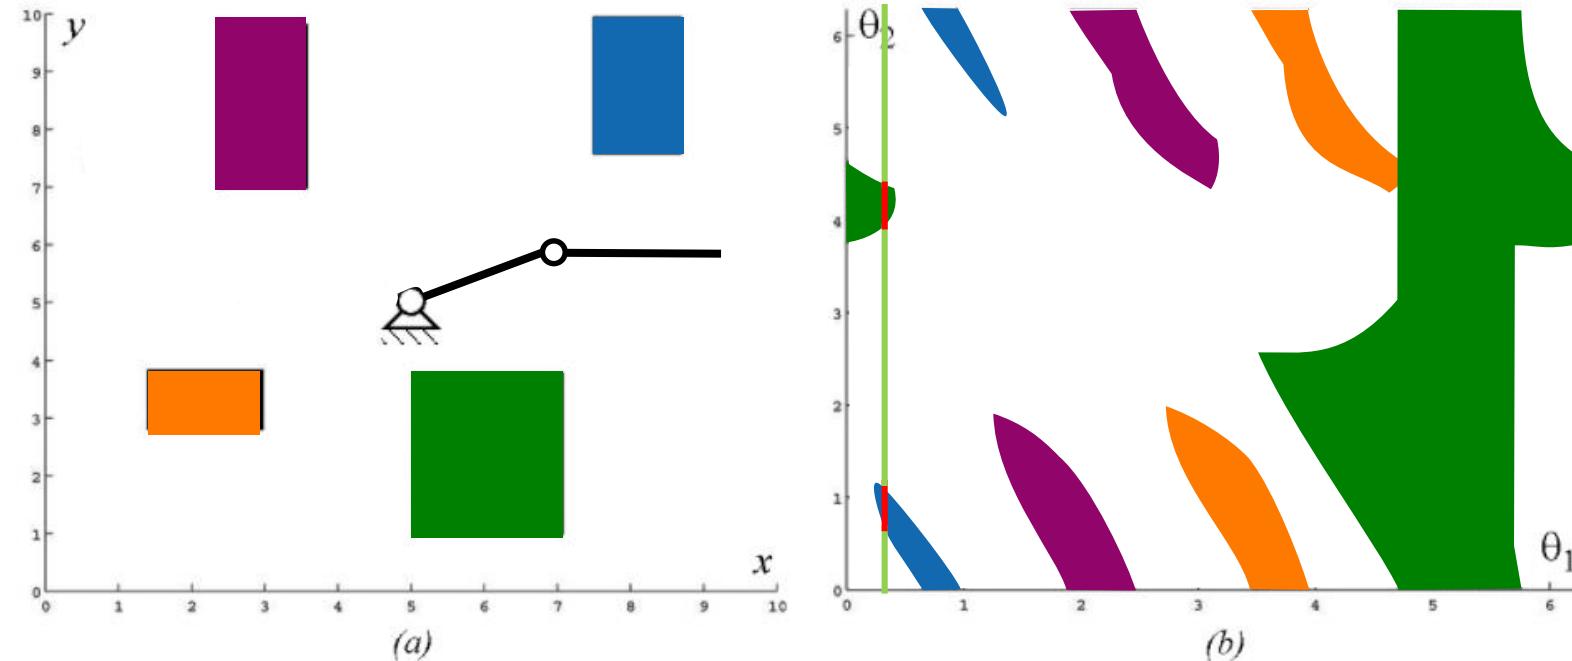


# Representation – configuration space

- A robot without rotational symmetry ( $x, y, \theta$ )



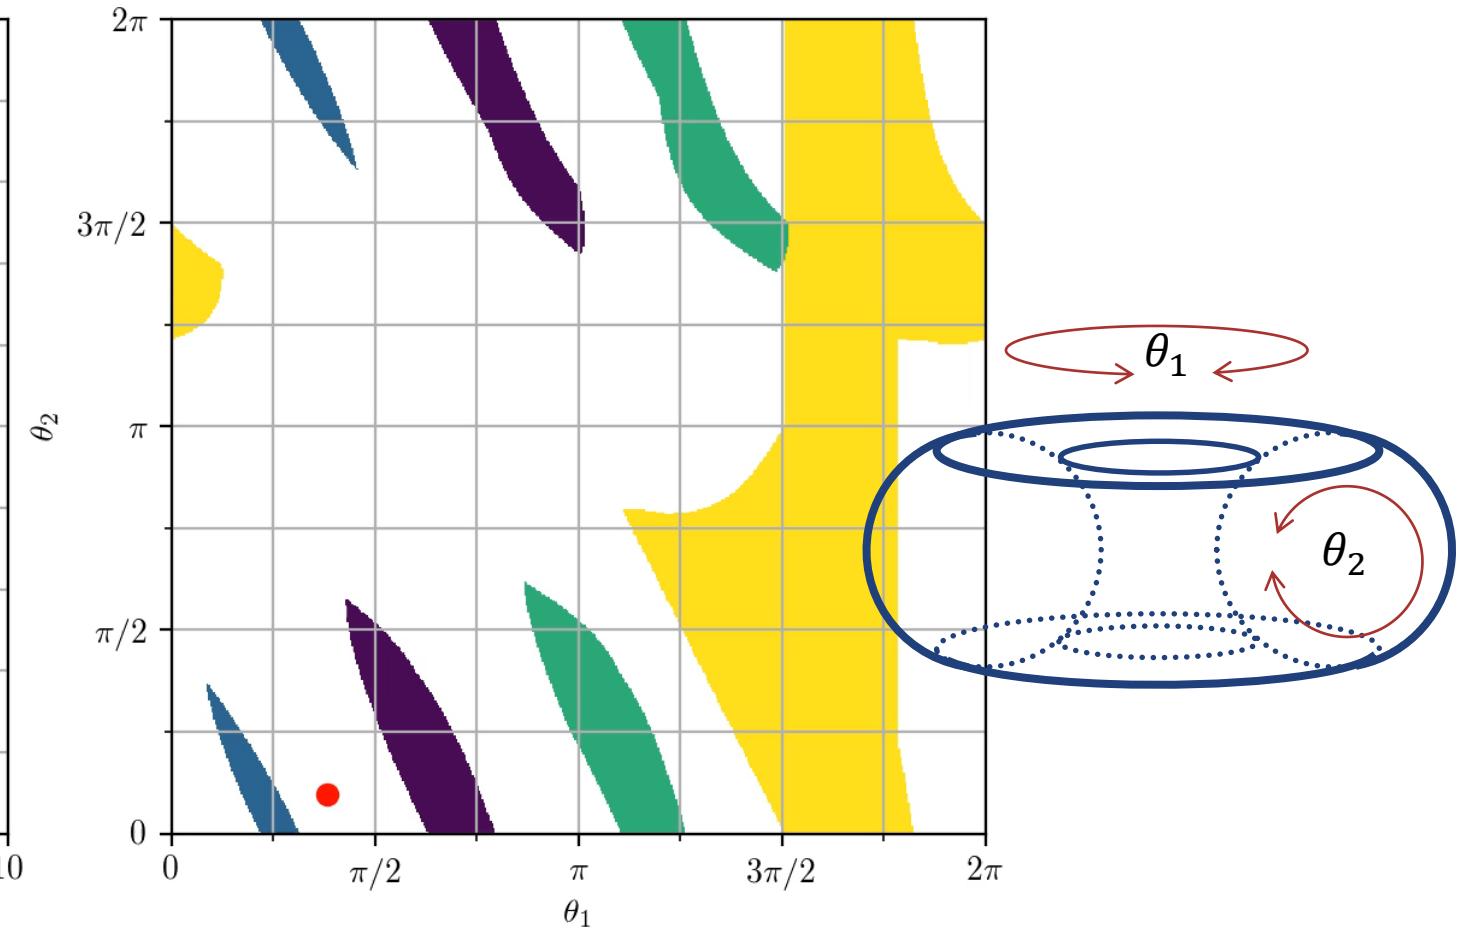
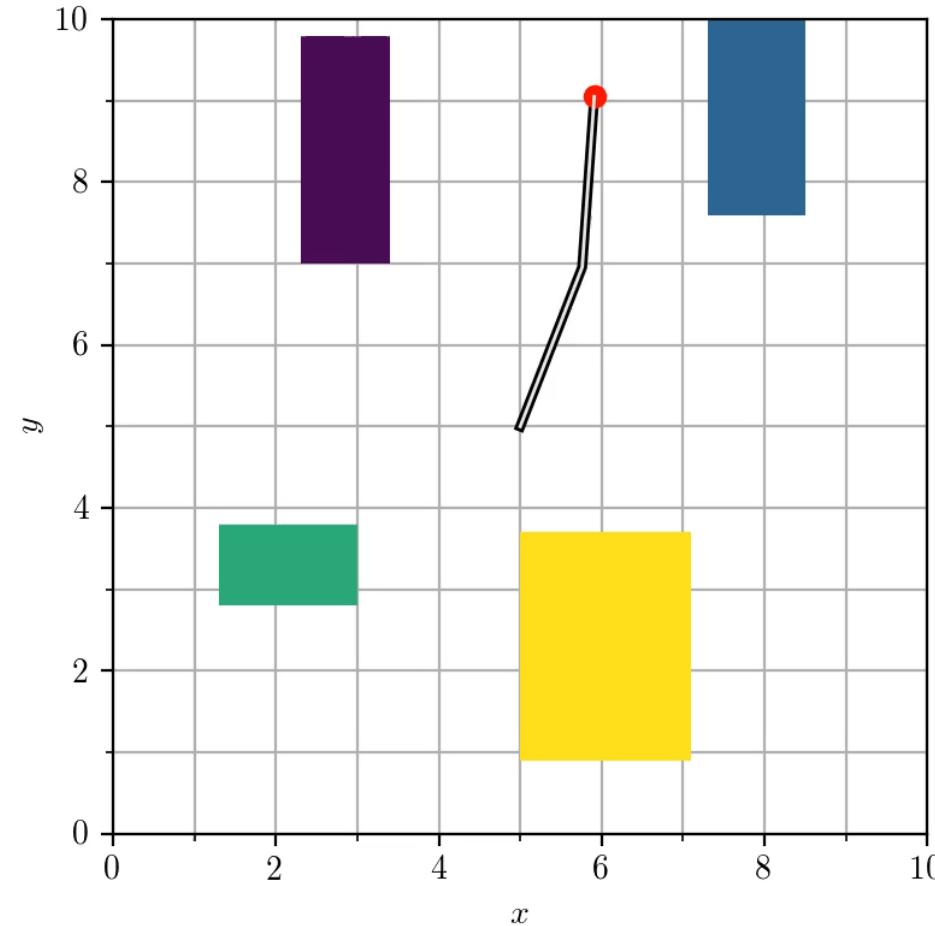
# Configuration space for alternative morphologies



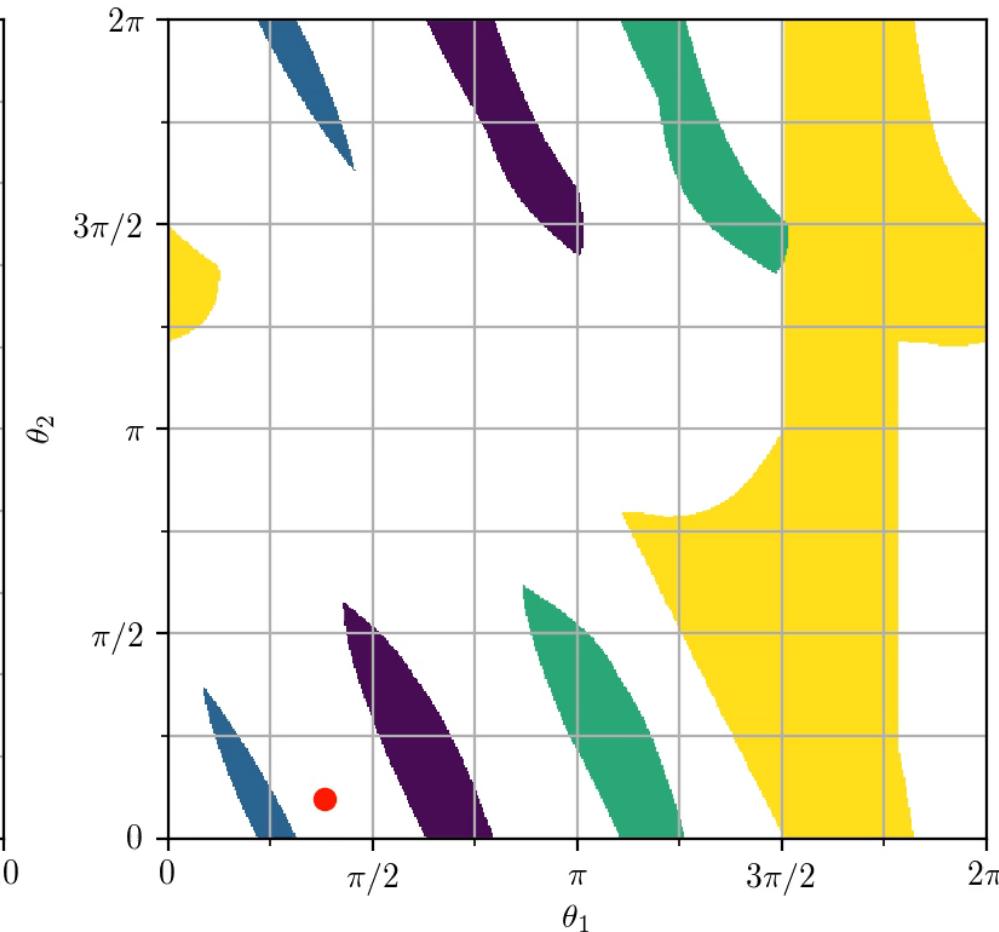
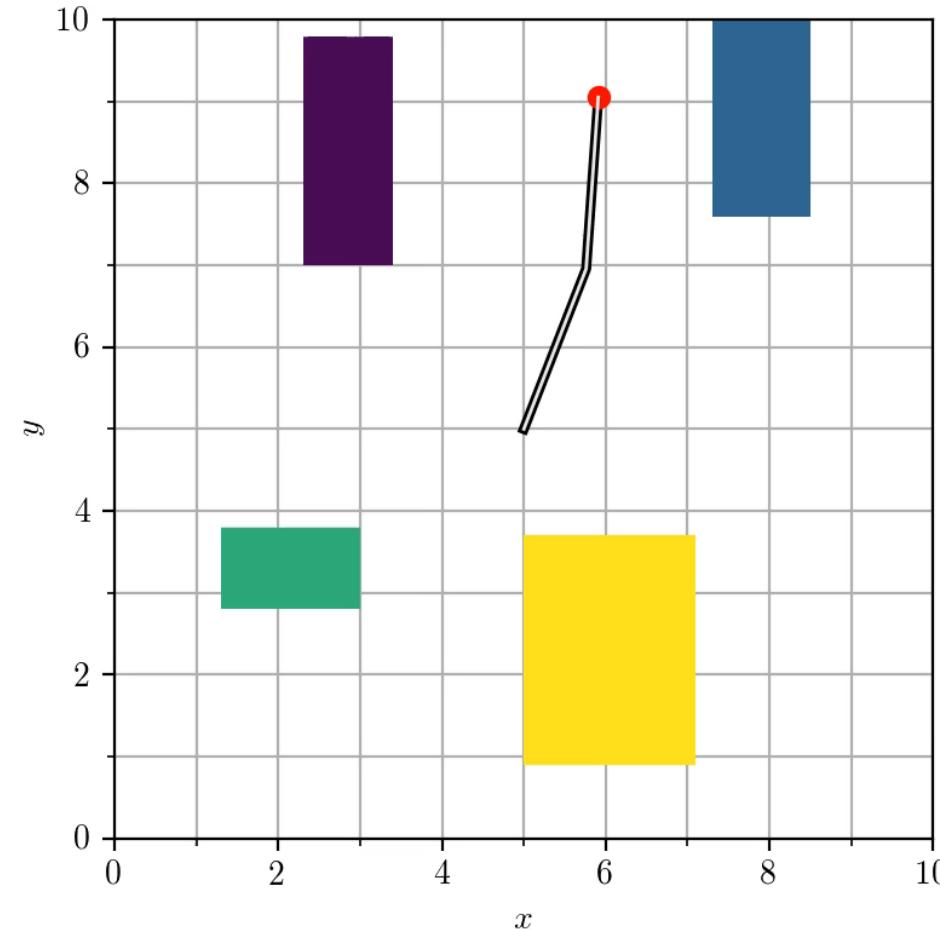
**Figure 6.1**

Physical space (a) and configuration space (b): (a) A two-link planar robot arm has to move from the configuration *start* to *end*. The motion is thereby constrained by the obstacles 1 to 4. (b) The corresponding configuration space shows the free space in joint coordinates (angle  $\theta_1$  and  $\theta_2$ ) and a path that achieves the goal.

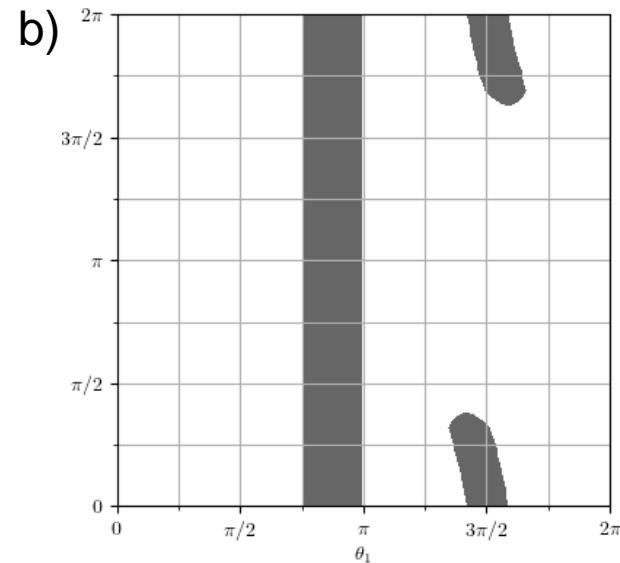
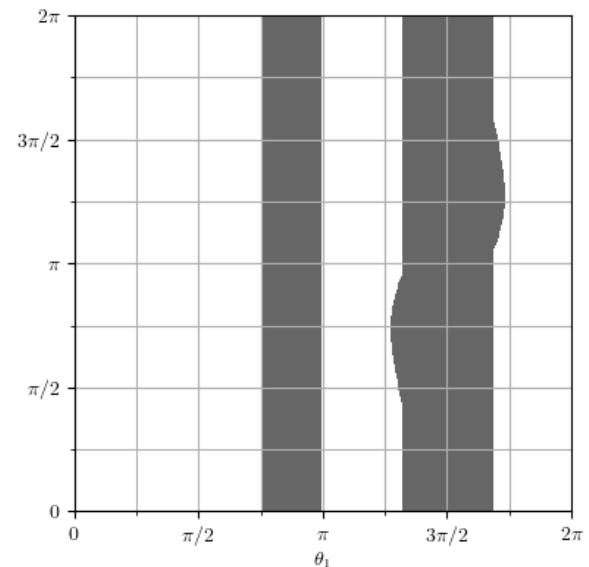
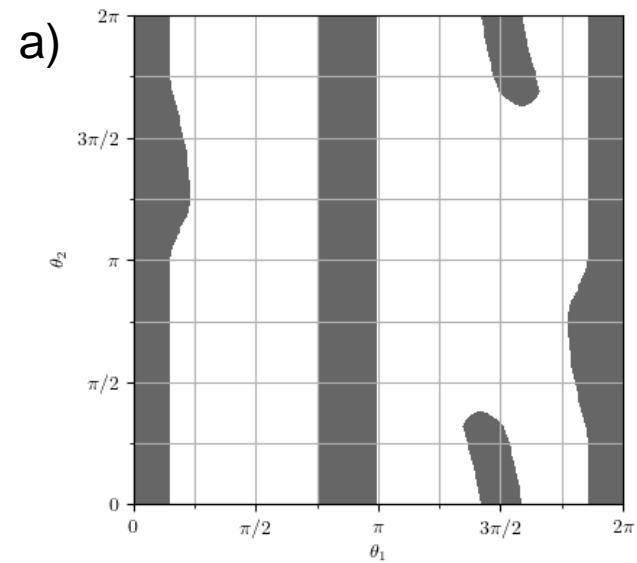
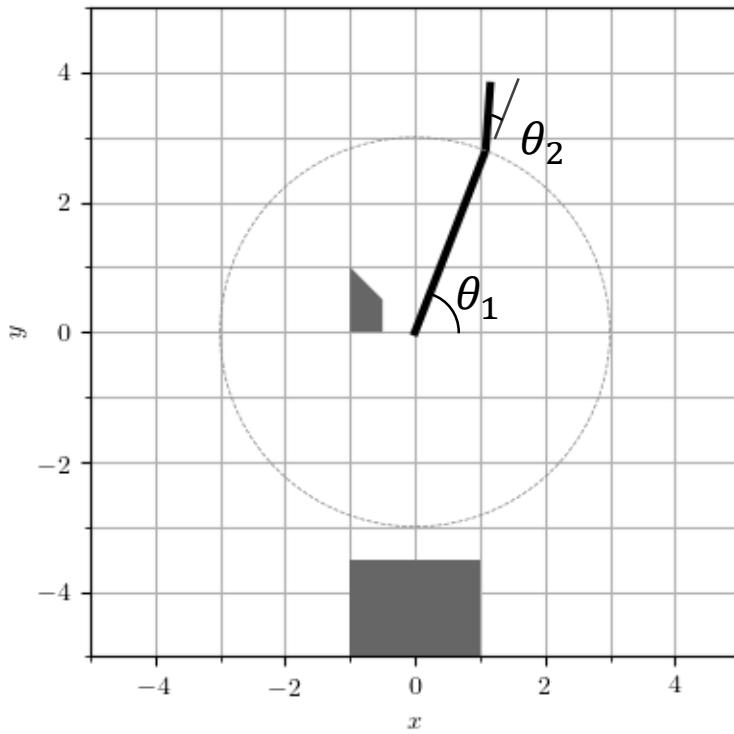
# Configuration space for alternative morphologies



# Configuration space for alternative morphologies

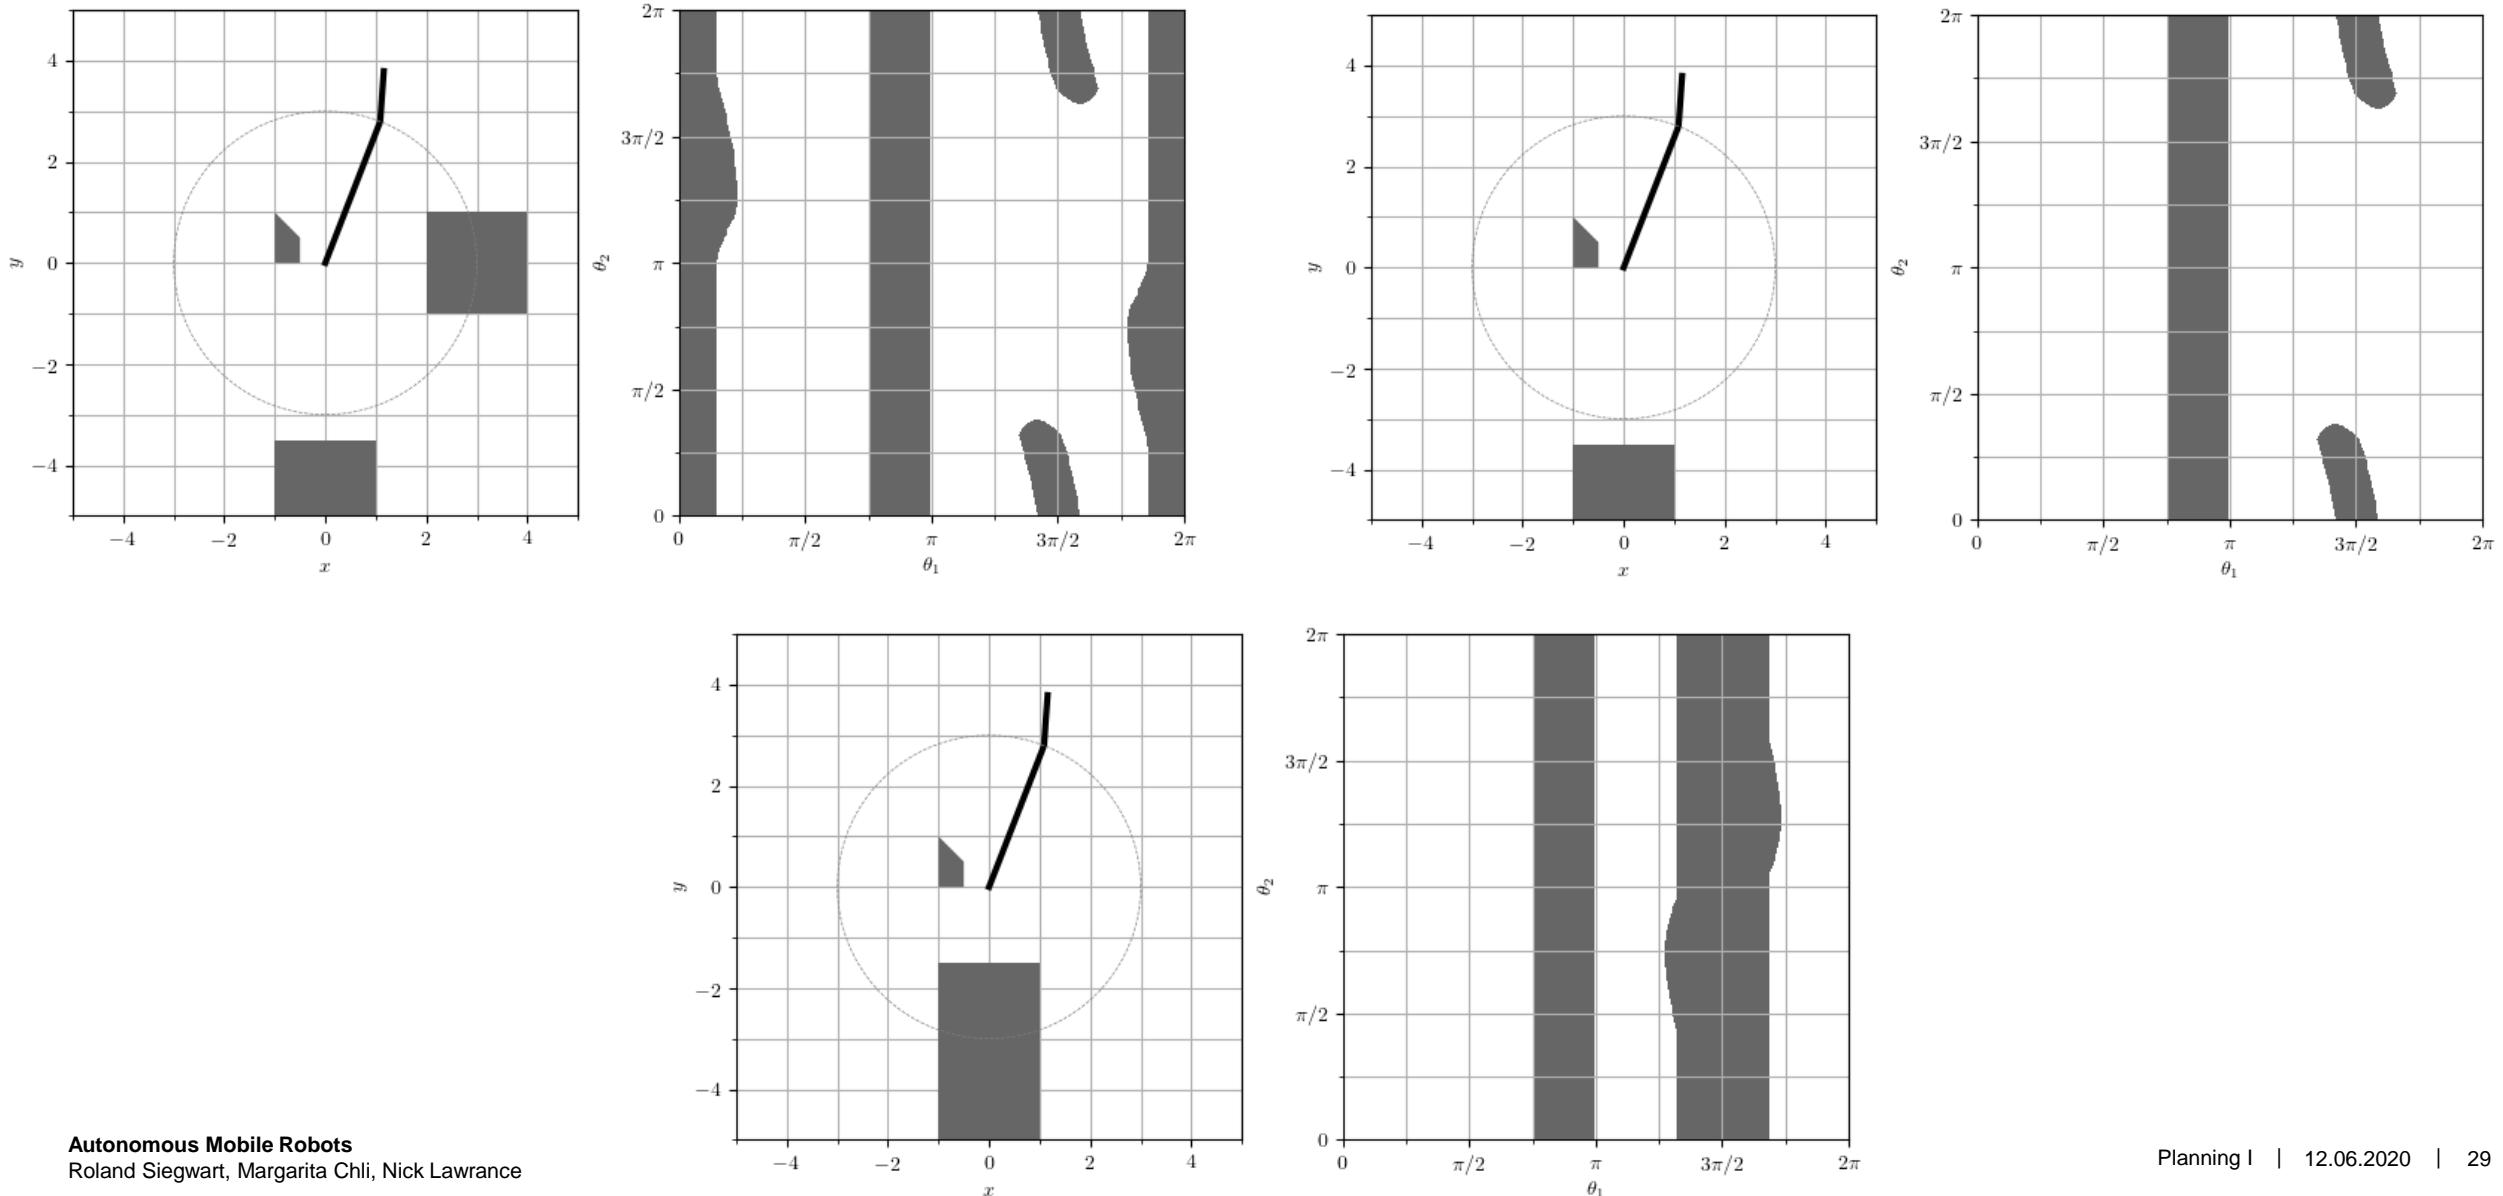


# Sample question – configuration space



<https://etc.ch/qbtw>





# Why use configuration space?

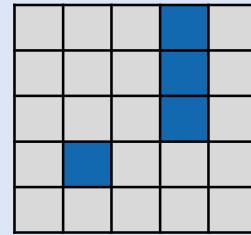
- Positions in configuration space tend be close together for the robot
- Can be easier to solve collision checks, and join nearby poses
- Allows a level of abstraction that means solution methods can solve a wider range of problems
- Sometimes helps with wraparound conditions (rotation joints)

# Continuous vs discrete state space representations

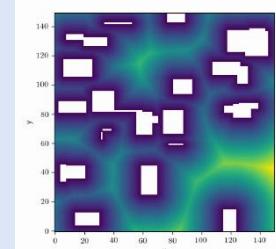
- Although continuous representations have some nice mathematical properties, they aren't always convenient
- Since computers store things digitally, there are some clever and efficient ways to create (approximate) discrete representations
- It is **very** common to convert a planning problem to some kind of (discrete) graph representation, then use one of a variety of existing search algorithms on the graph

# Structured

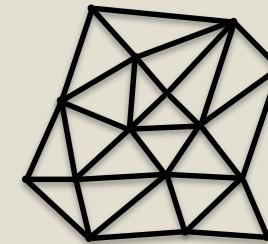
Occupancy grids



Distance fields

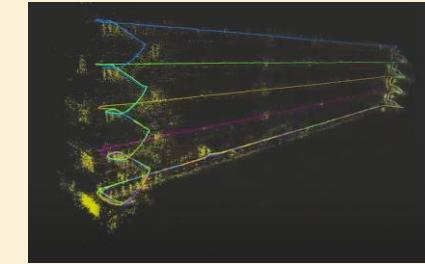


Graphs

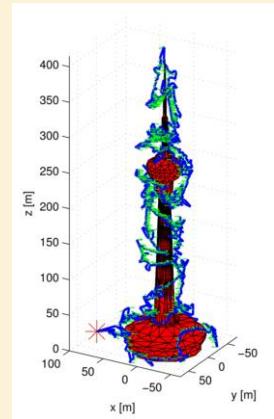


# Unstructured

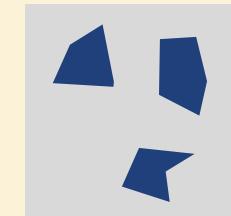
Point clouds



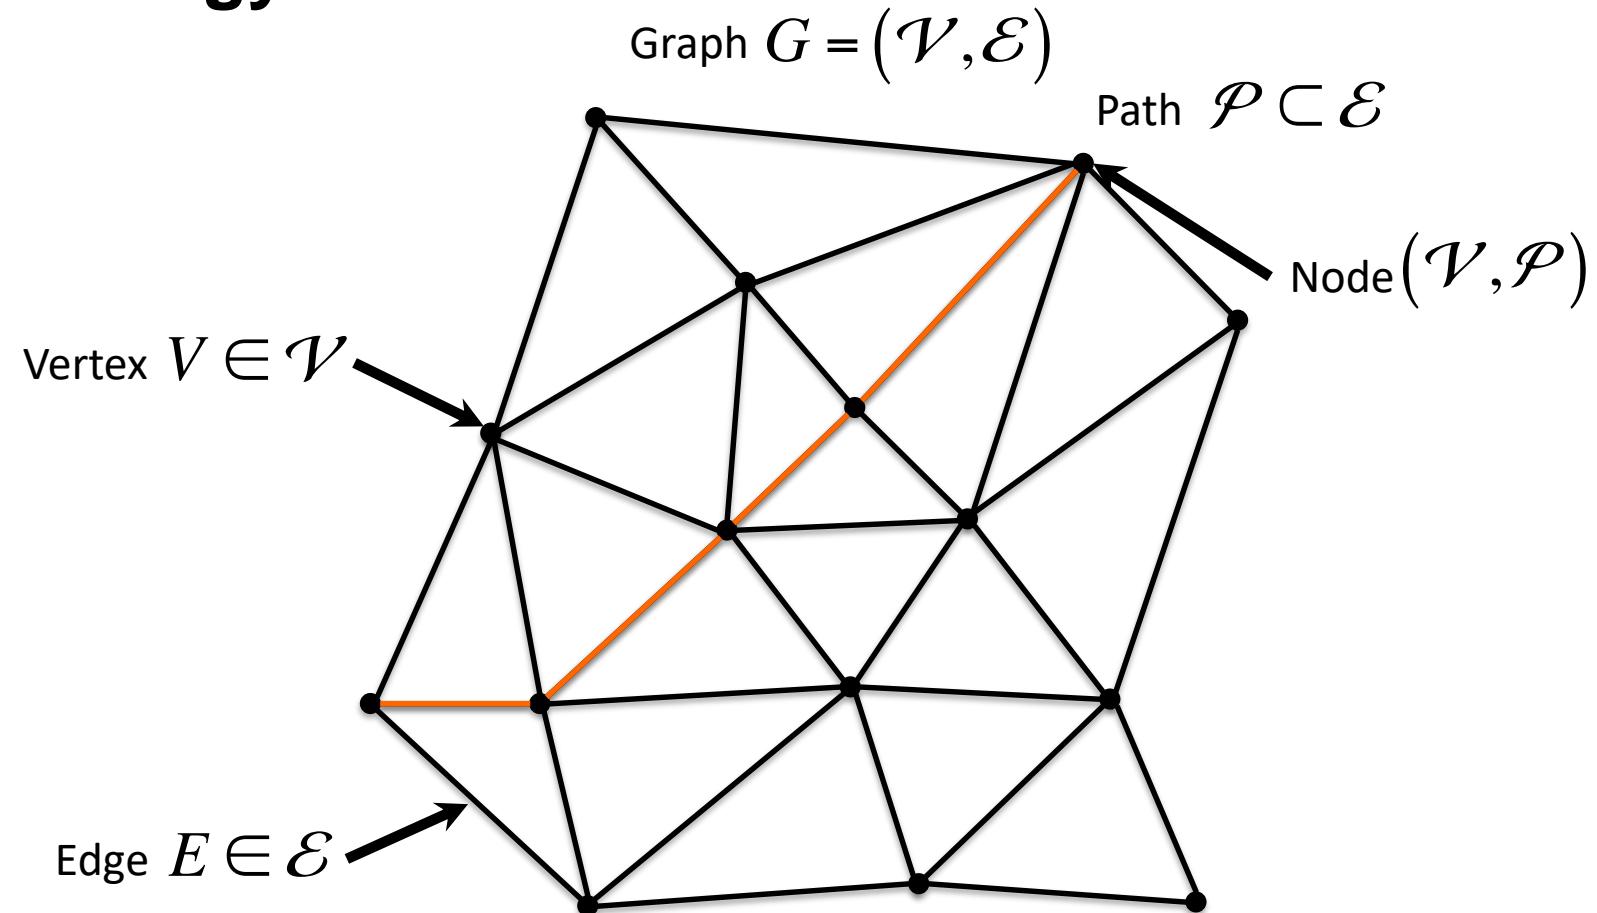
Meshes



Exact



# Some Terminology



**Directed** graph: edges have direction

**Weighted** graph: edges have costs

# Discrete state space representation

- Reduce continuous state space to a finite set of discrete **states**

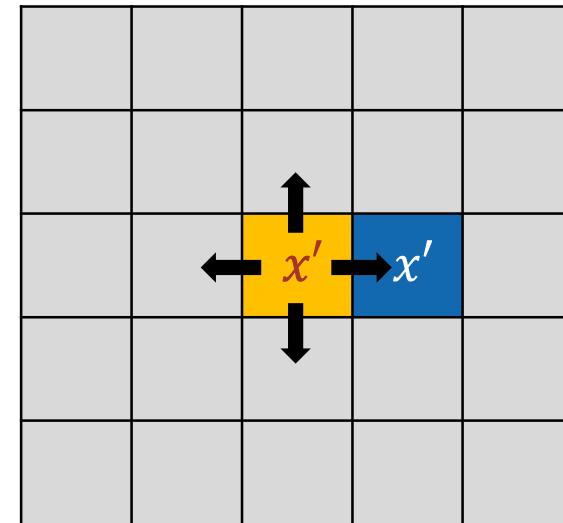
$$x \in X$$

- Also define feasible **actions** from each state

$$A(x) = \{a_0, a_1, \dots, a_n\}$$

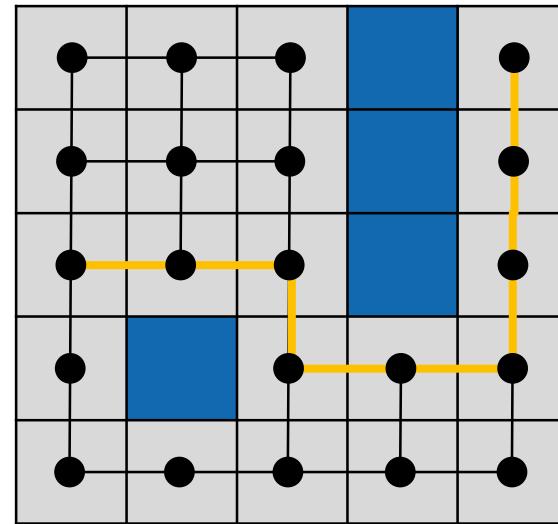
- And an associated **transition function**

$$f(x, a) = x'$$



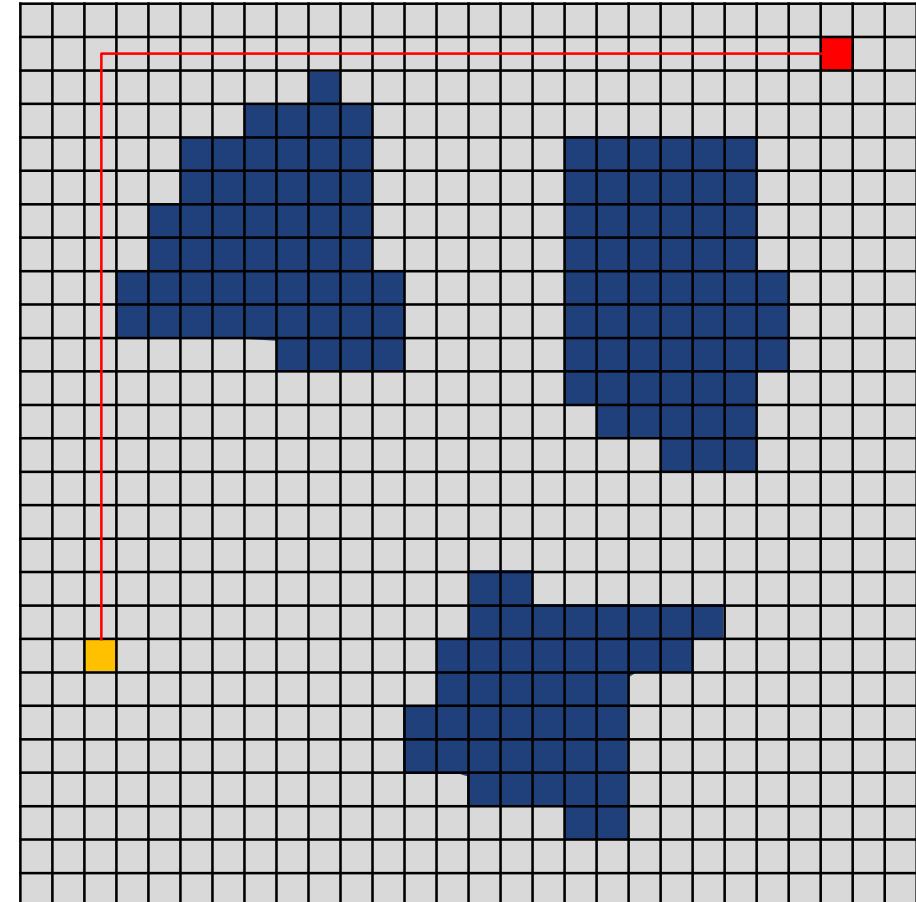
# Grid → Graph

- Consider:
  - States as vertices
  - Transitions as directed edges
- The result is a graph
- Add:
  - Start node,  $x_s$
  - Goal node,  $x_g$
  - Cost function  $C: X \times A \rightarrow \mathbb{R}^+$
- Finding the shortest path can be treated as a graph search problem



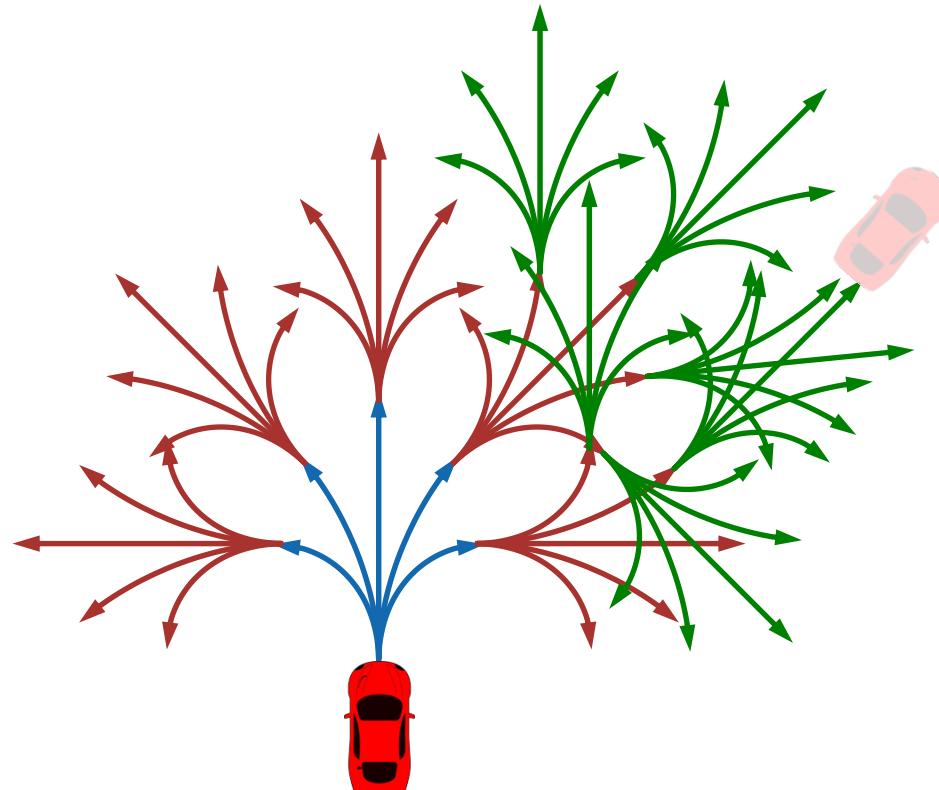
# Issues with grid-based representations

- Usually suffer some loss of precision
- Selecting an appropriate grid resolution can be a challenge (multi-resolution mapping)
- Can limit the type of output path
- Suffer from poor scaling in higher dimensions



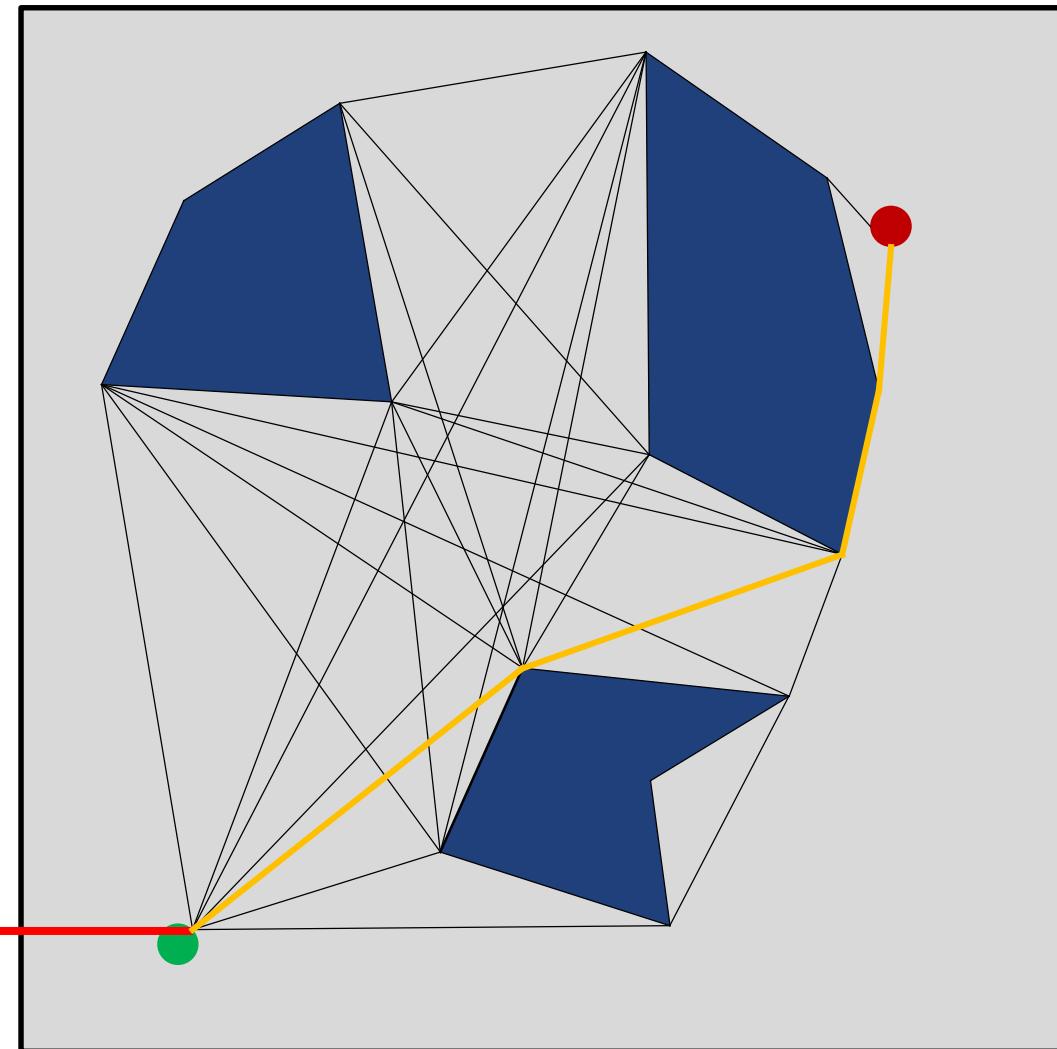
## Brief aside – other graph-based representations

- Grid lattice – create a set of feasible motion primitives, and construct a tree (graph) that chains the motions into a sequence (plan)



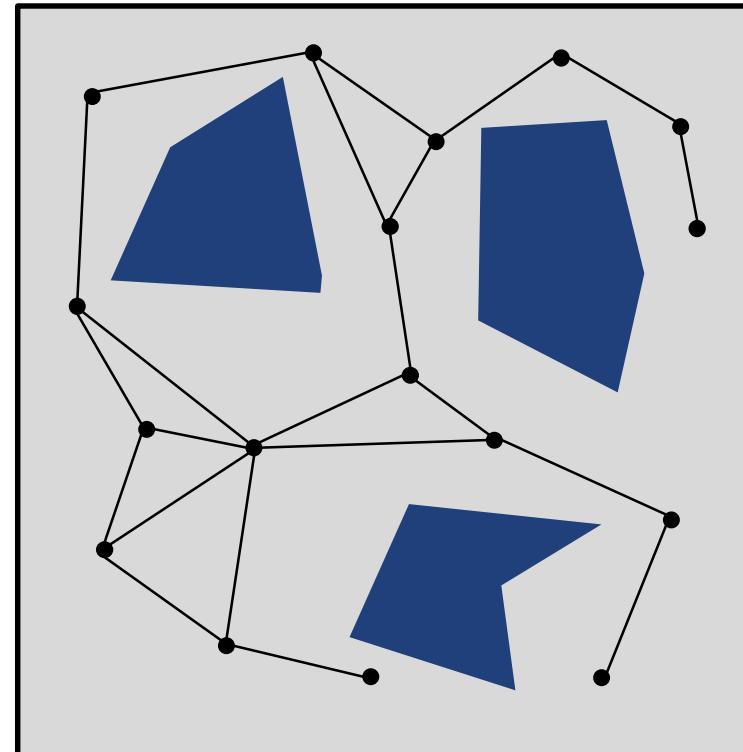
# Visibility graph

- Create edges between all pairs of mutually visible vertices
- Search resulting graph
- Optimal plan!
- Limited to straight motion, 2D, polygonal obstacles



# Randomly-sampled graphs

- Especially popular for sample-based methods (next lecture)
- Require careful consideration to construct graphs with guarantees



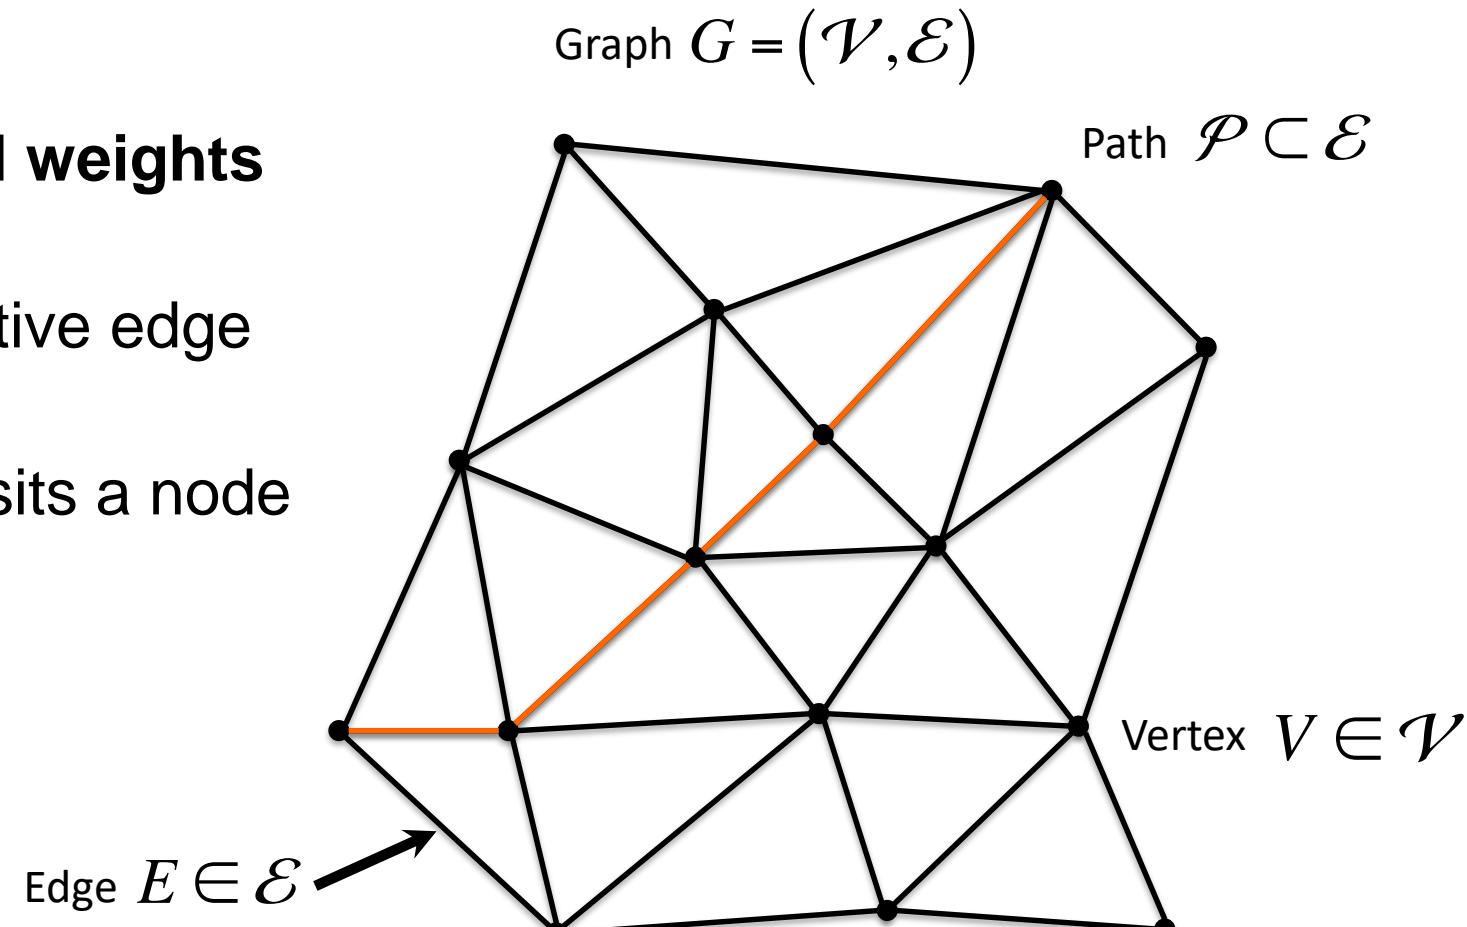
# Planning with graphs

- Given a representation, a start, a goal, and a motion model, how do we actually generate a plan?
- Many planning approaches use graphs, because we know how to search graphs and computers are good at it
- Solve your planning problem in three easy steps:
  1. Convert problem to a graph
  2. Search the graph
  3. Profit!

# GRAPH SEARCH METHODS

# Graph Search

- Edges can also have associated **weights** (cost of traversing the edge)
- We generally only consider positive edge weights
- A minimum cost path never revisits a node
  
- Graph search methods
  - Breadth-first search (BFS)
  - Depth-first search (DFS)
  - Dijkstra's Algorithm
  - A\*

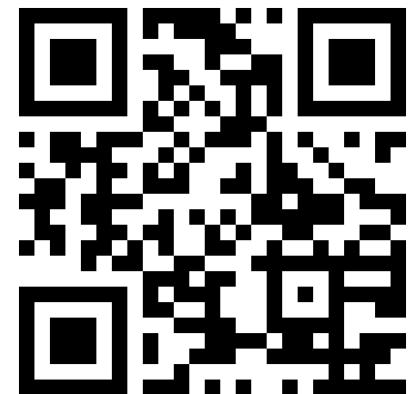


# Prior knowledge feedback

- Which graph search methods do you already know?

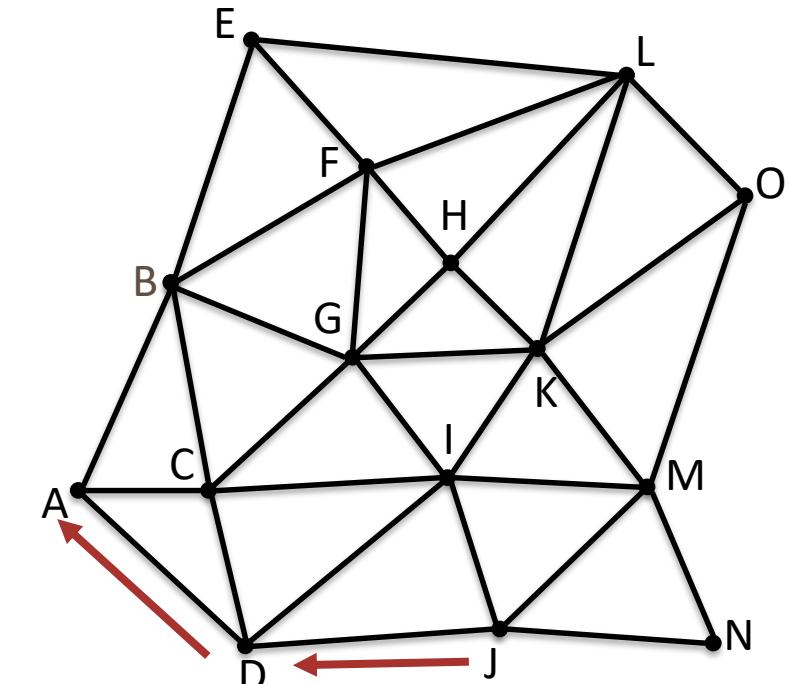
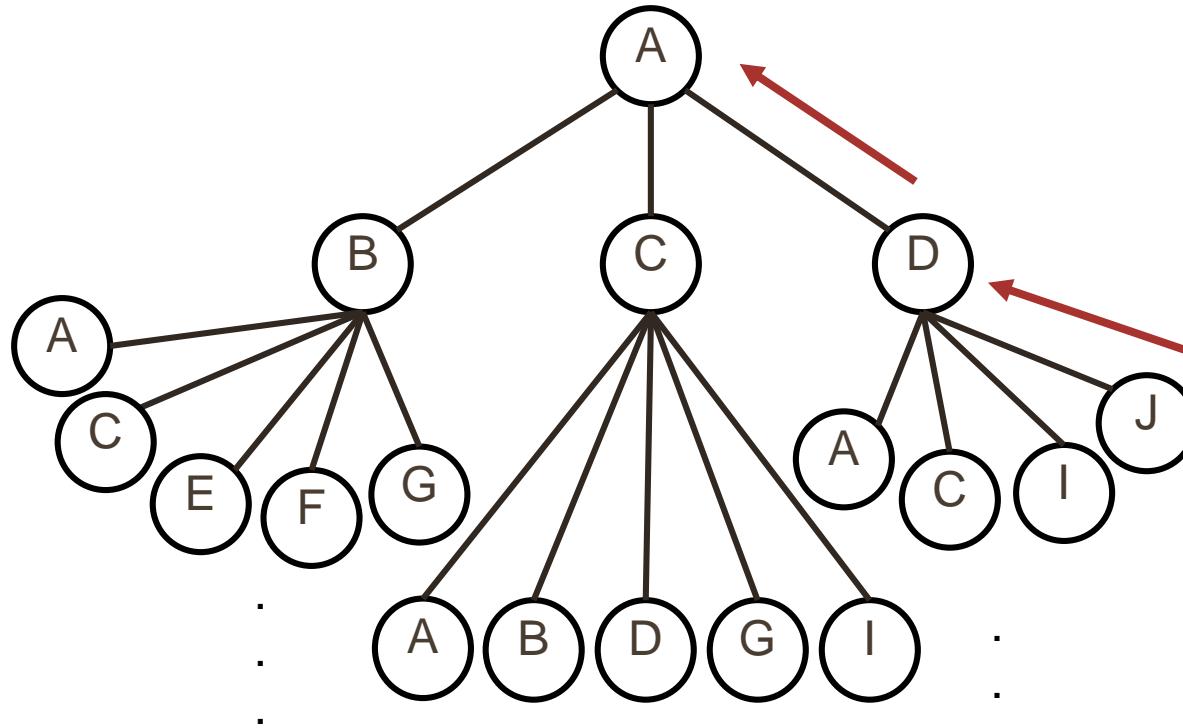
1. Breadth-first search
2. Depth-first search
3. Dijkstra's algorithm
4. A\*

<http://etc.ch/qbtw>



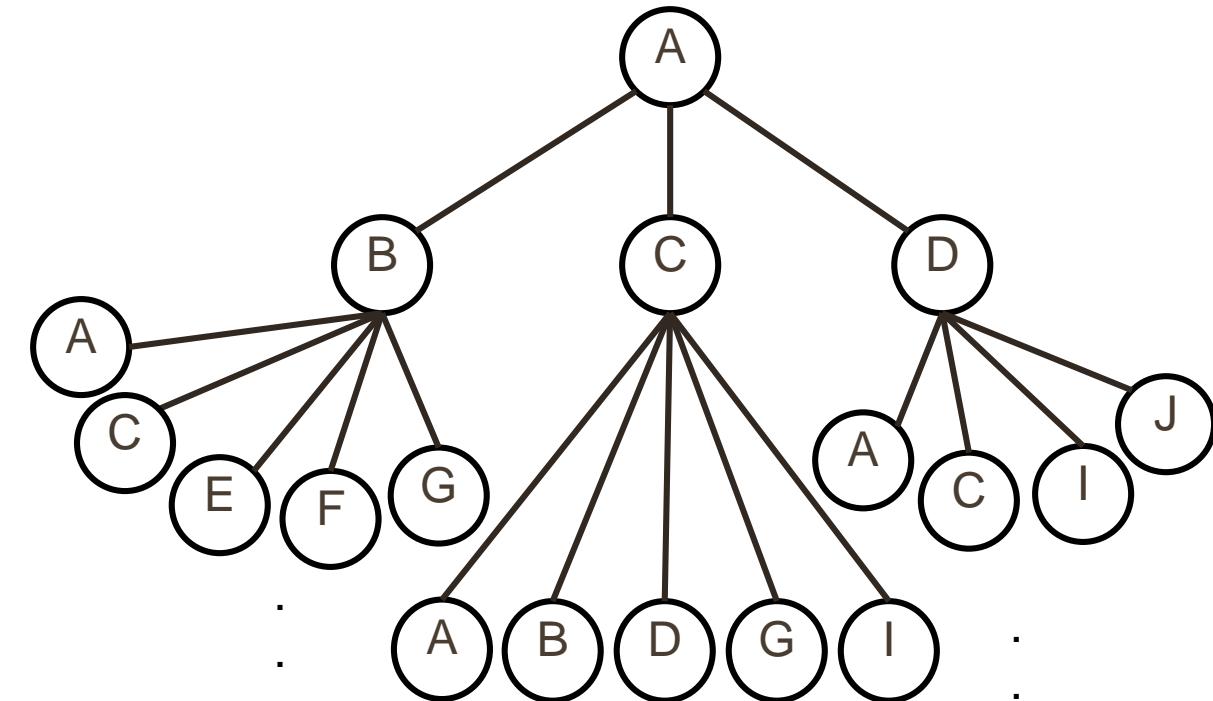
# Search Trees

- We construct a “tree” through which we can search for optimal paths through the environment



# Search Trees

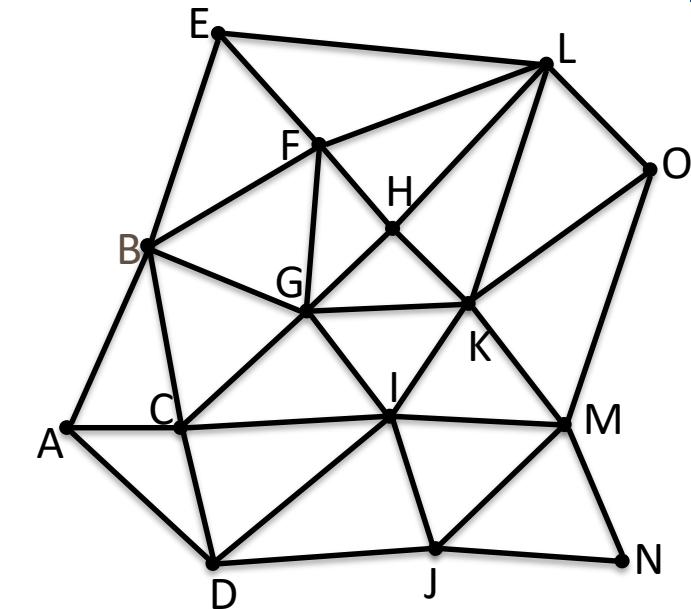
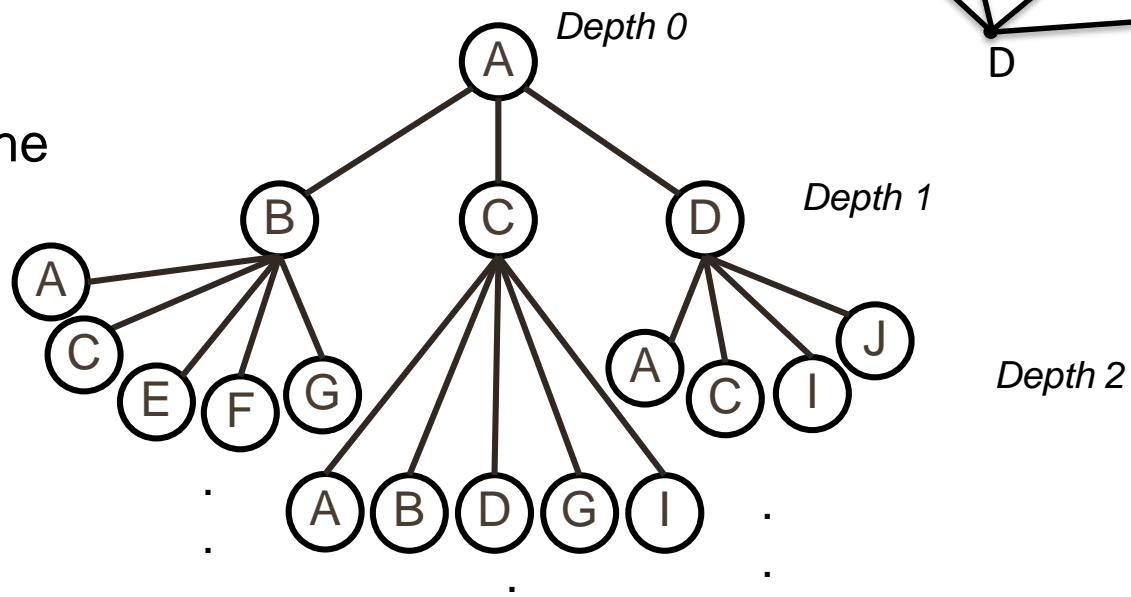
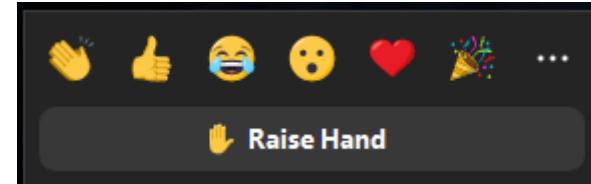
- A search tree:
  - Start state at the root node
  - Children correspond to successors
  - A node corresponds to a unique **plan** from start to that state (follow up tree from node)
  - For most problems, we want to avoid building the whole tree
  - Use search algorithms to efficiently traverse tree



# Search Trees

- Questions:

- Does depth in this tree correspond to distance (number of edges)?
- Would an optimal path ever visit the same node twice?
- Can there be multiple optimal solutions?



# Basics of Forward Search

- Generally, start from the start, grow tree until you find a solution (path to goal)
- Expanding a node refers to adding children to the tree, pushing them onto the open set
- Try to expand as few tree nodes as possible
- **Open** set maintains a list of frontier (unexpanded) plans
  - Keeps track of what nodes to expand next
  - Often stored as a priority queue
  - For each node in the open list, we know of at least one path to it from the start
- **Closed** set keeps track of nodes that have been expanded
  - For each node in the closed list, we've already found the lowest-cost path to it from the start

# Forward Search algorithm from LaValle

---

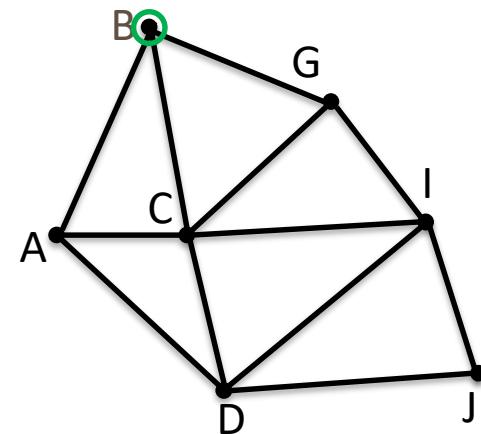
```
FORWARD_SEARCH
1    $Q.Insert(x_I)$  and mark  $x_I$  as visited
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x'$  not visited
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE
```

---

Figure 2.4: A general template for forward search.

LaValle, Steven M. *Planning algorithms*. Cambridge university press, 2006, p. 33

# Breadth-First Search Example




---

## FORWARD SEARCH

```

1   Q.Insert( $x_I$ ) and mark  $x_I$  as visited
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x'$  not visited
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE

```

---

Open ( $Q$ ):

{**B**}

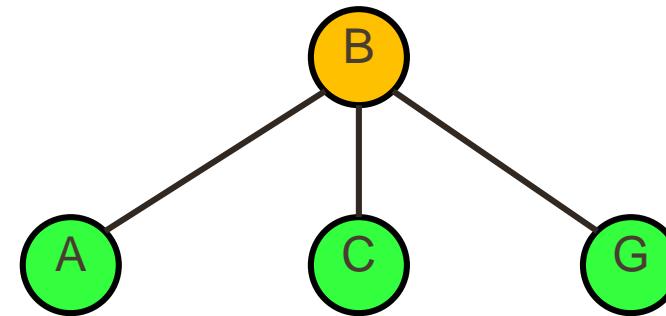
Visited:

{}

Our (BFS) queue will be FIFO:  

- push ( $Q.Insert$ ) onto the end
- pop ( $Q.GetFirst$ ) from the front

# Breadth-First Search Example

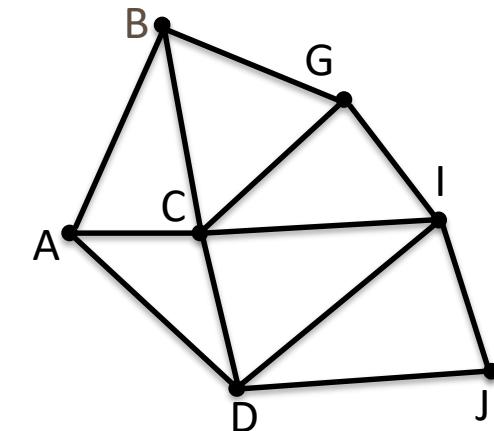



---

FORWARD\_SEARCH

- 1     $Q.Insert(x_I)$  and mark  $x_I$  as visited
- 2    while  $Q$  not empty do
- 3        $x \leftarrow Q.GetFirst()$
- 4       if  $x \in X_G$
- 5           return SUCCESS
- 6       forall  $u \in U(x)$
- 7            $x' \leftarrow f(x, u)$
- 8           if  $x'$  not visited
- 9              Mark  $x'$  as visited
- 10              $Q.Insert(x')$
- 11       else
- 12              Resolve duplicate  $x'$
- 13   return FAILURE

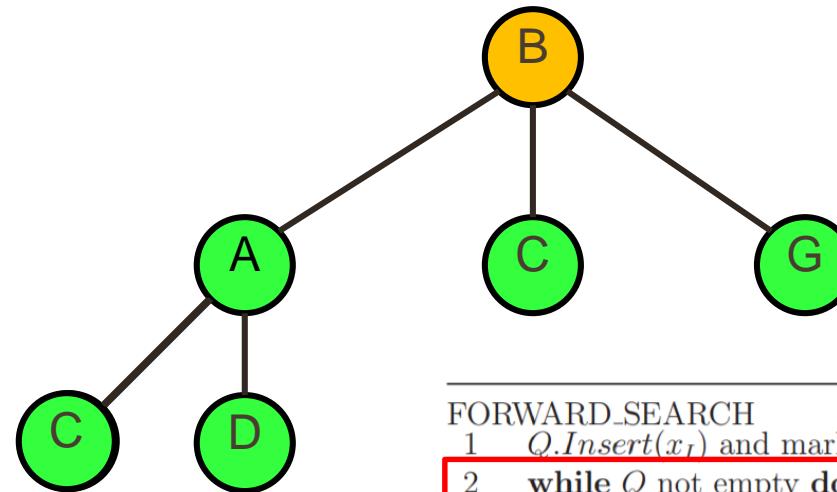
---



Open ( $Q$ ):  
 $\{A, C, G\}$

Visited:  
 $\{B, A, C, G\}$

# Breadth-First Search Example



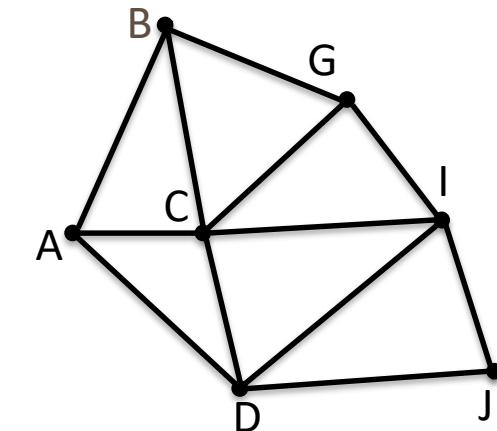

---

```

FORWARD_SEARCH
1   Q.Insert( $x_I$ ) and mark  $x_I$  as visited
2   while Q not empty do
3        $x \leftarrow Q.GetFirst()$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x'$  not visited
9               Mark  $x'$  as visited
10              Q.Insert( $x'$ )
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE

```

---



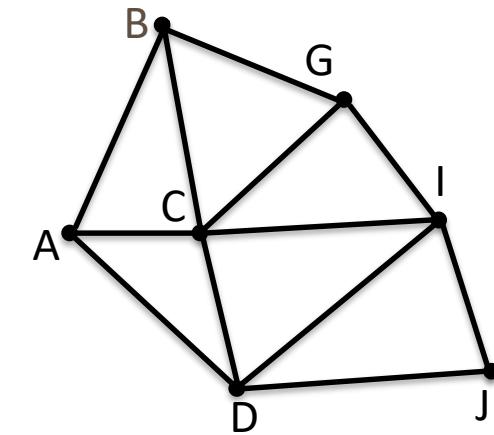
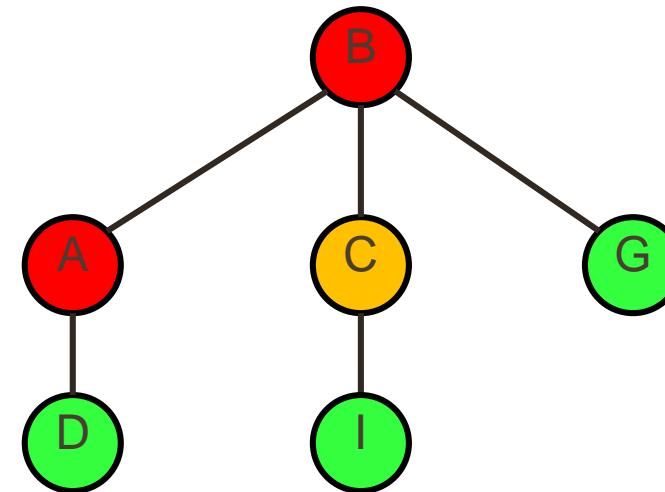
Open (Q):

{C,G,D}

Visited:

{B,A,C,G,D}

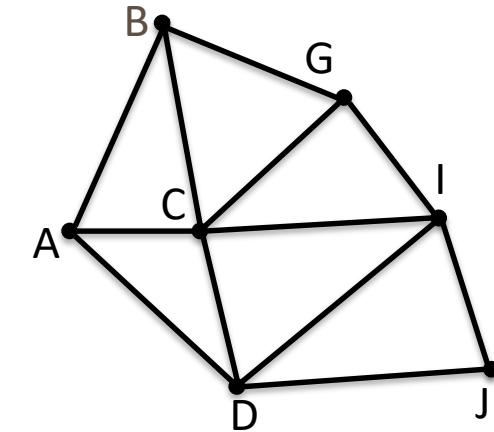
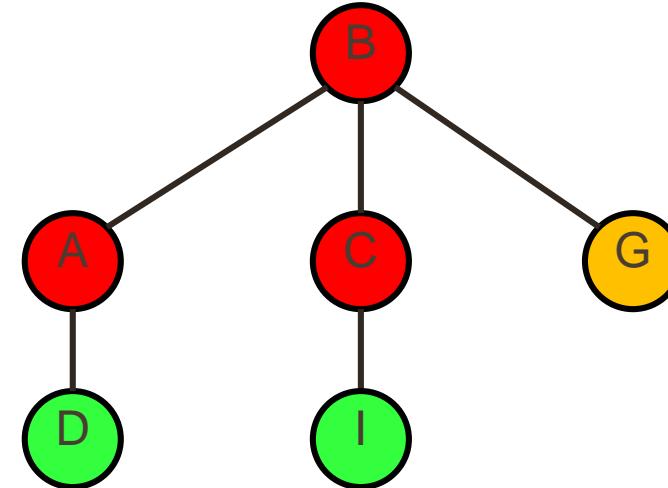
# Breadth-First Search Example



Open (Q):  
 $\{G, D, I\}$

Visited:  
 $\{B, A, C, G, D, I\}$

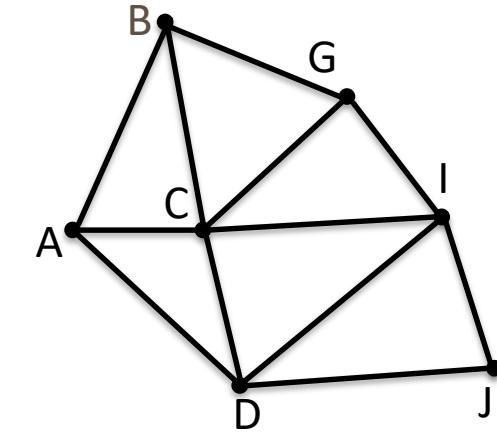
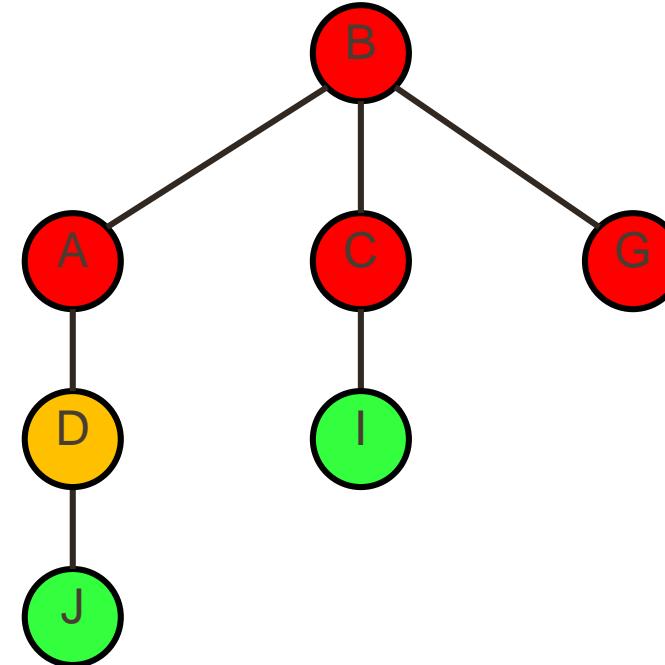
# Breadth-First Search Example



Open (Q):  
{D,I}

Visited:  
{B,A,C,G,D,I}

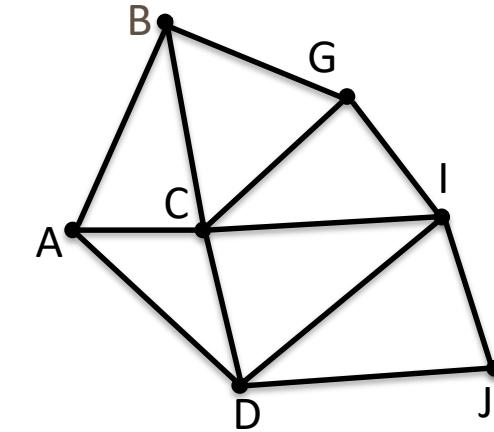
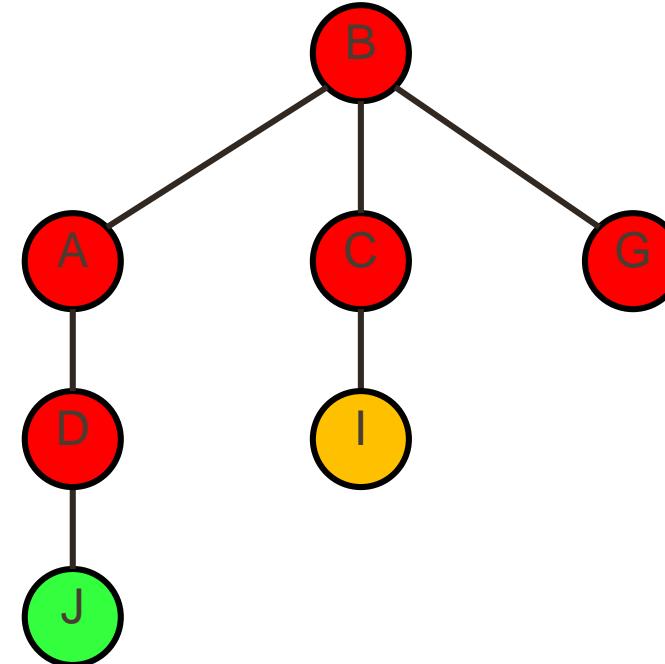
# Breadth-First Search Example



Open (Q):  
 $\{I, J\}$

Visited:  
 $\{B, A, C, G, D, I, J\}$

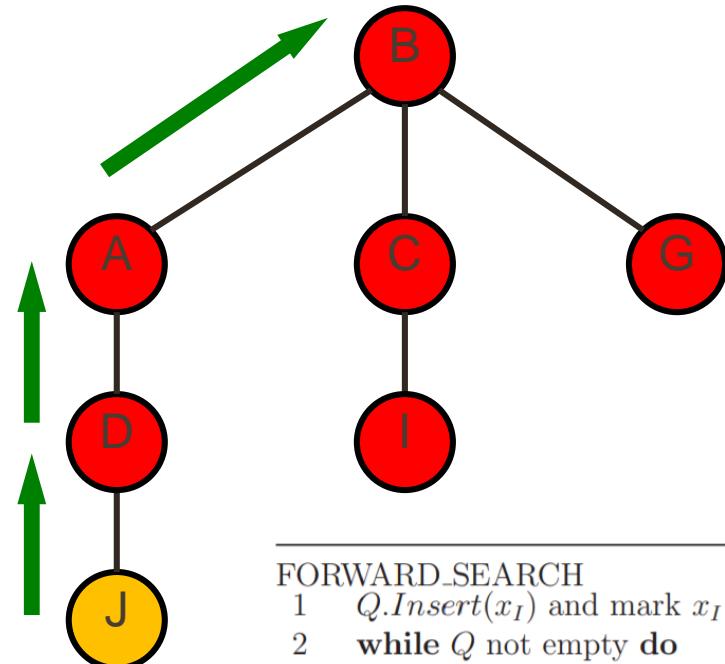
# Breadth-First Search Example



Open (Q):  
{J}

Visited:  
{B,A,C,G,D,I,J}

# Breadth-First Search Example

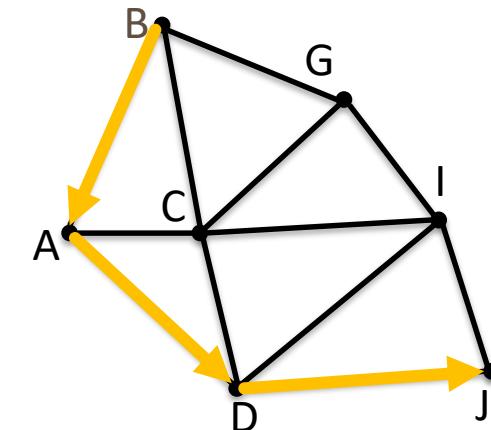



---

FORWARD\_SEARCH

- 1     $Q.Insert(x_I)$  and mark  $x_I$  as visited
- 2    **while**  $Q$  not empty **do**
- 3         $x \leftarrow Q.GetFirst()$
- 4        **if**  $x \in X_G$
- 5            **return** SUCCESS
- 6        **forall**  $u \in U(x)$
- 7             $x' \leftarrow f(x, u)$
- 8            **if**  $x'$  not visited
- 9                Mark  $x'$  as visited
- 10               $Q.Insert(x')$
- 11        **else**
- 12                Resolve duplicate  $x'$
- 13    **return** FAILURE

---



Open ( $Q$ ):     {}                  Visited:  
 $\{B, A, C, G, D, I, J\}$

Final path solution:  $B \rightarrow A \rightarrow D \rightarrow J$

Other solutions may exist but have the same number or more transitions

# Breadth-First Search

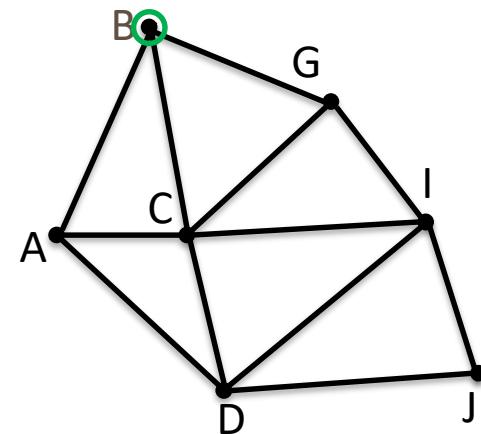
- Complete (will find the solution if it exists)
- Guaranteed to find the shortest (number of edges) path
  - First solution found is the optimal path
- What about non-uniform edge weights? (... Dijkstra)
- Time complexity  $O(|V|+|E|)$
- Consider another approach: Depth-first search

# Depth-first search

- Instead of searching across levels of the tree, DFS starts at the root node and explores as far as possible along each branch before backtracking
- Similar implementation to BFS, but with a stack (last-in first-out) queue

# Depth-First Search Example

B




---

## FORWARD SEARCH

```

1    $Q.Insert(x_I)$  and mark  $x_I$  as visited
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x'$  not visited
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11           else
12               Resolve duplicate  $x'$ 
13   return FAILURE

```

---

Open (Q):

{B}

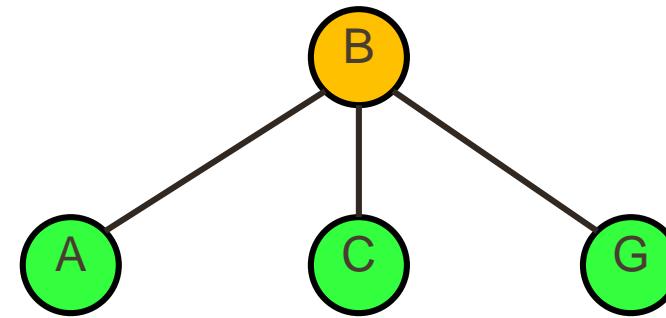
Visited:

{}

Our (DFS) queue will be LIFO:  

- push ( $Q.Insert$ ) onto the front
- pop ( $Q.GetFirst$ ) from the front

# Depth-First Search Example

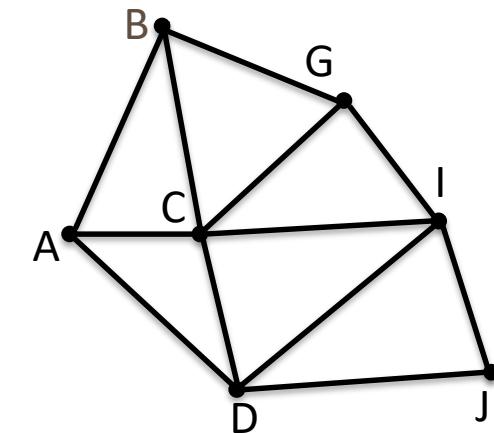



---

FORWARD\_SEARCH

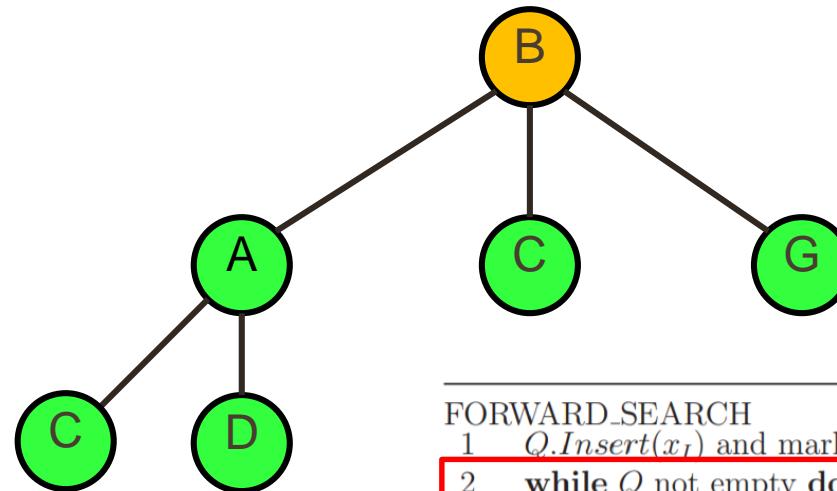
- 1     $Q.Insert(x_I)$  and mark  $x_I$  as visited
- 2    while  $Q$  not empty do
- 3        $x \leftarrow Q.GetFirst()$
- 4       if  $x \in X_G$
- 5           return SUCCESS
- 6       forall  $u \in U(x)$
- 7            $x' \leftarrow f(x, u)$
- 8           if  $x'$  not visited
- 9              Mark  $x'$  as visited
- 10           $Q.Insert(x')$
- 11       else
- 12              Resolve duplicate  $x'$
- 13   return FAILURE

---



Open ( $Q$ ):      Visited:  
 $\{A, C, G\}$        $\{B, A, C, G\}$

# Depth-First Search Example

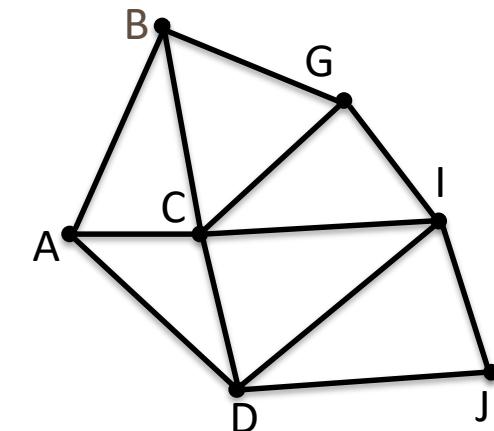



---

**FORWARD\_SEARCH**

- 1     $Q.Insert(x_I)$  and mark  $x_I$  as visited
- 2    while  $Q$  not empty do
- 3        $x \leftarrow Q.GetFirst()$
- 4       if  $x \in X_G$
- 5           return SUCCESS
- 6       forall  $u \in U(x)$
- 7            $x' \leftarrow f(x, u)$
- 8           if  $x'$  not visited
- 9              Mark  $x'$  as visited
- 10           $Q.Insert(x')$
- 11       else
- 12              Resolve duplicate  $x'$
- 13   return FAILURE

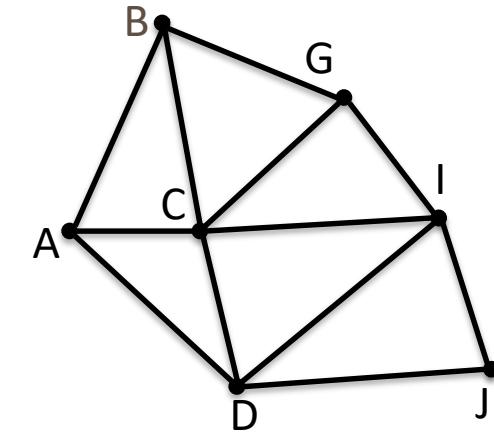
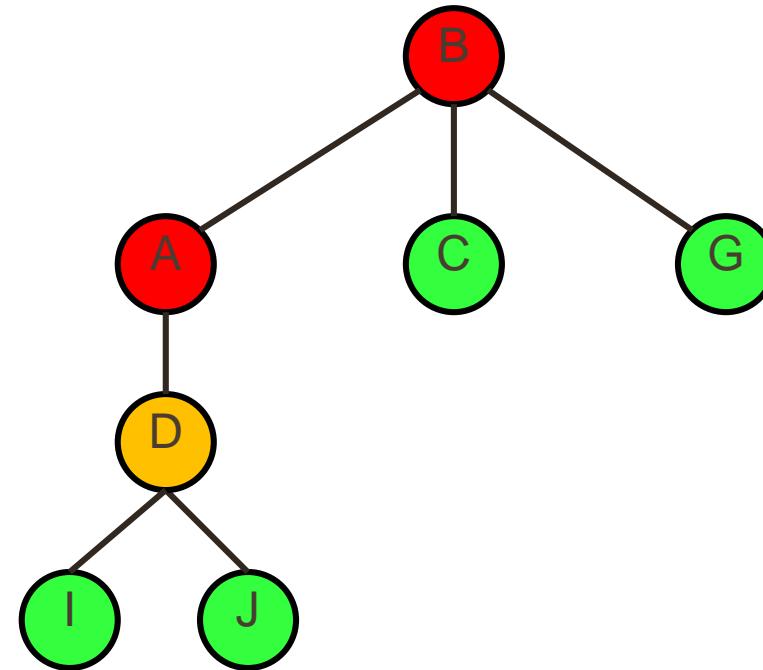
---



Open (Q):  
 $\{D, C, G\}$

Visited:  
 $\{B, A, C, G, D\}$

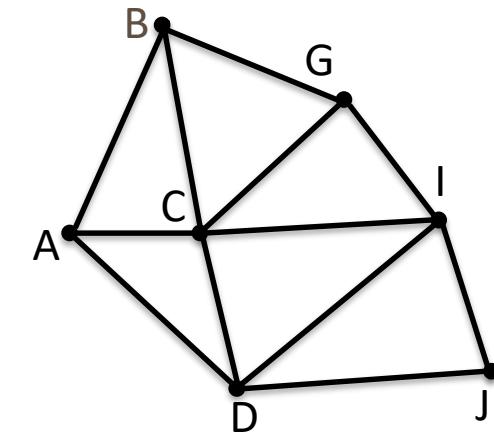
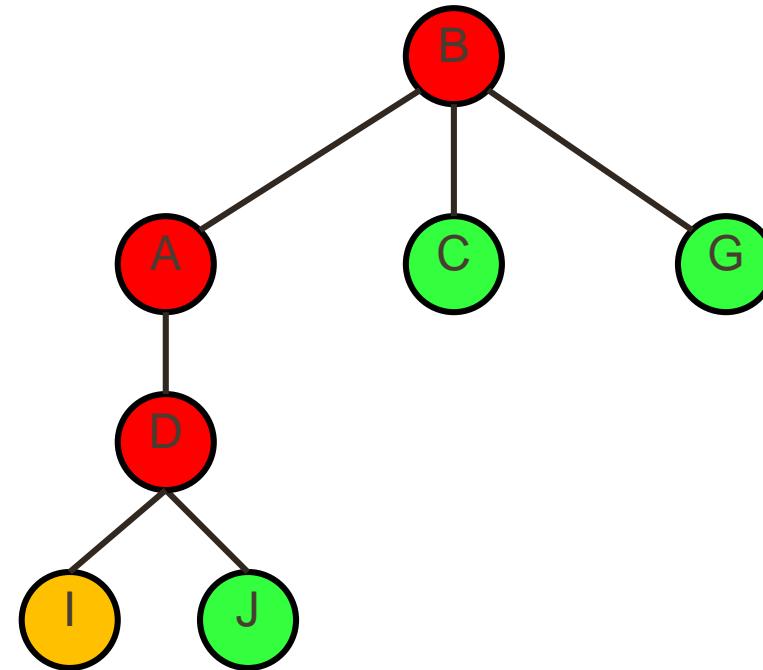
# Depth-First Search Example



Open (Q):  
 $\{I, J, C, G\}$

Visited:  
 $\{B, A, C, G, D, I, J\}$

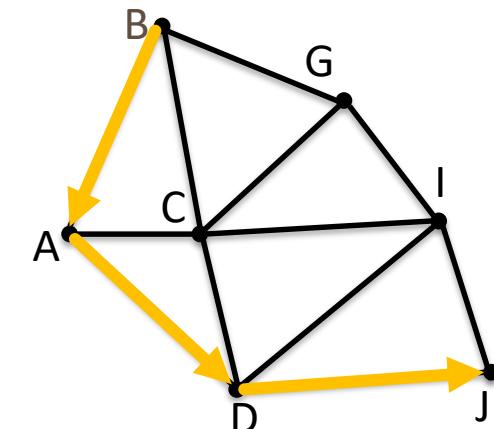
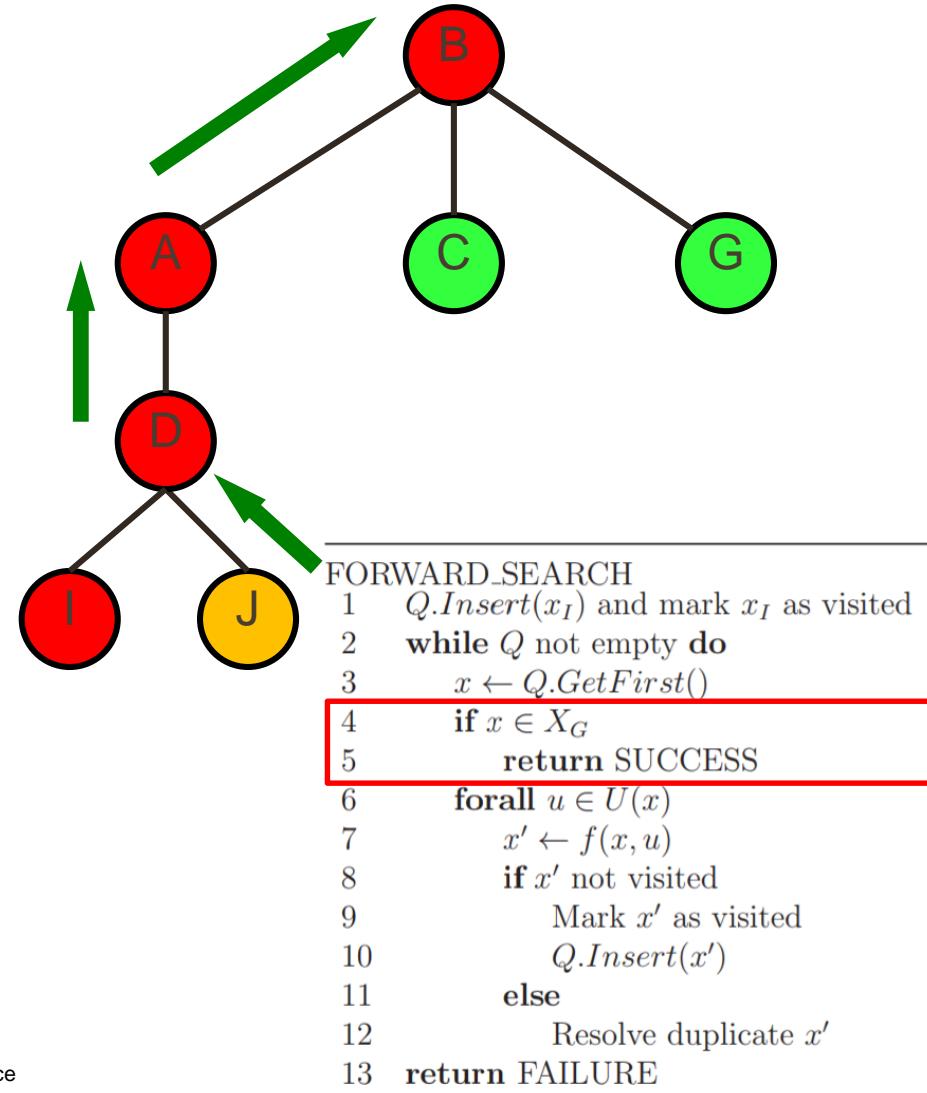
# Depth-First Search Example



Open (Q):  
 $\{J, C, G\}$

Visited:  
 $\{B, A, C, G, D, I, J\}$

# Breadth-First Search Example

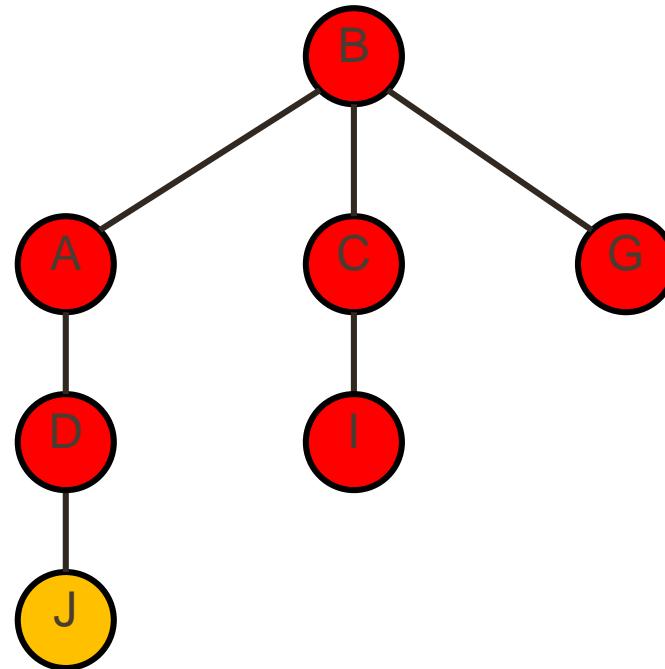


Open ( $Q$ ):  
 {}

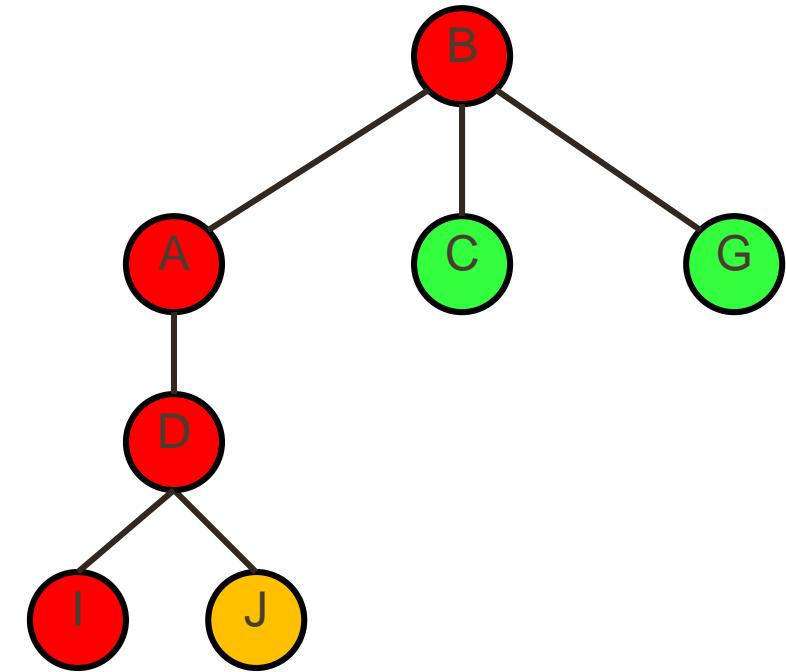
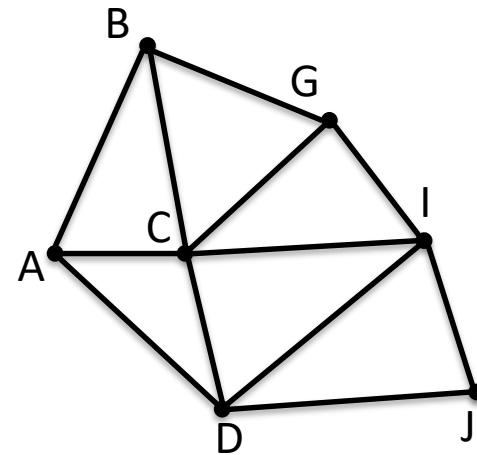
Visited:  
 {B,A,C,G,D,I,J}

Final path solution: B → A → D → J

# Search tree comparison



BFS



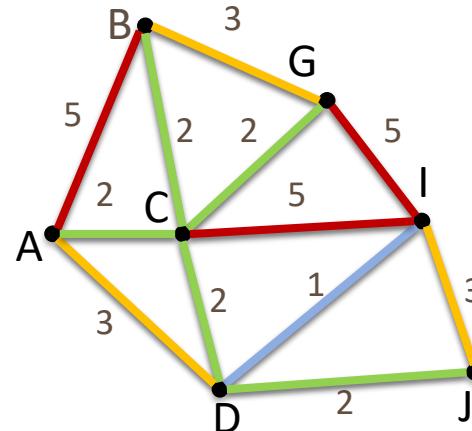
DFS

# Depth-First Search

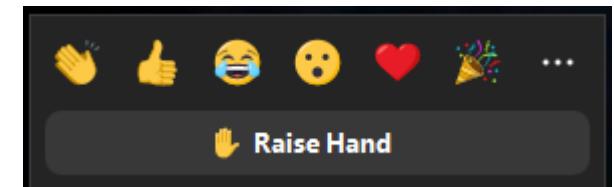
- Lower memory footprint than BFS with high-branching
- Not often used for path search, sometimes used to completely explore a graph
- Both BFS and DFS are simple to implement, but might be inefficient. More complex algorithms are faster, but generally more difficult to implement
- Seems like we want a compromise, search promising paths while we can, then go back up if they aren't working out
- DFS not complete for infinite trees (may explore an incorrect branch infinitely deep, never come back up, BFS *is* complete)

# Costs on Actions

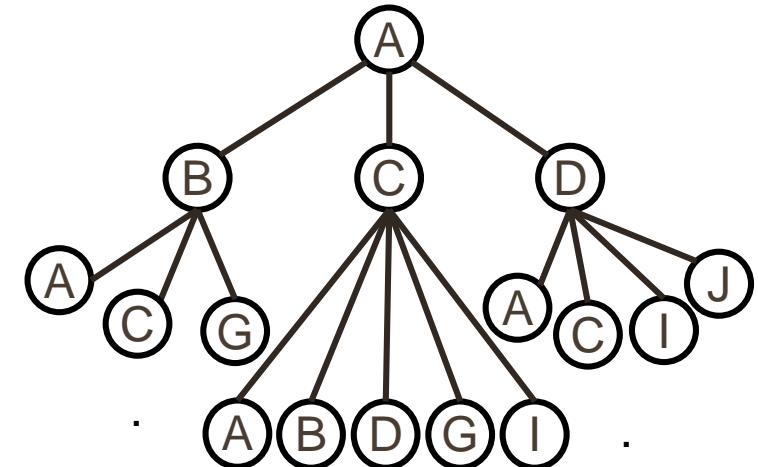
- What about non-uniform edge weights (costs)?



- Dijkstra's Algorithm and A\* search



Does depth in this tree correspond to distance on the graph?



# Dijkstra's Algorithm

- Published by Edsger Dijkstra in 1959
- Basic idea of expanding in order of closest to start (BFS with edge costs)
- One of the most commonly used routing algorithms in graph traversal problems
- Asymptotically the fastest known single-source shortest path algorithm for arbitrary directed graphs
- **Open queue is ordered according to currently known best cost to arrive**

# Dijkstra's Algorithm Example

B(0)

---

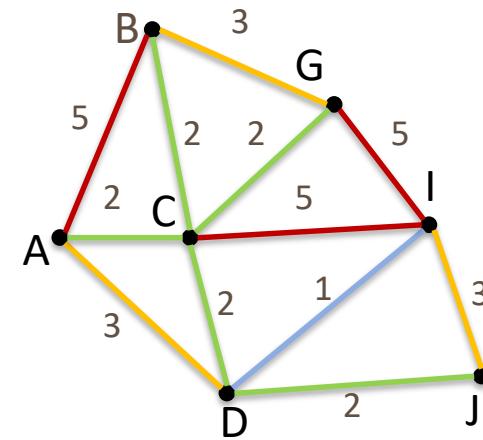
## FORWARD SEARCH

```

1  Q.Insert( $x_1$ )
2  while  $Q$  not empty do
3       $x \leftarrow Q.GetFirst()$ , R.Insert( $x$ )
4      if  $x \in X_G$ 
5          return SUCCESS
6      forall  $u \in U(x)$ 
7           $x' \leftarrow f(x, u)$ 
8          if  $x' \notin Q$  or R
9              Mark  $x'$  as visited
10             Q.Insert( $x'$ )
11         else
12             Resolve duplicate  $x'$ 
13     return FAILURE

```

---



Open (Q):

{B(0)}

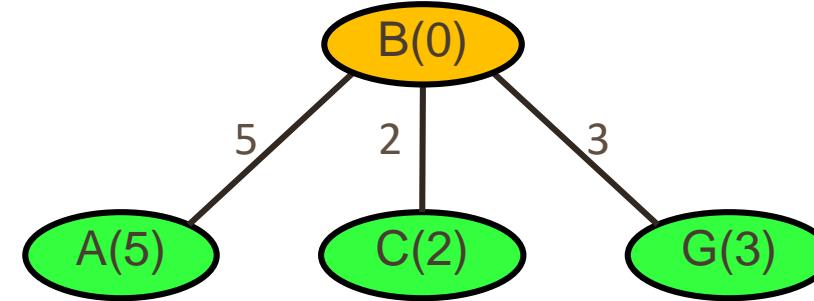
Closed (R):

{B(0)}

Our Dijkstra queue will be ordered by cost to arrive:

- push ( $Q.Insert$ ) by cost
- pop ( $Q.GetFirst$ ) from the front, and add it to the closed list

# Dijkstra's Algorithm Example



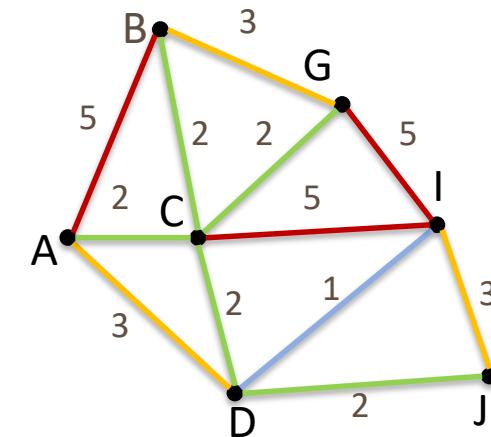

---

```

FORWARD_SEARCH
1   Q.Insert( $x_I$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ ,  $R.Insert(x)$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or  $R$ 
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11          else
12              Resolve duplicate  $x'$ 
13  return FAILURE

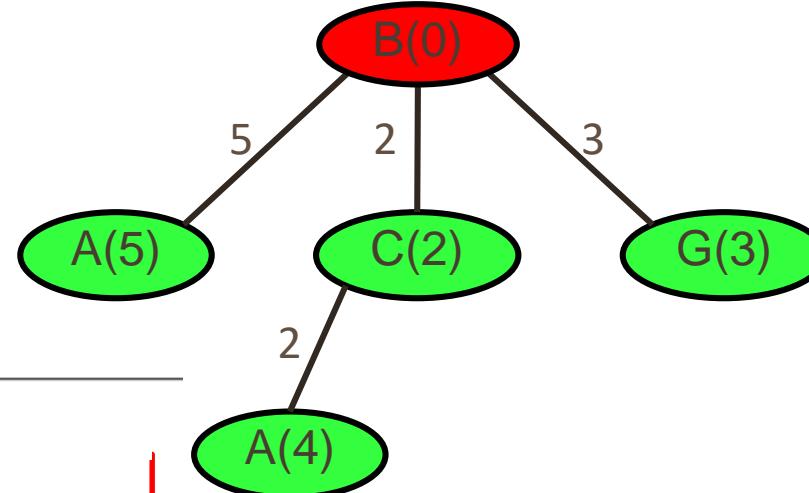
```

---



Open ( $Q$ ):      Closed:  
{ C (2),            { B (0) }  
  G (3),  
  A (5) }

# Dijkstra's Algorithm Example



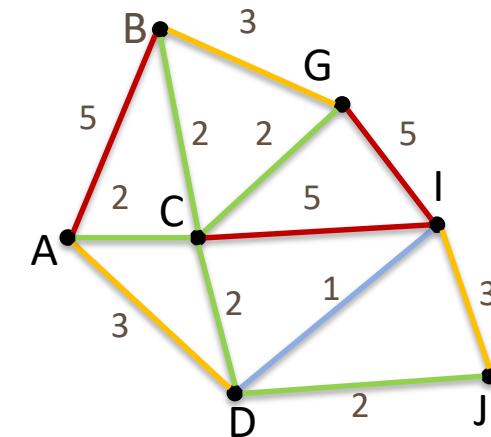

---

```

FORWARD_SEARCH
1   Q.Insert( $x_1$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ ,  $R.Insert(x)$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or  $R$ 
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11          else
12              Resolve duplicate  $x'$ 
13  return FAILURE

```

---



Open ( $Q$ ):      Closed:  
{ G (3),      { B (0),  
  A (4) }      C (2) }

# Dijkstra's Algorithm Example

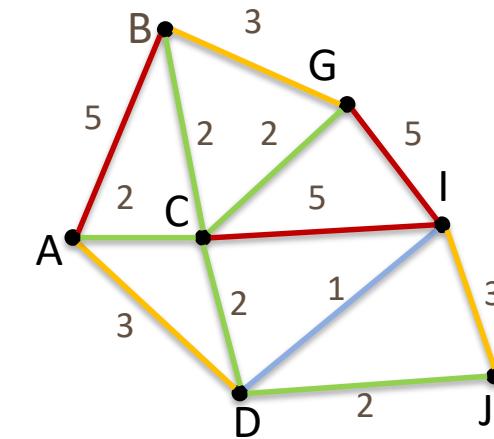
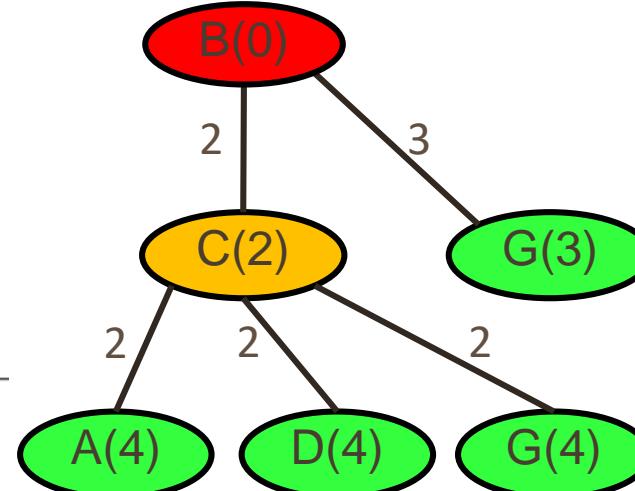
---

```

FORWARD_SEARCH
1   Q.Insert( $x_I$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ ,  $R.Insert(x)$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or  $R$ 
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11          else
12              Resolve duplicate  $x'$ 
13  return FAILURE

```

---



Open ( $Q$ ):  
{ G (3),  
 A (4),  
 D (4) }

Closed:  
{ B (0),  
 C (2) }

# Dijkstra's Algorithm Example

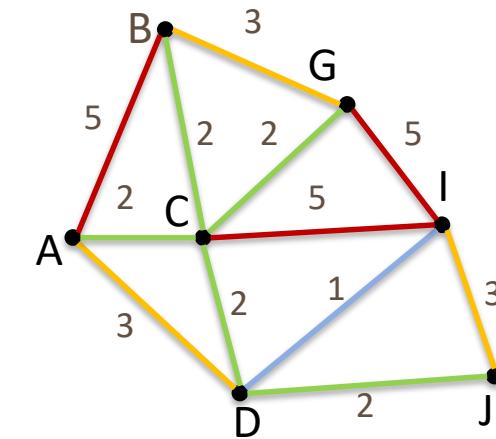
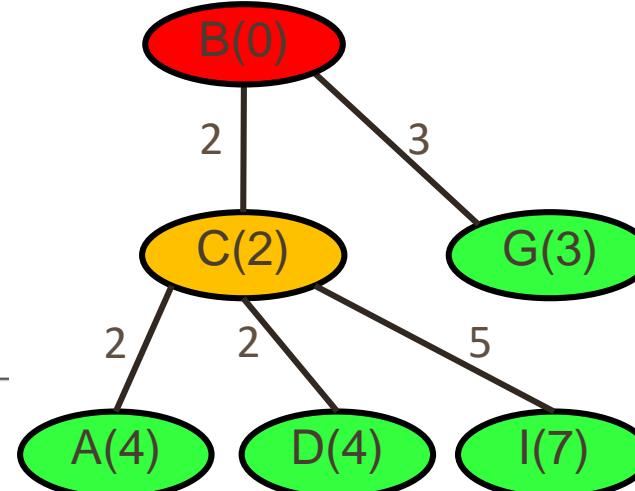
---

```

FORWARD_SEARCH
1   Q.Insert( $x_I$ )
2   while  $Q$  not empty do
3        $x \leftarrow Q.GetFirst()$ ,  $R.Insert(x)$ 
4       if  $x \in X_G$ 
5           return SUCCESS
6       forall  $u \in U(x)$ 
7            $x' \leftarrow f(x, u)$ 
8           if  $x' \notin Q$  or  $R$ 
9               Mark  $x'$  as visited
10               $Q.Insert(x')$ 
11          else
12              Resolve duplicate  $x'$ 
13  return FAILURE

```

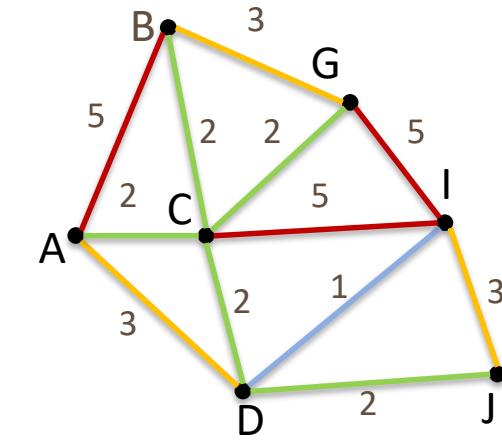
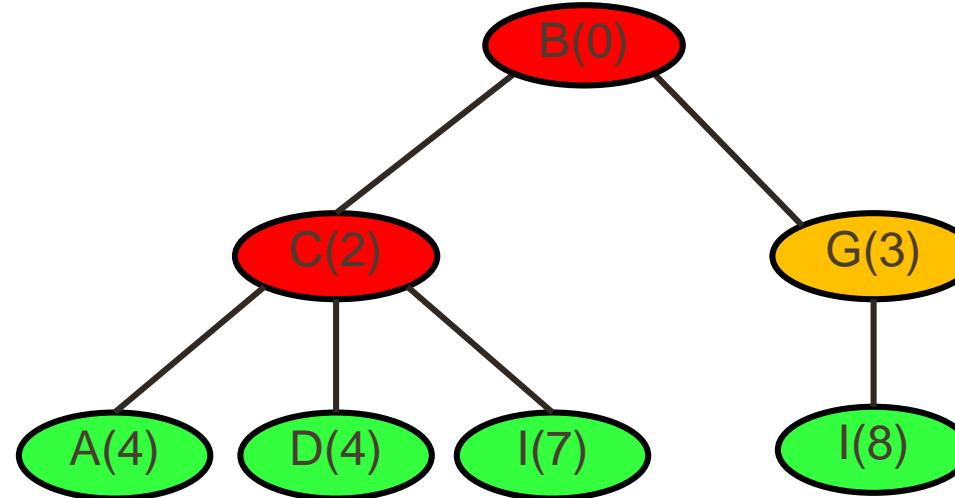
---

Open ( $Q$ ):
 $\{ G(3),$   
 $A(4),$   
 $D(4),$   
 $I(7) \}$ 

Closed:

 $\{ B(0),$   
 $C(2) \}$

# Dijkstra's Algorithm Example



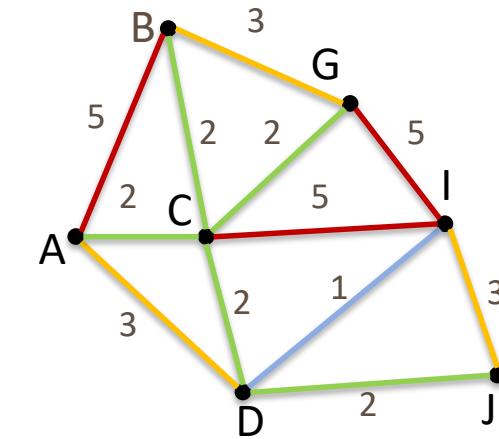
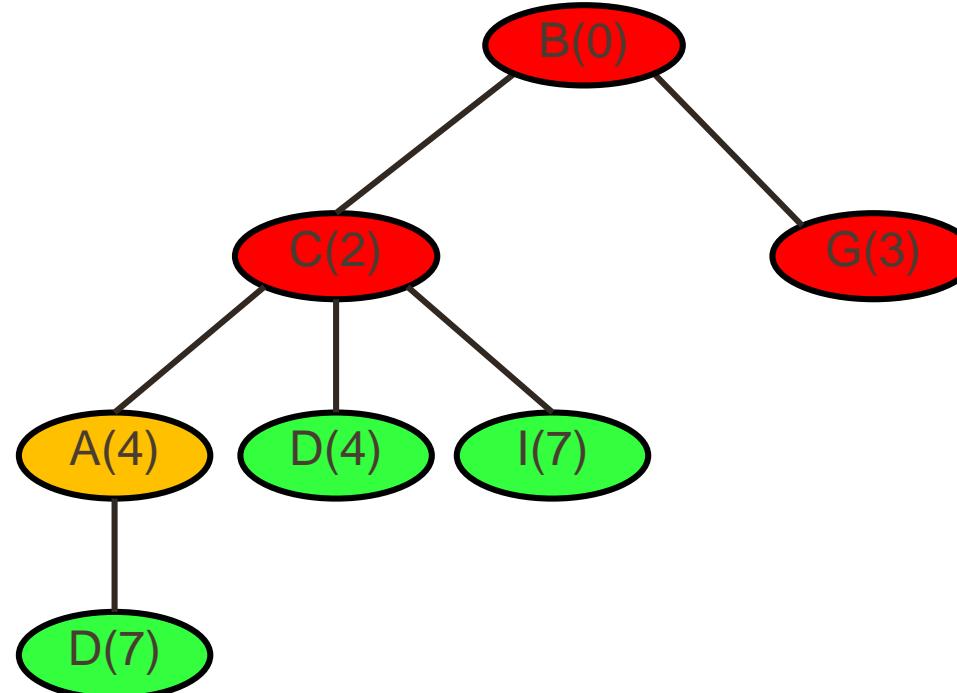
Open (Q):

{ A (4),  
D (4),  
I (7) }

Closed:

{ B (0),  
C (2),  
G (3) }

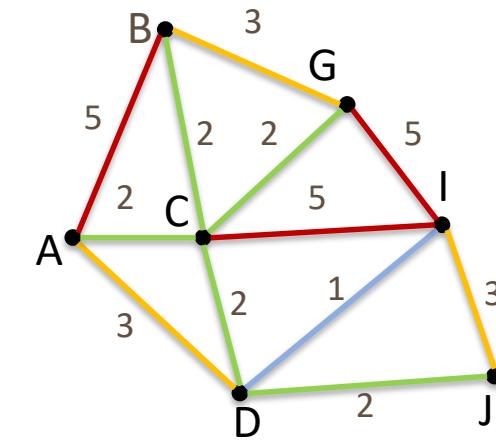
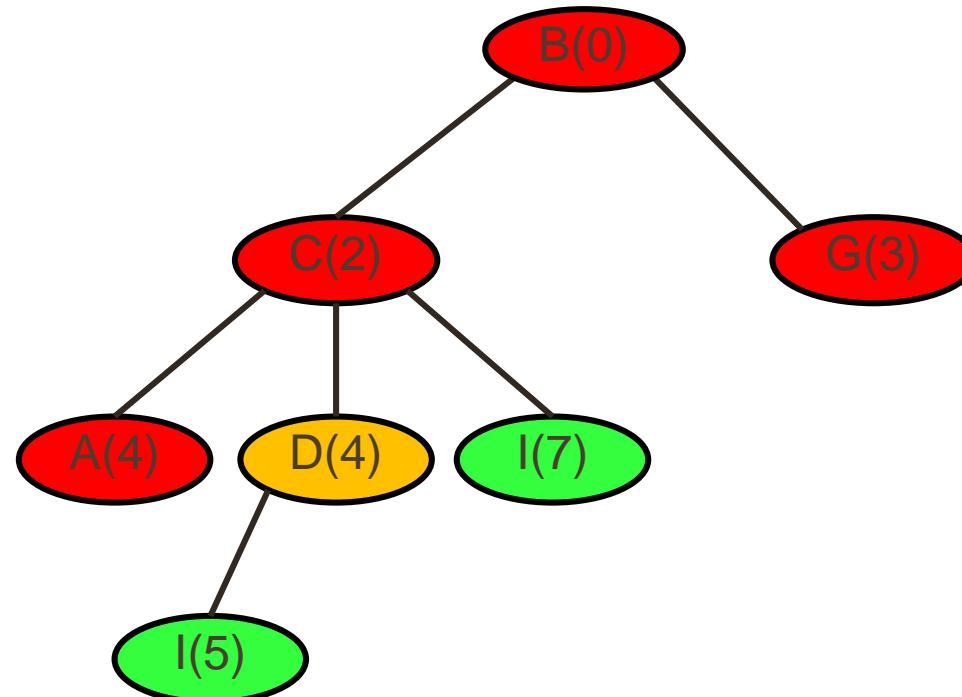
# Dijkstra's Algorithm Example



Open (Q):  
{ D (4),  
I (7) }

Closed:  
{ B (0),  
C (2),  
G (3),  
A (4) }

# Dijkstra's Algorithm Example



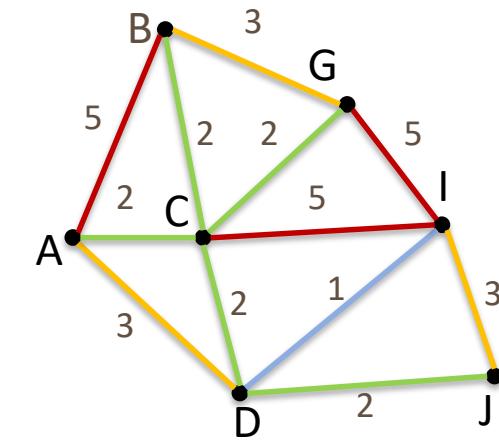
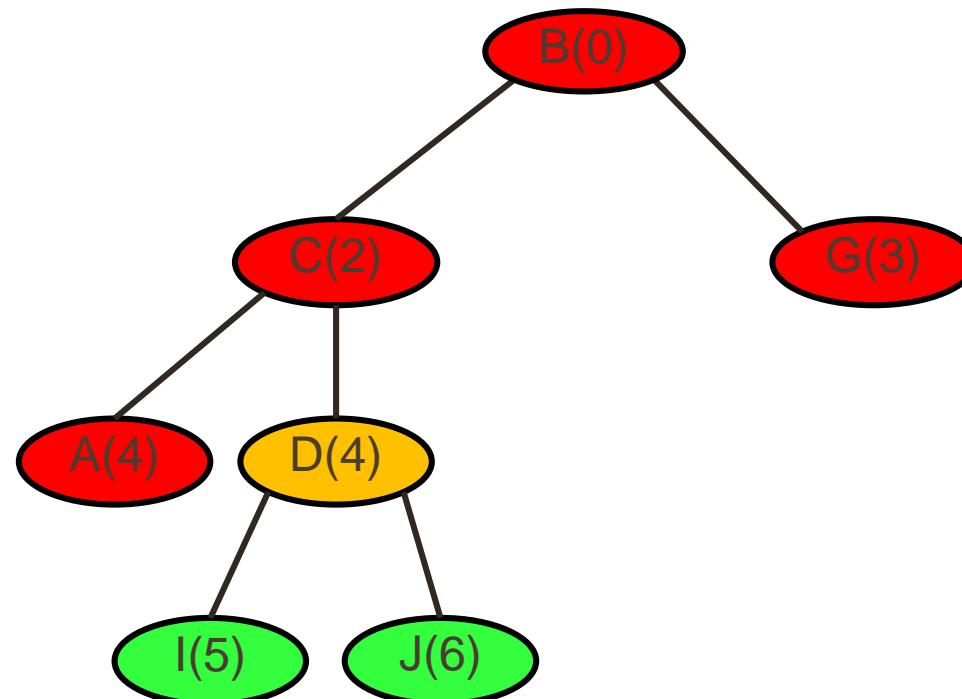
Open (Q):

{ I (5) }

Closed:

{ B (0),  
C (2),  
G (3),  
A (4),  
D (4) }

# Dijkstra's Algorithm Example



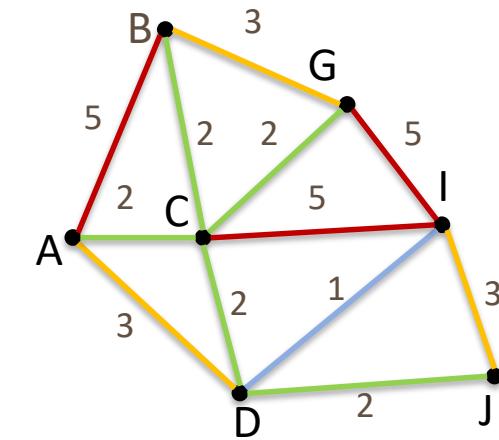
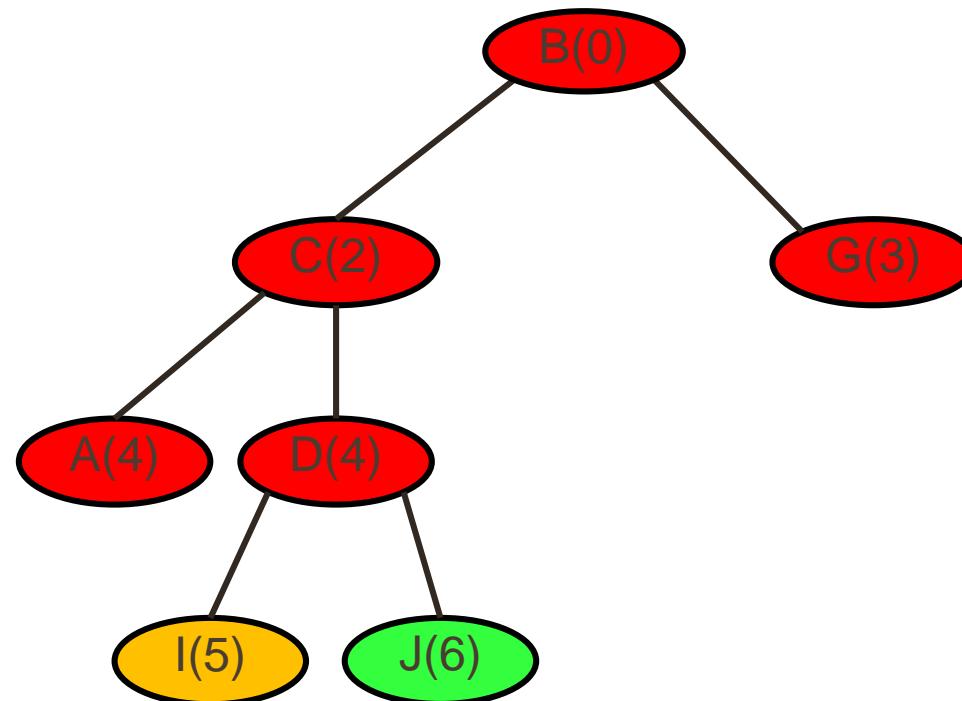
Open (Q):

{ I (5) ,  
J (6) }

Closed:

{ B (0),  
C (2),  
G (3),  
A (4),  
D (4) }

# Dijkstra's Algorithm Example



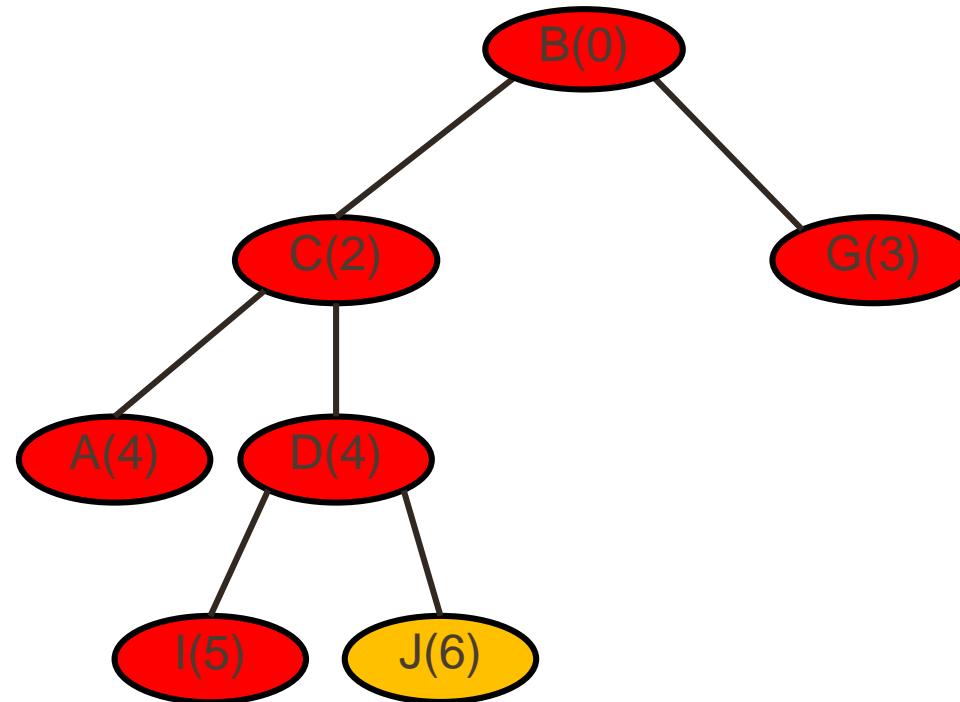
Open (Q):

{ J (6) }

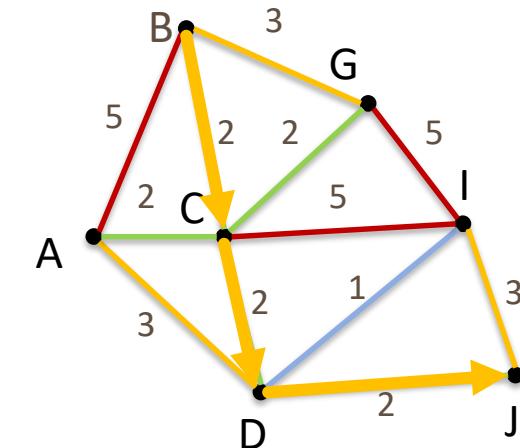
Closed:

{ B (0),  
C (2),  
G (3),  
A (4),  
D (4),  
I (5) }

# Dijkstra's Algorithm Example



Final path solution: B → C → D → J  
with path cost 6



Open (Q):  
{ }

Closed:  
{ B (0),  
C (2),  
G (3),  
A (4),  
D (4),  
I (5),  
J (6) }

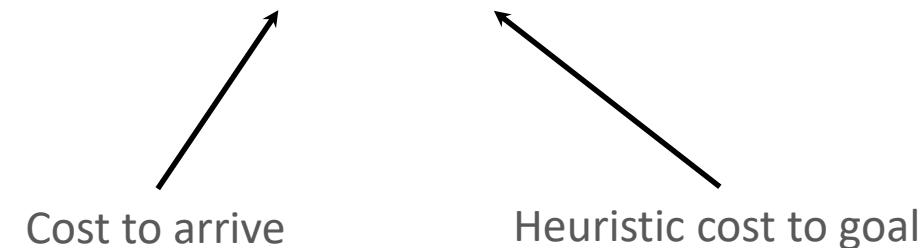
# Dijkstra's Algorithm

- At the end, we can recover the lowest-cost route from the start to any node (or any node with cost  $<$  goal if we terminate at a goal)
- Quite easy to implement, but requires a little bit of careful management with the priority queue
- Doesn't really know the goal exists until it reaches it
  - Could we guide the search to expand nodes that are closer to the goal earlier?
  - Can we do it without breaking the condition that a node is only accepted with its lowest cost of arrival?

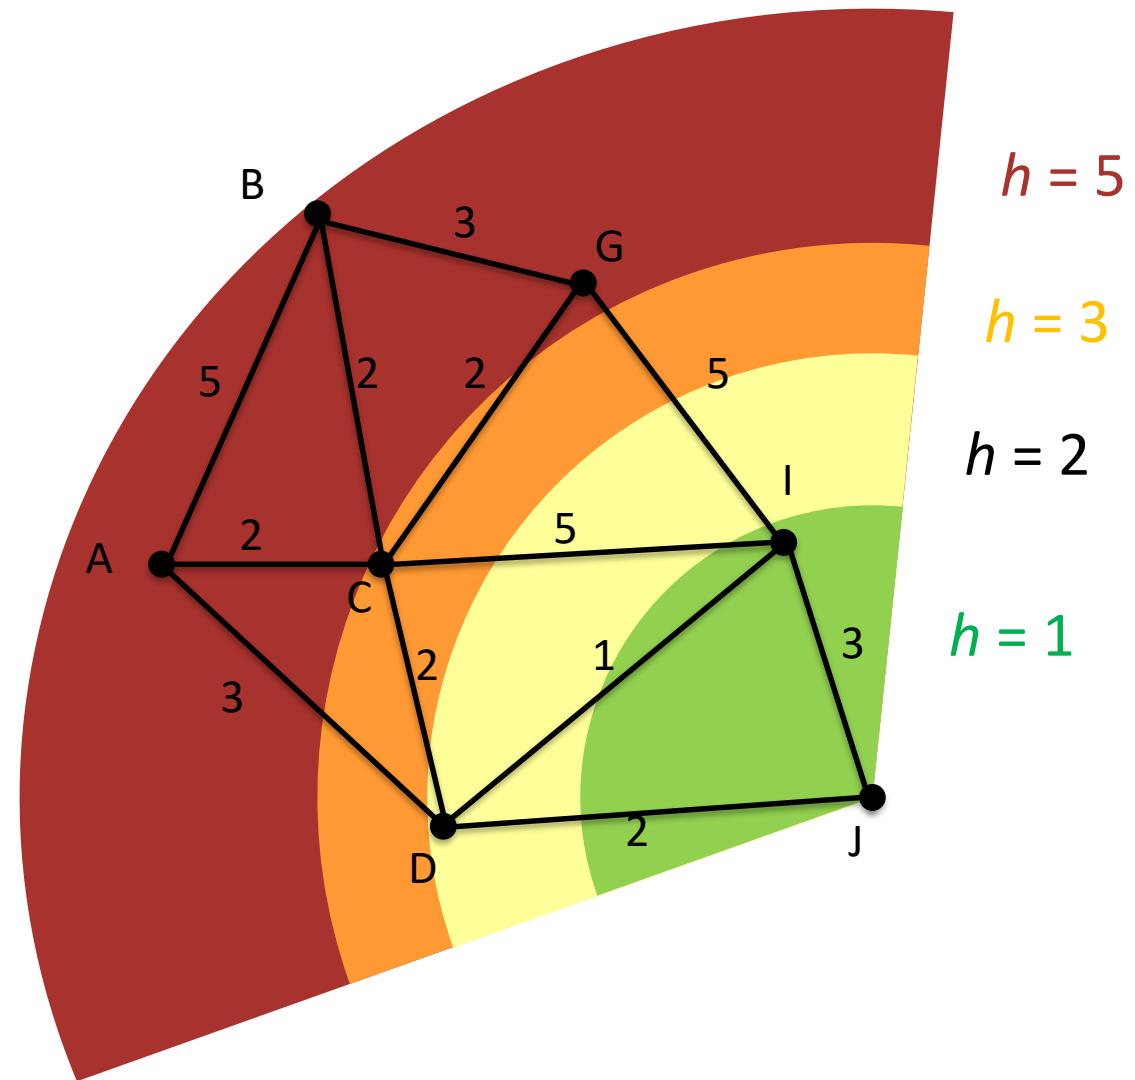
# A\* Heuristic Search

- Heuristic:
  - Any *optimistic estimate* of how close a state is to a goal
  - Designed for a particular search problem
  - Examples: Manhattan distance, Euclidean distance
- A\* Priority:

$$f(n) = g(n) + h(n)$$

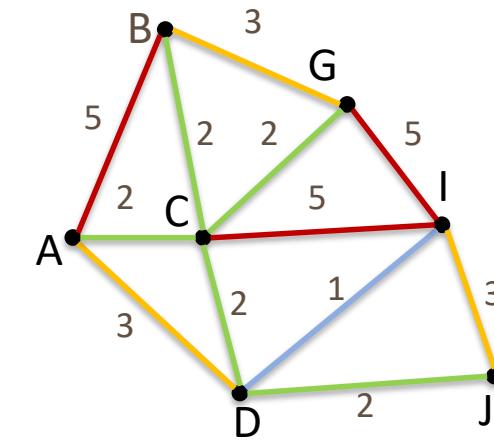


# A\* Heuristic



# A\* Algorithm Example

B(0)



Open (Q):

{B(0)}

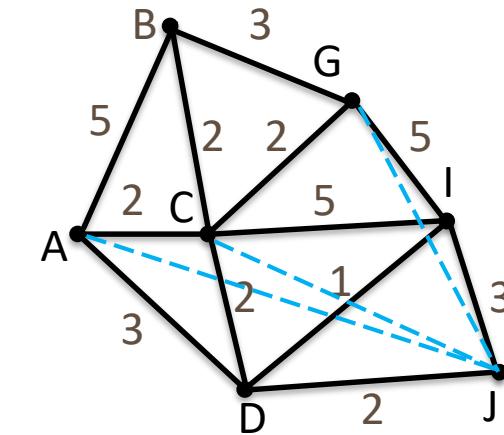
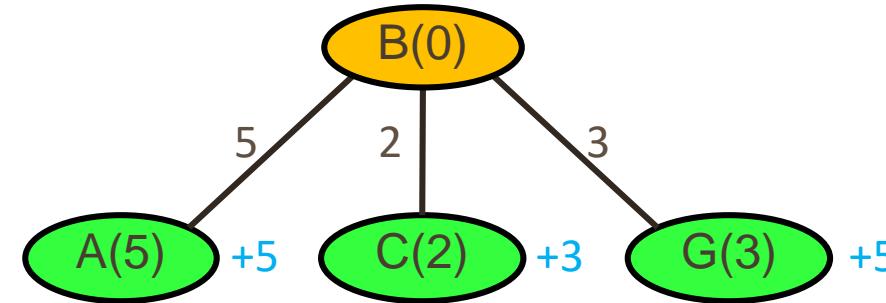
Closed:

{B(0)}

Our A\* queue will be ordered by cost to arrive + heuristic:

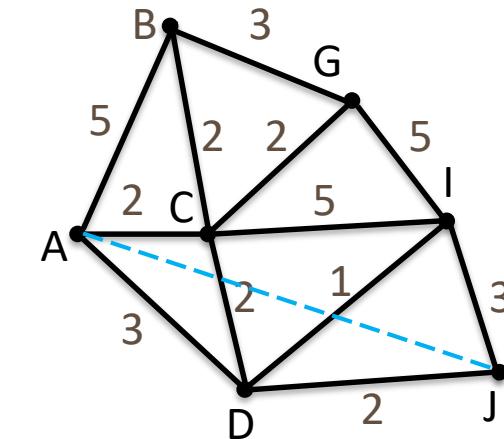
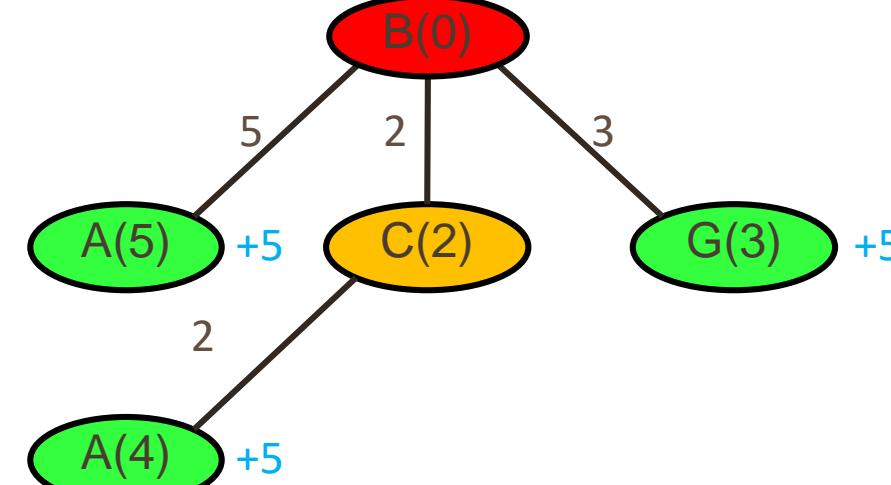
- push ( $Q.Insert$ ) by A\* priority,  $f(n)$
- pop ( $Q.GetFirst$ ) from the front, and add it to the closed list

# A\* Algorithm Example



Open (Q):      Closed:  
{ C (2+3),      { B (0) }  
  G (3+5),  
  A (5+5) }

# A\* Algorithm Example



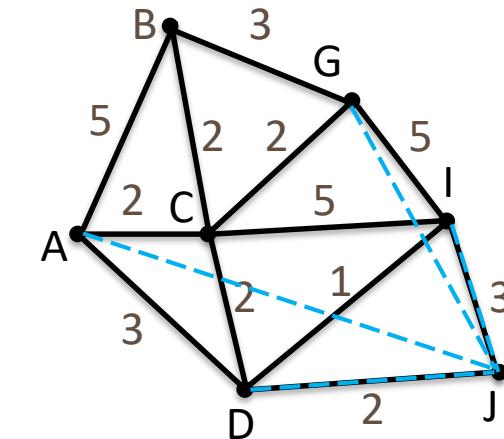
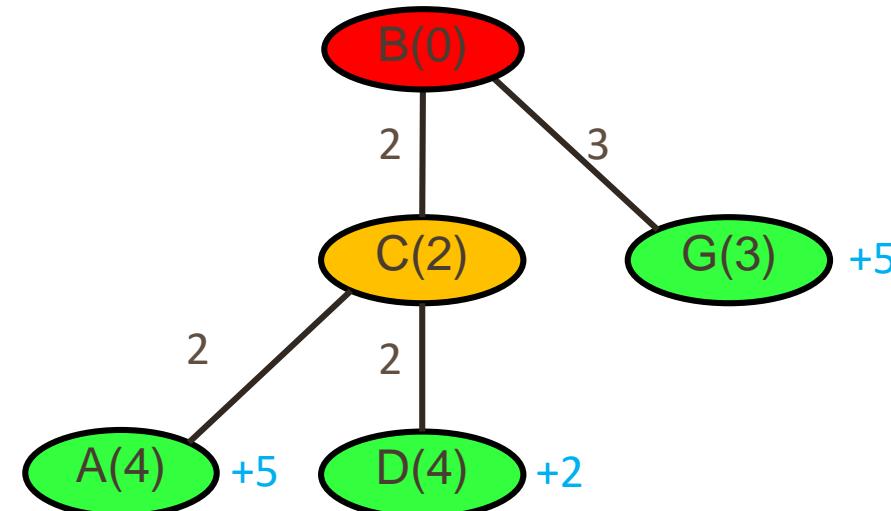
Open (Q):

{ G (3+5),  
A (4+5) }

Closed:

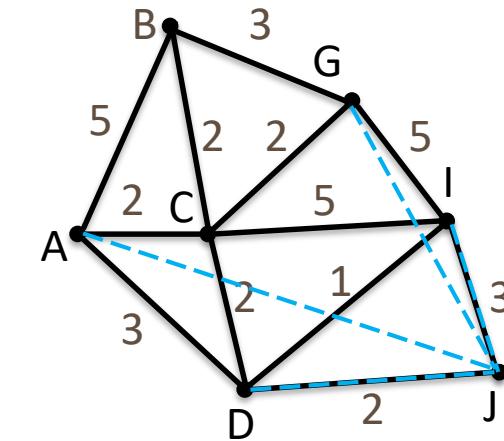
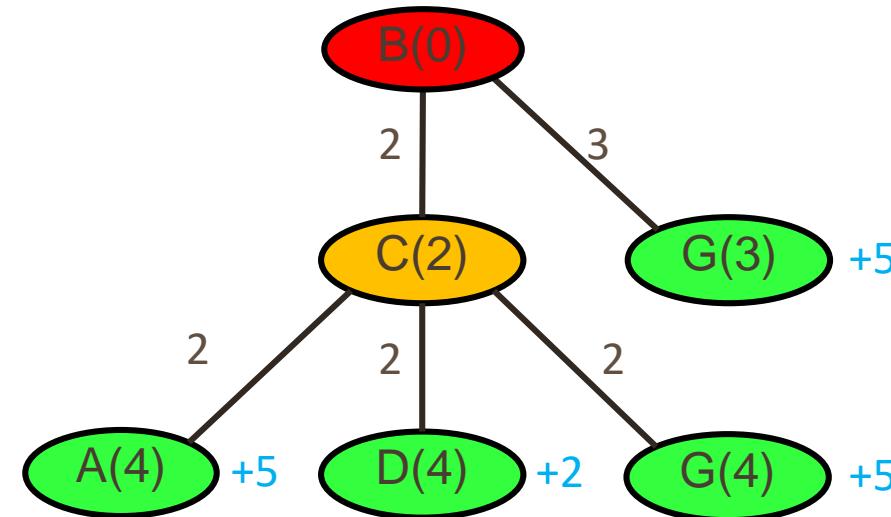
{ B (0) ,  
C (2) }

# A\* Algorithm Example



Open (Q):      Closed:  
 $\{ D (4+2),$   
 $G (3+5),$   
 $A (4+5) \}$

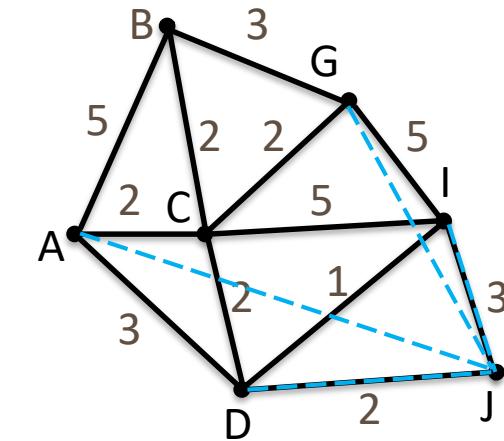
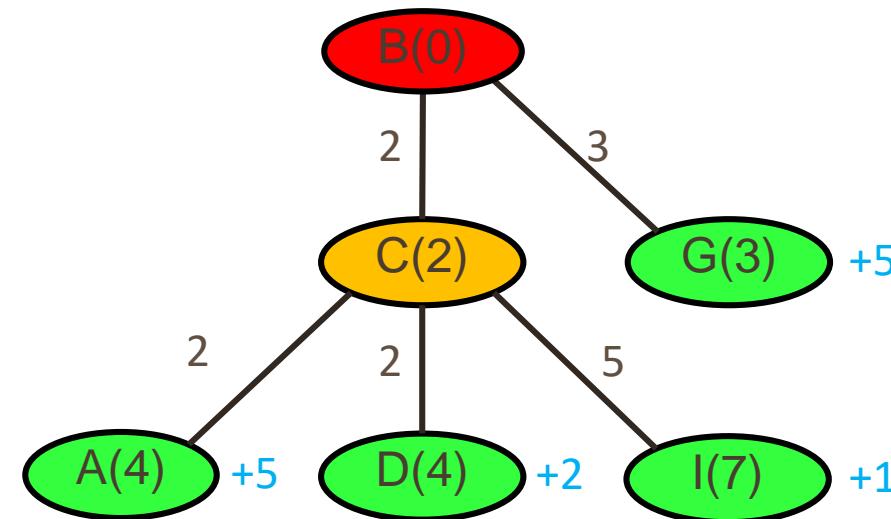
# A\* Algorithm Example



Open (Q):  
{ D (4+2),  
G (3+5),  
A (4+5) }

Closed:  
{ B (0) ,  
C (2) }

# A\* Algorithm Example



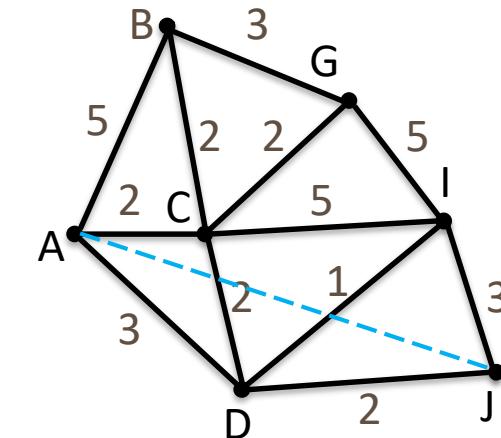
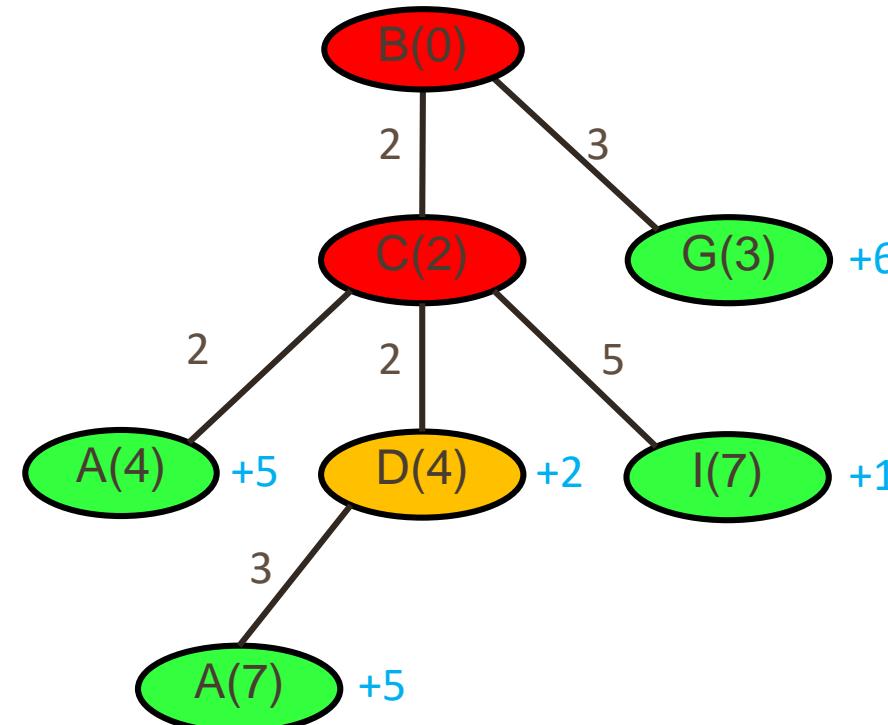
Open (Q):

{ D (4+2),  
I (7+1),  
G (3+5),  
A (4+5) }

Closed:

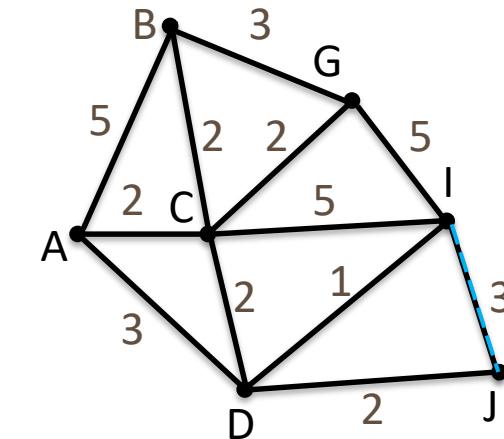
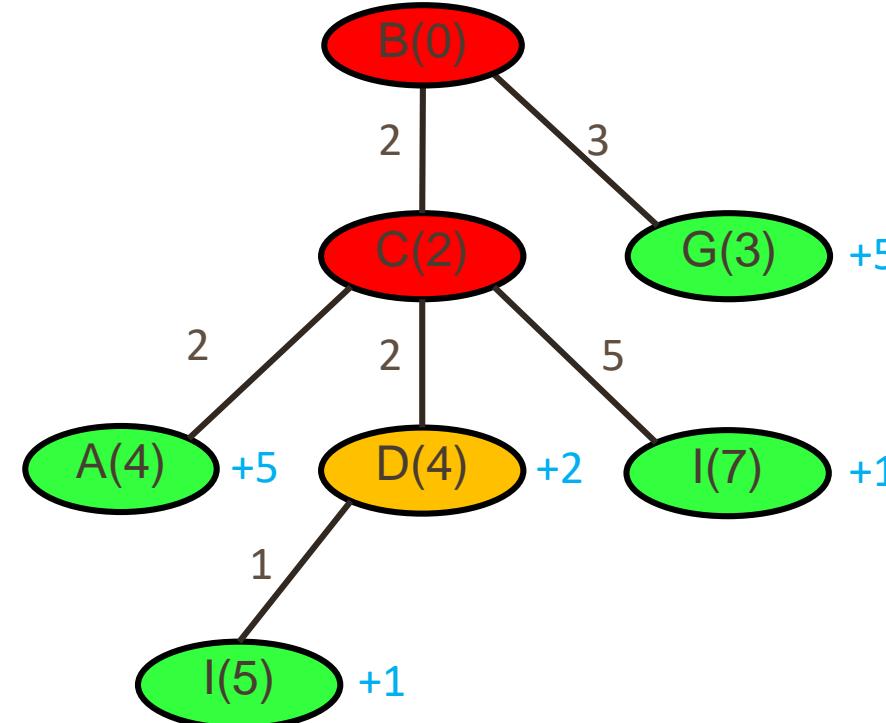
{ B (0) ,  
C (2) }

# A\* Algorithm Example



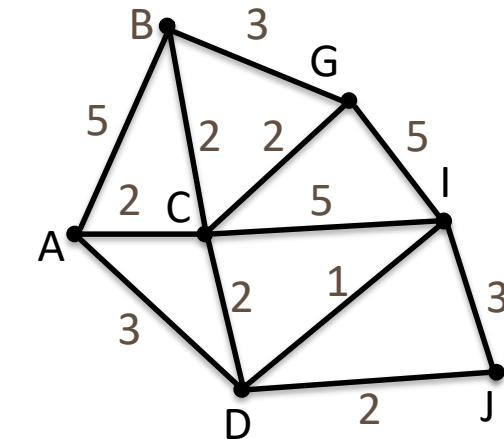
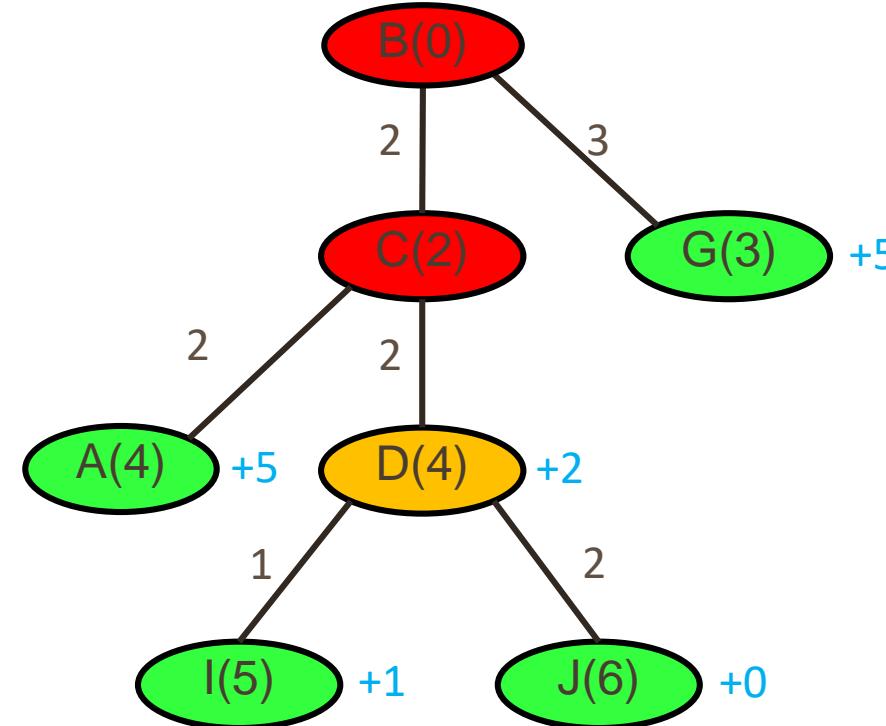
Open (Q): { I (5+1),  
G (3+5),  
A (4+5) }      Closed: { B (0) ,  
C (2),  
D (4) }

# A\* Algorithm Example



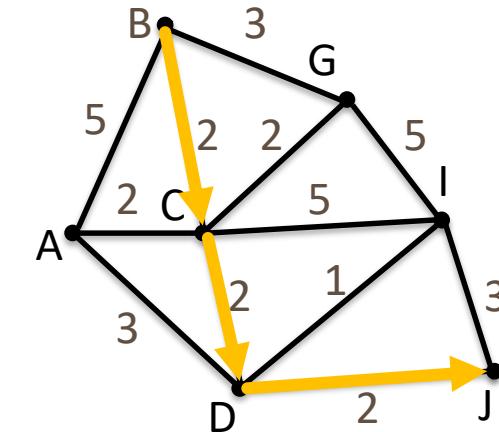
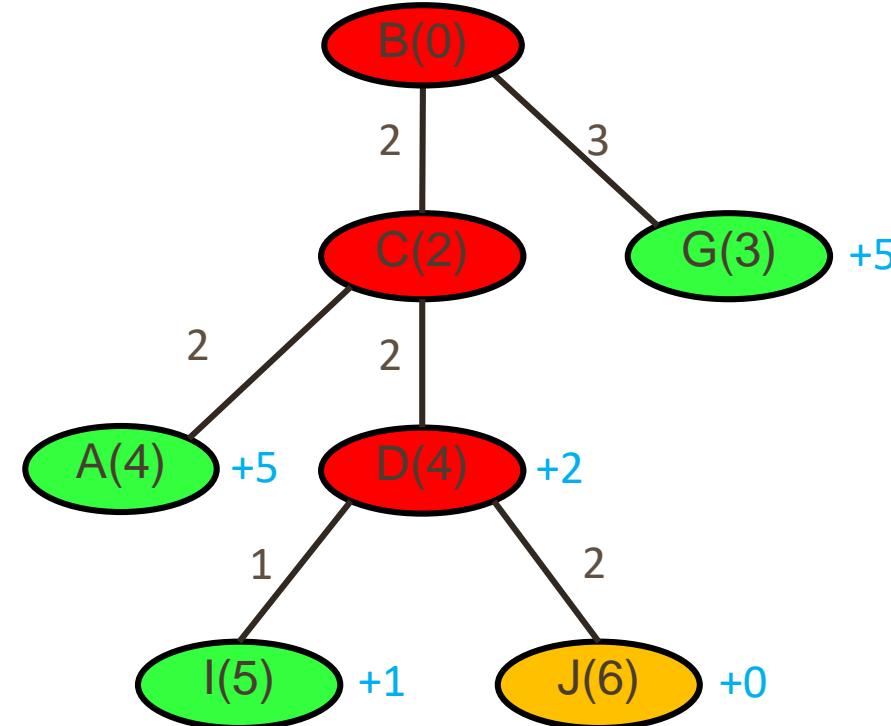
Open (Q):      Closed:  
{ I (5+1),      { B (0) ,  
  G (3+5),      C (2),  
  A (4+5) }      D (4) }

# A\* Algorithm Example



Open (Q):      Closed:  
{ J (6+0),      { B (0) ,  
  I (5+1),      C (2),  
  G (3+5),      D (4) }  
  A (4+5) }

# A\* Algorithm Example



Open (Q): { I (5+1),  
G (3+5),  
A (4+5) }  
Closed: { B (0) ,  
C (2),  
D (4),  
J (6) }

Final path solution: B → C → D → J  
with path cost 6

# A\* Heuristic

- The heuristic must be **admissible**
  - It never overestimates the cost

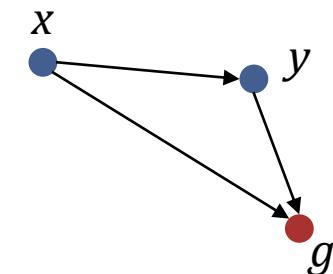
True cost to goal

$$h(x) \leq d(x, \text{goal})$$

- The heuristic must be **consistent**
  - For any pair of adjacent nodes  $x$  and  $y$ , where  $d(x,y)$  is the cost of edge between them

$$h(x) \leq d(x, y) + h(y)$$

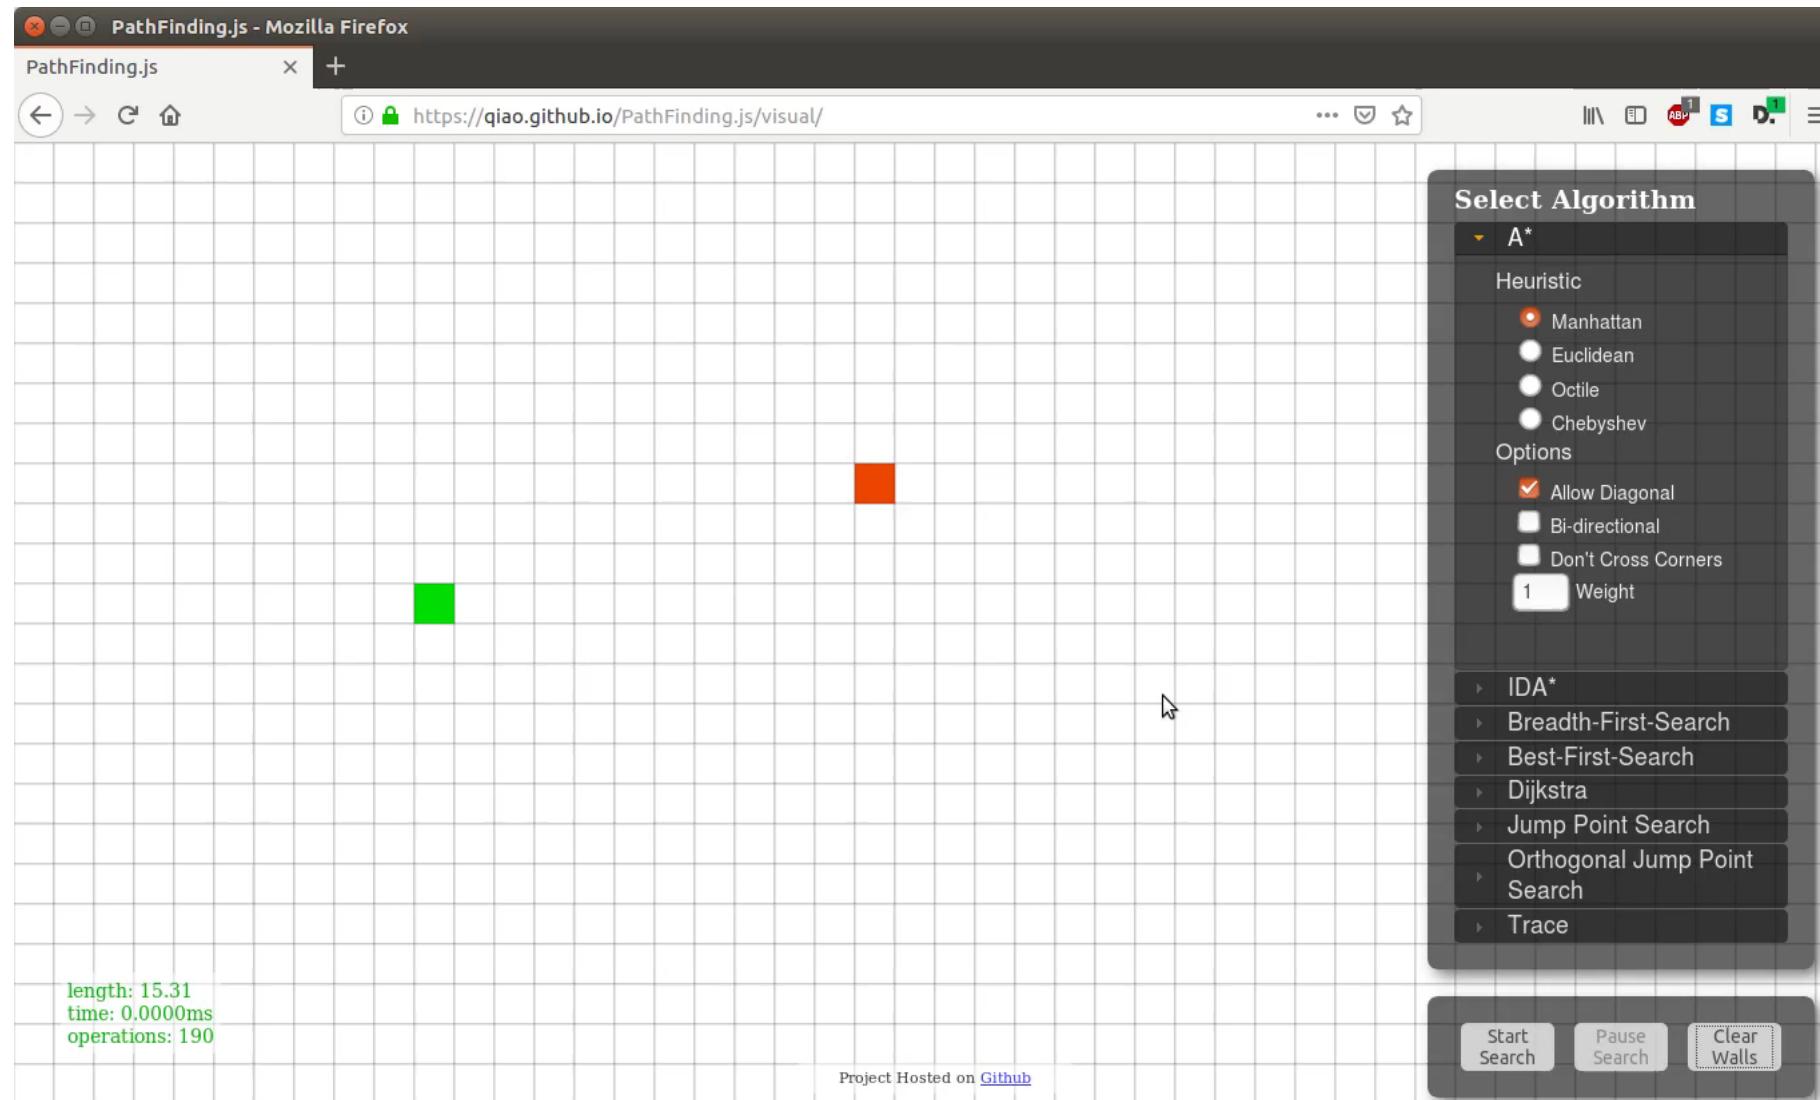
- Typical valid heuristics:
  - Euclidean distance
  - Manhattan distance
  - Zero (Dijkstra's algorithm)



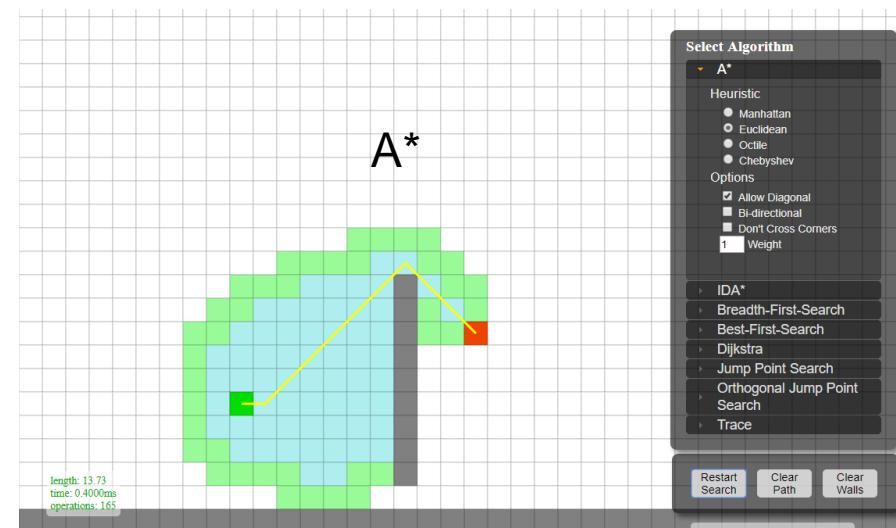
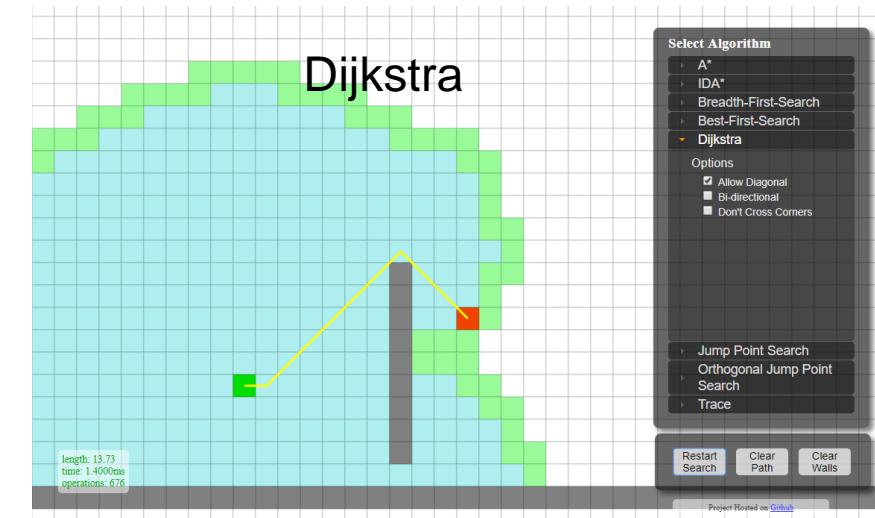
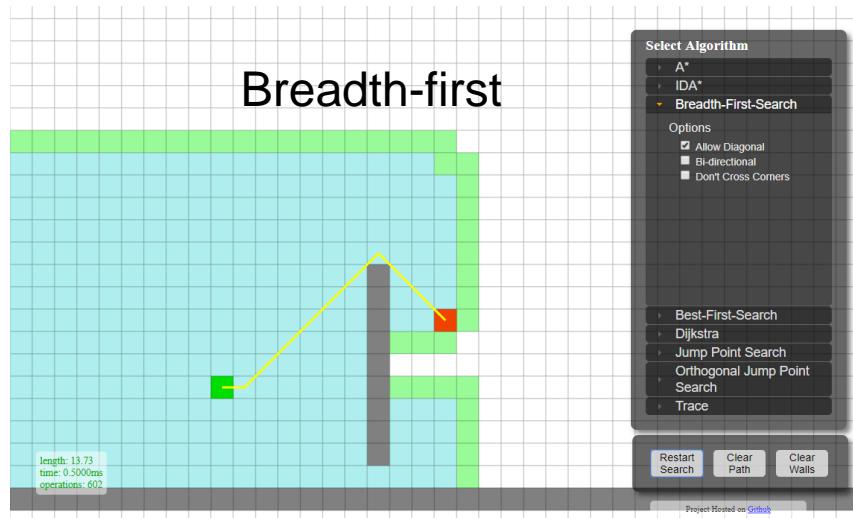
# A\* Search Algorithm

- A\* is an extension of Dijkstra's algorithm, and achieves faster performance by using heuristics
- **Best-first search:** A\* traverses a graph following a path of lowest expected total cost or distance
- The cost function is a sum of two functions:
  - Past path-cost function, which is a **known** cost from the starting node to the current node
  - Future path-cost function, which is a “heuristic estimate” of the distance from the current node to the goal

# <https://qiao.github.io/PathFinding.js/visual/>

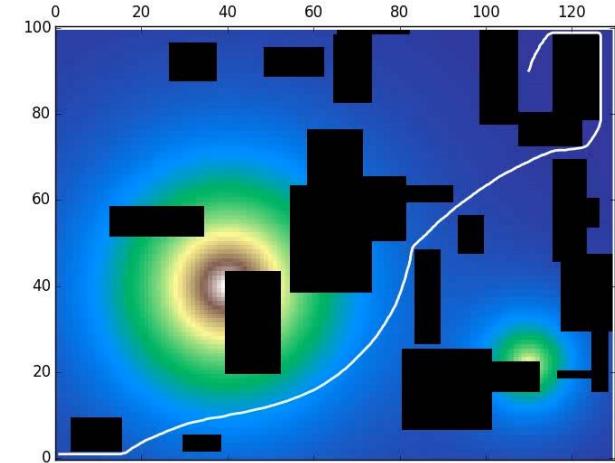
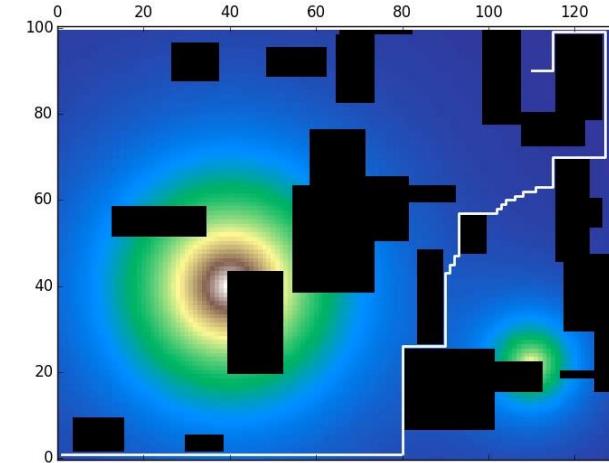


# <https://qiao.github.io/PathFinding.js/visual/>



# Limitations

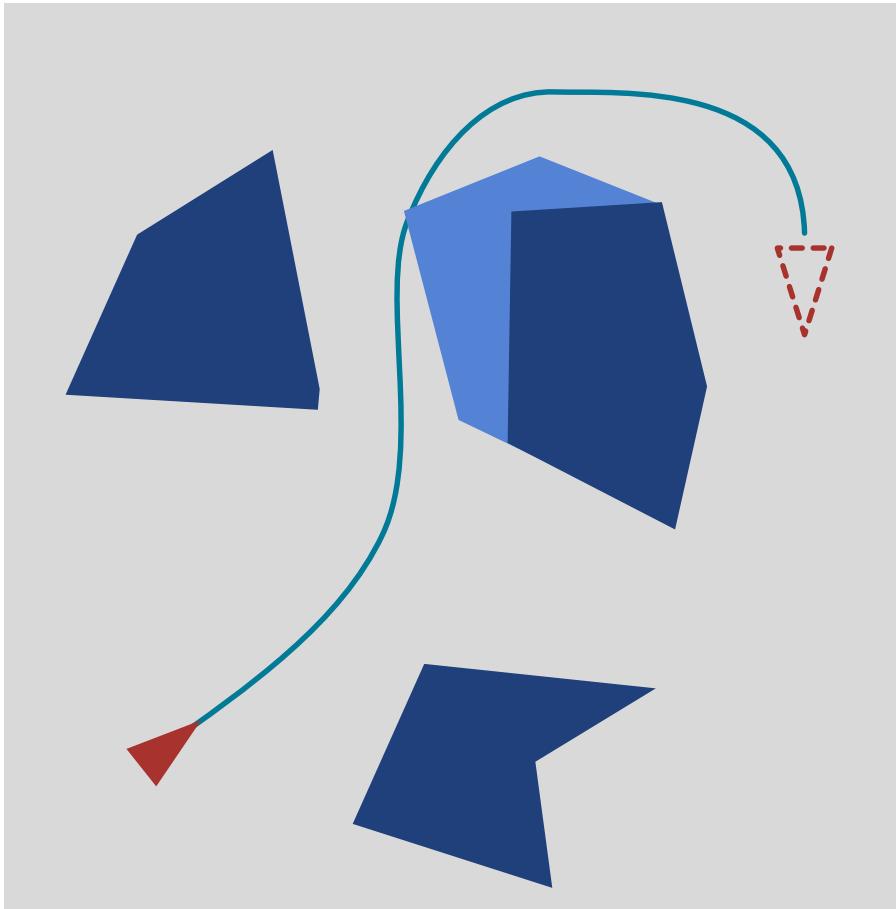
- A\* is very commonly used in robot planning, especially for low-dimensional state spaces
- Limitations:
  - You need to construct a graph
  - Sometimes an admissible heuristic function is difficult to find (as hard as the problem)
  - A grid may not be a good representation of your problem



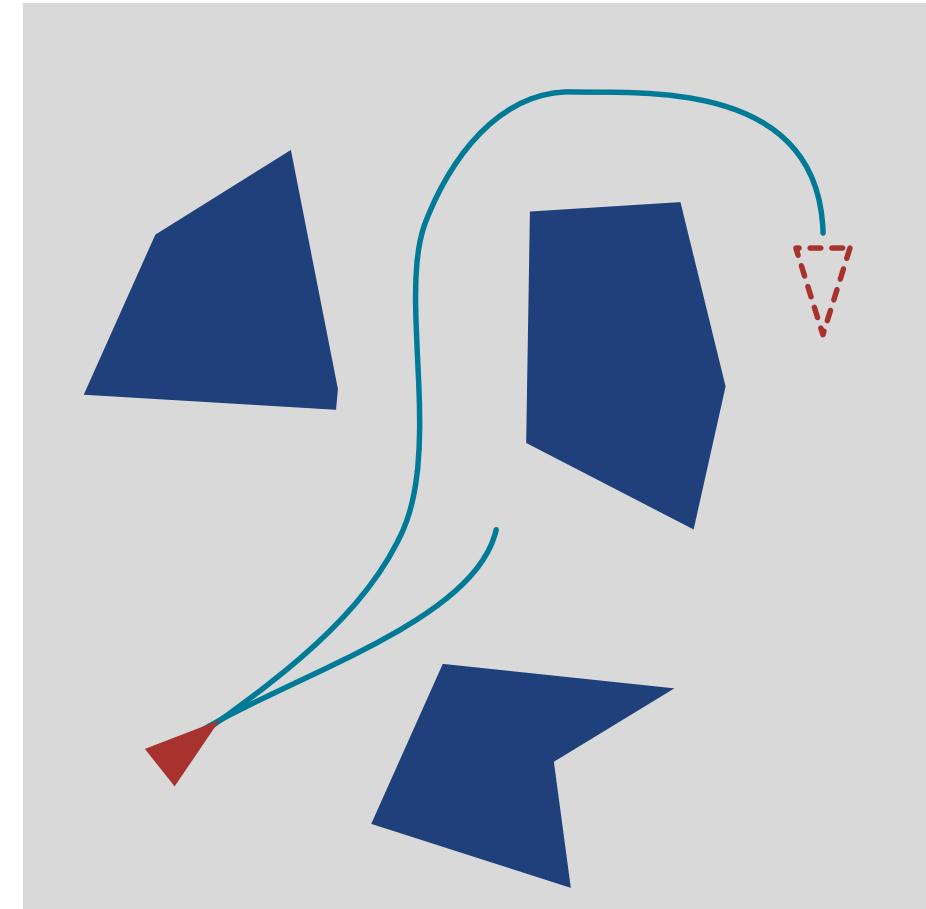
# What happens when it doesn't all go to plan?

- So far, we've basically ignored the plan **execution**
- Our plans are a series of states that we assume the robot is capable of visiting sequentially and reliably
- However, there are many reasons that, in practice, this may not always be the case

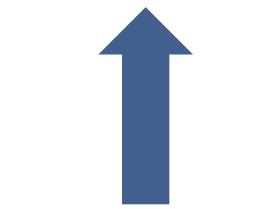
- Environment uncertainty



- Motion uncertainty



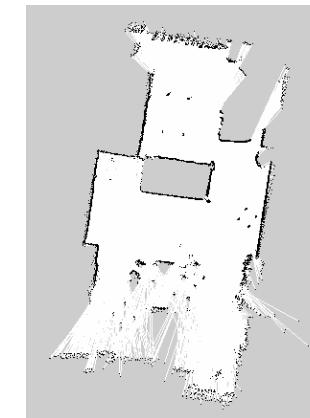
# Planning hierarchy



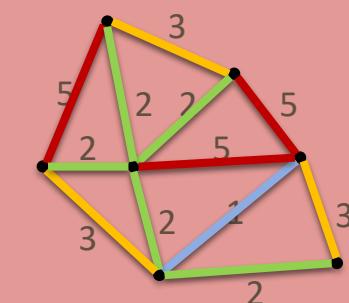
Sensor data



Map

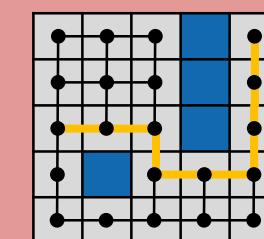


Graph



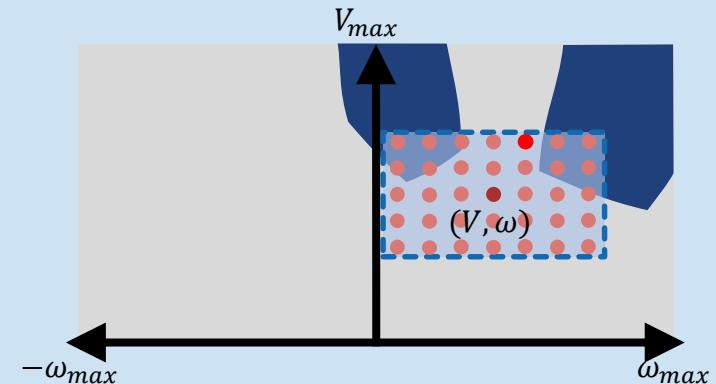
*Difficult problem,  
low-frequency updates*

Search



GLOBAL PLANNER

*Easier problem,  
high-frequency updates*

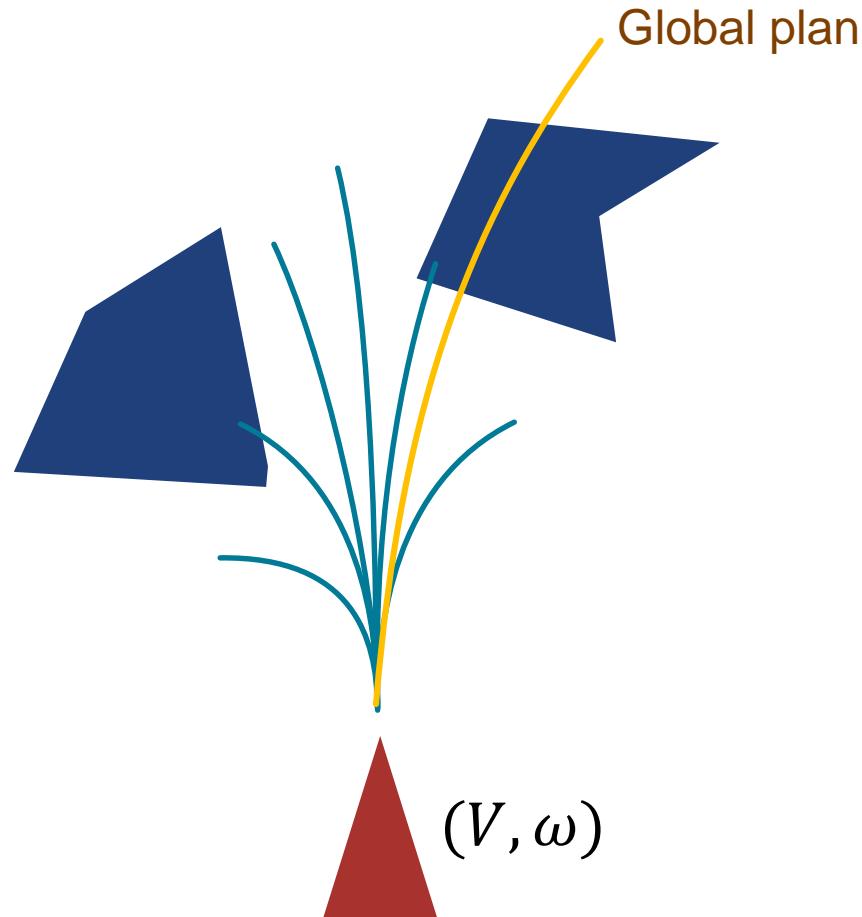


Global plan

( $V, \omega$ )

LOCAL PLANNER

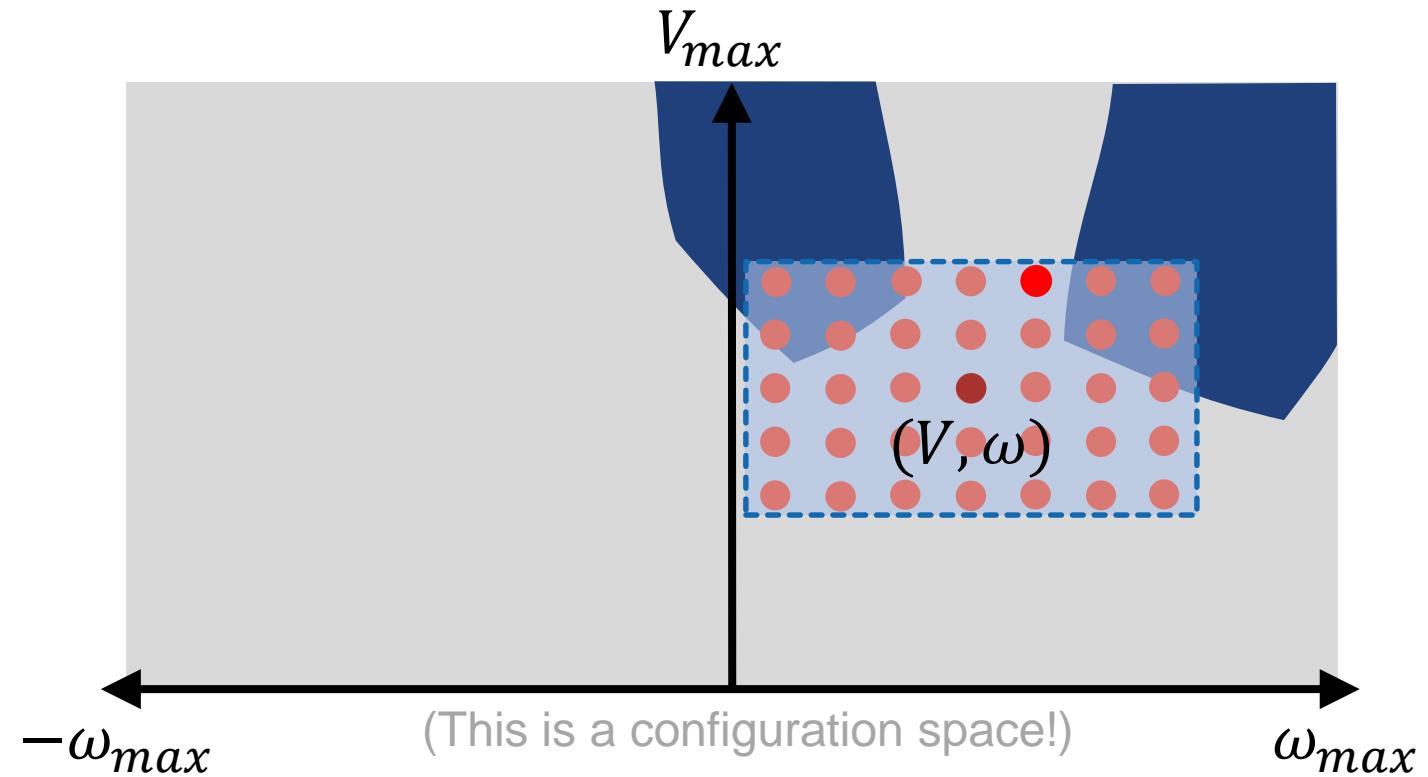
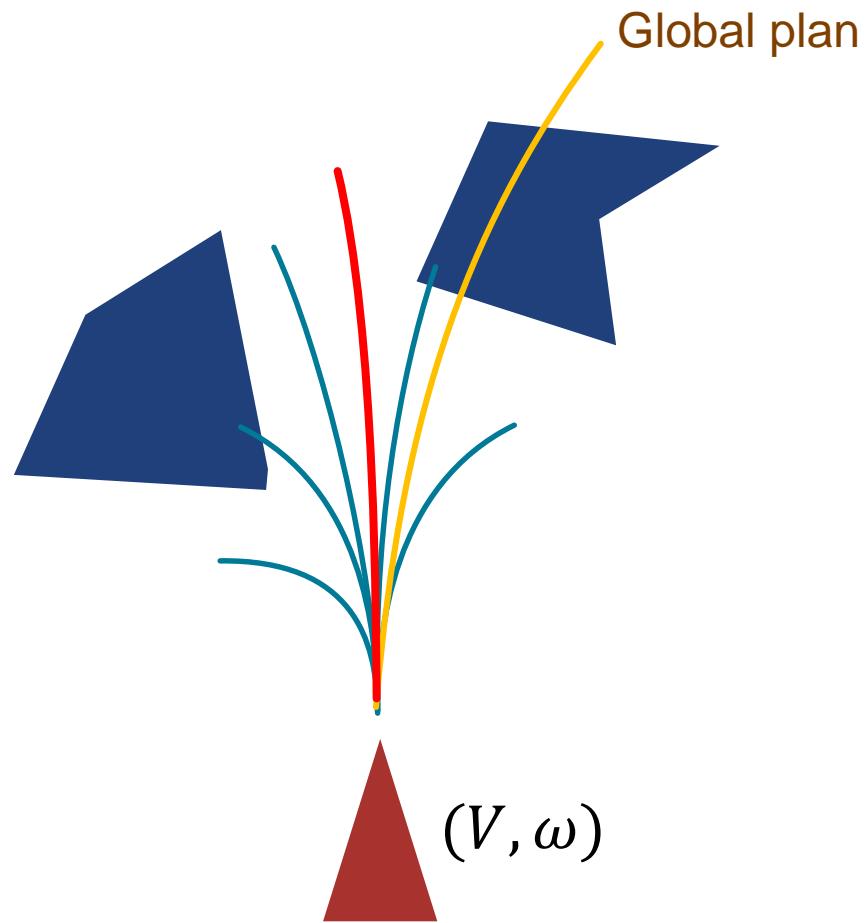
# Dynamic Window Approach Collision Avoidance



- Generating full time-varying trajectories for  $V(t)$  and  $\omega(t)$  is still very challenging
- If we assume  $(V, \omega)$  are constant for a fixed  $\Delta t$ , each local path in the future is a circular arc segment
- This can be easily considered as a type of *velocity configuration space*

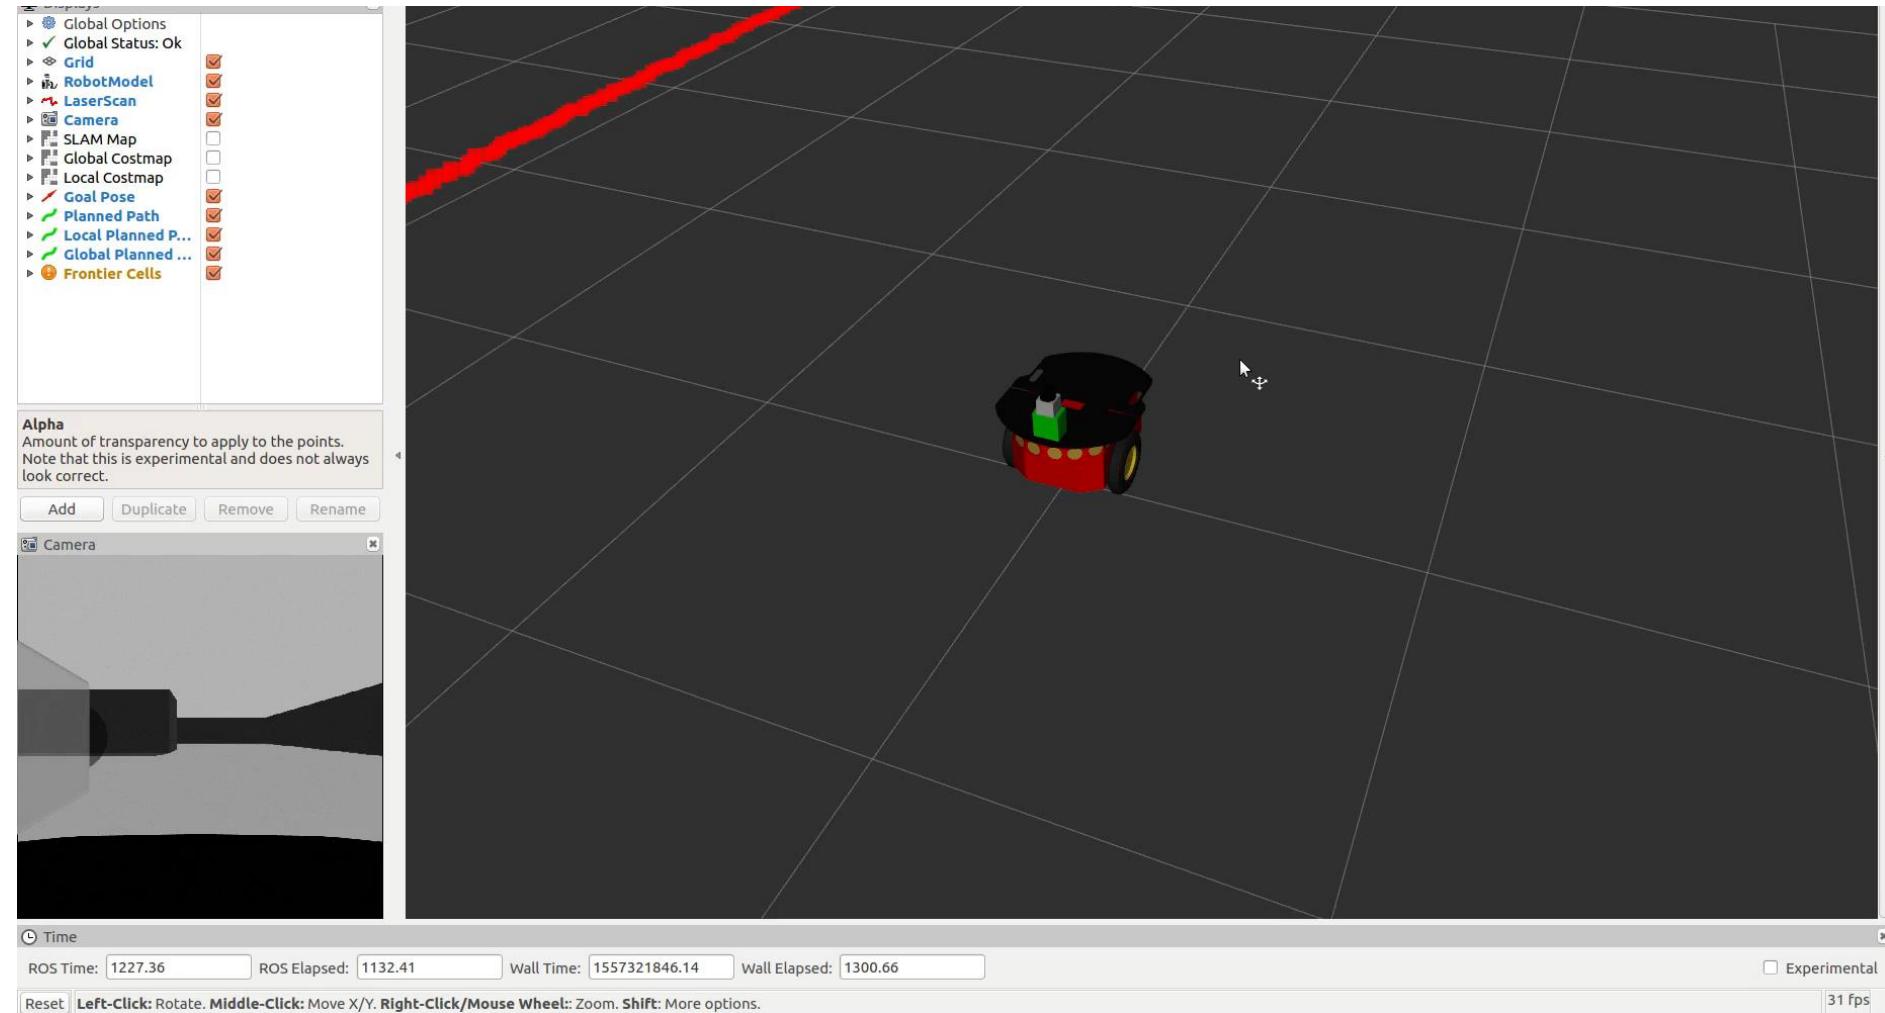
Fox, D., Burgard, W., & Thrun, S. (1997). The dynamic window approach to collision avoidance. *IEEE Robotics & Automation Magazine*, 4(1), 23-33.

# Dynamic Window Collision Avoidance



- Maximise utility metric (usually maximise speed, minimise distance to goal, maximise distance from obstacles) across configuration samples

# ROS Navigation Stack (*Dijkstra global + DWA local*)



# Collision Avoidance

- Sometimes, you may prefer to avoid the hierarchy and consider both global and local planning simultaneously
- A conceptually simple and relatively common approach for this type of problem are **potential field** methods

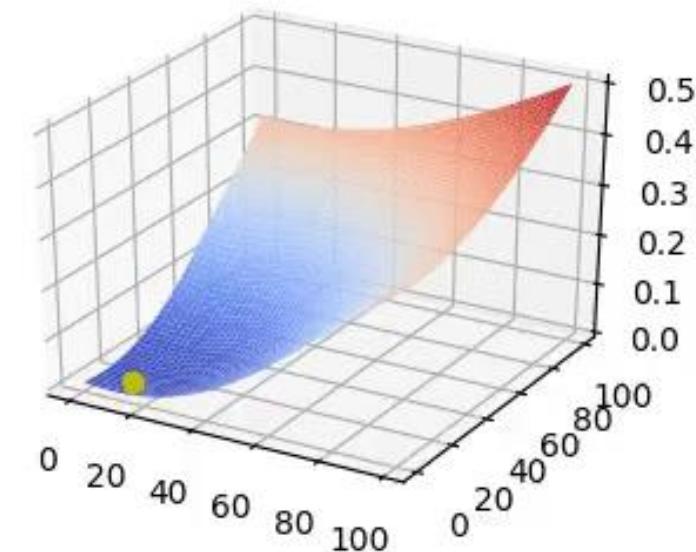
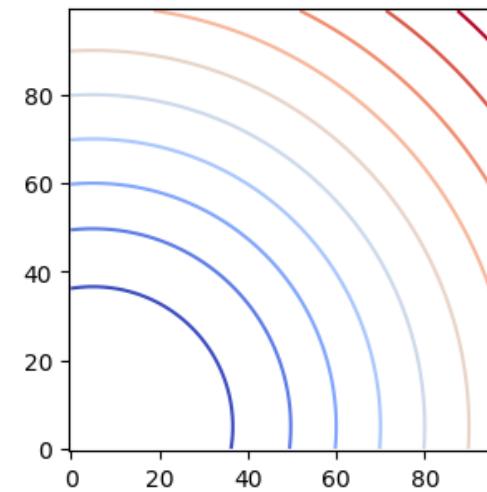
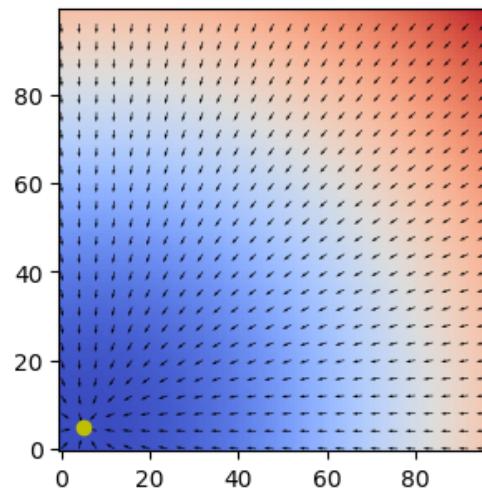
# Potential Field Methods – Global potential

- Imagine we had a function that related an **elevation** with some kind of distance to the goal (a ‘potential’ function, as in potential energy)
- By taking control actions that direct the robot in the direction of **maximum gradient (down)**, the robot should ‘fall’ towards the goal (minimum energy state)
- The global potential function should ‘attract’ the robot towards the goal, from any valid state

# Potential Field Methods – Global potential

- We want a smooth, differentiable function so that it is easy to calculate the target vector

$$U_{goal}(x) = \begin{cases} \frac{1}{2}\zeta\|x - x_{goal}\|^2, & \|x - x_{goal}\| < d^* \\ d^*\zeta\left(\|x - x_{goal}\|^2 - \frac{1}{2}d^*\right), & \text{otherwise} \end{cases}$$

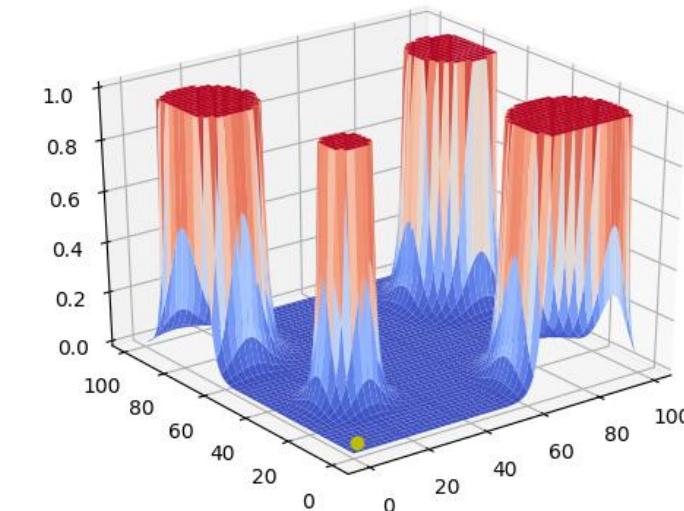
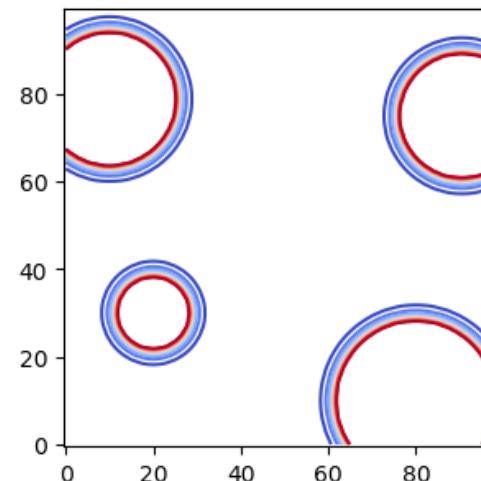
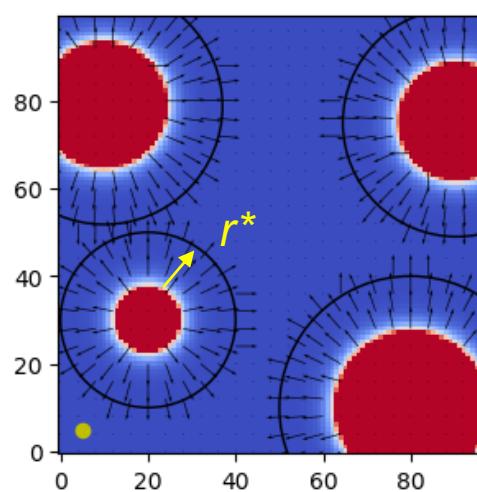


# Potential Field Methods – Obstacles

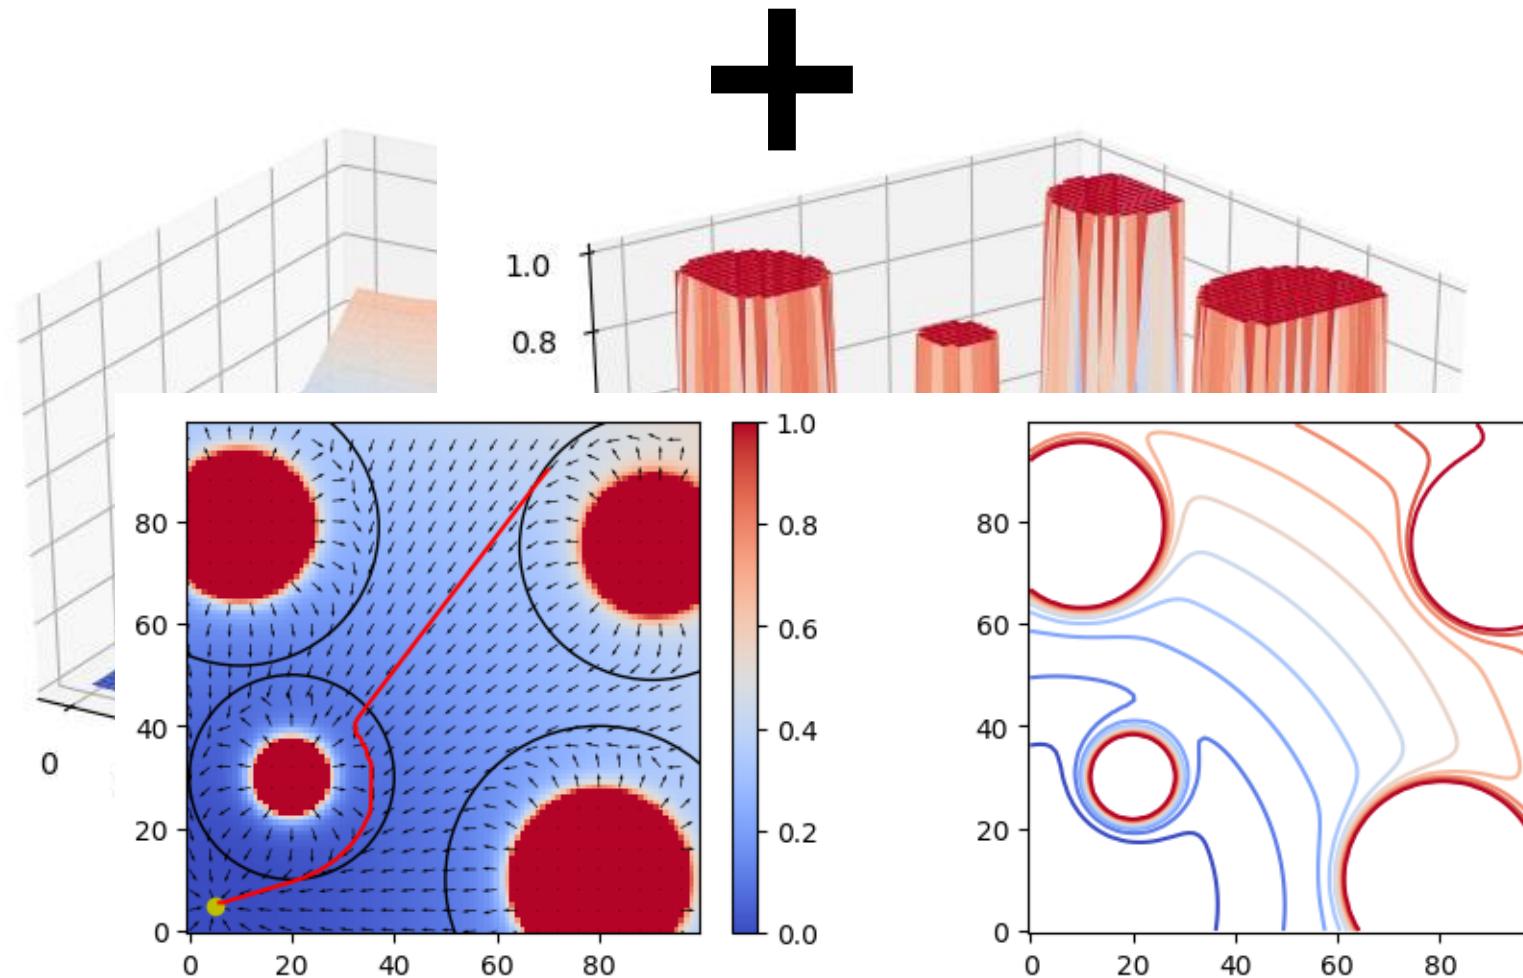
- We also want to avoid obstacles, so we add an additional component that ‘repels’ from obstacles

$$U_{obs}(x) = \begin{cases} \frac{1}{2}\eta \left( \frac{1}{D(x)} - \frac{1}{r^*} \right)^2, & D(x) \leq r^* \\ 0, & \text{otherwise} \end{cases}$$

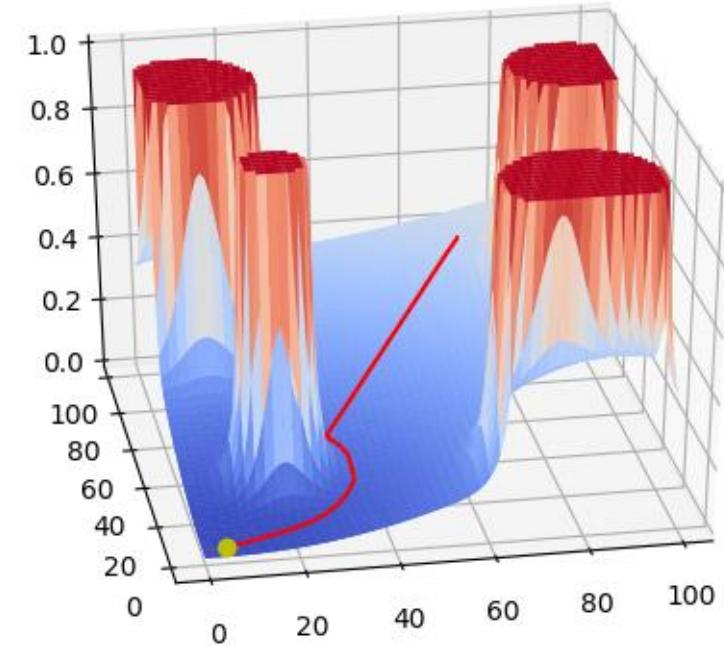
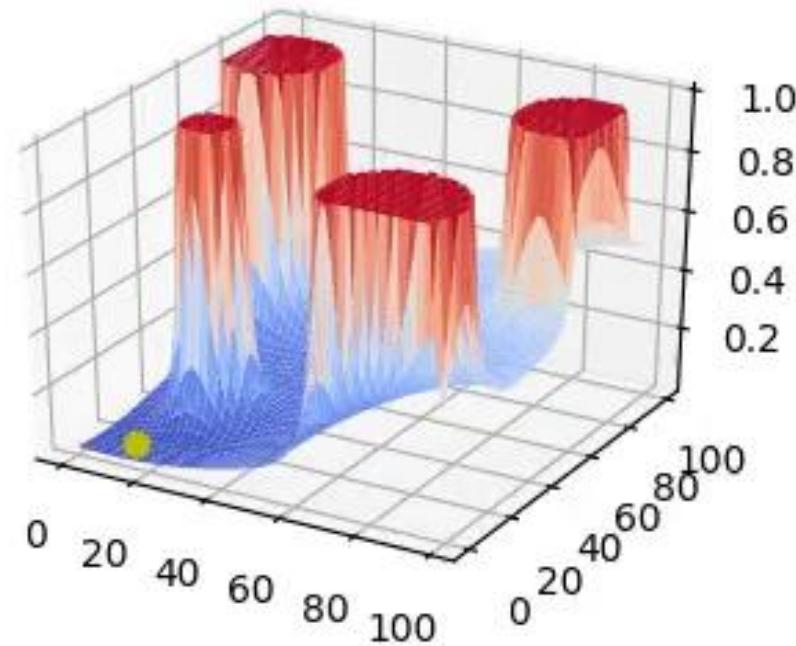
$*D(x)$  is the distance to the nearest obstacle boundary



# Potential Field

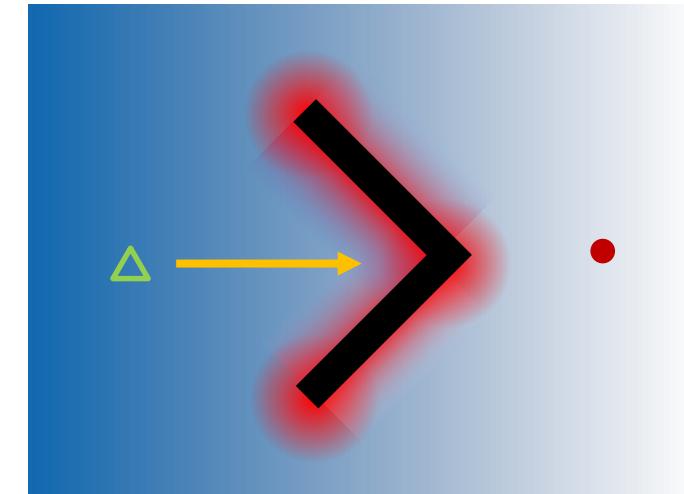


# Potential Field



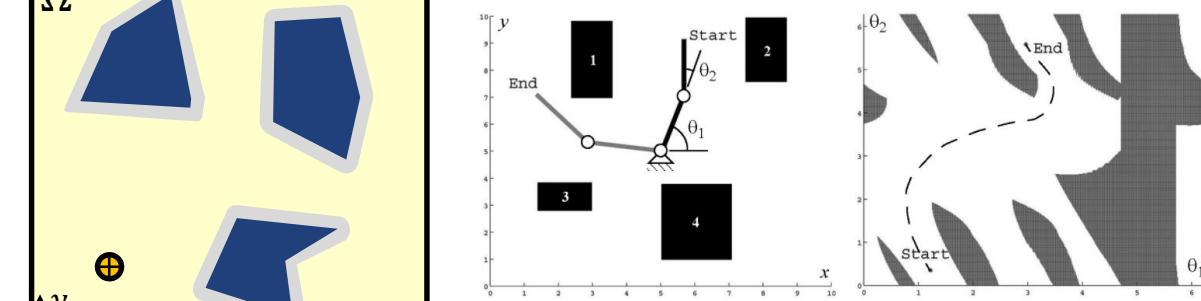
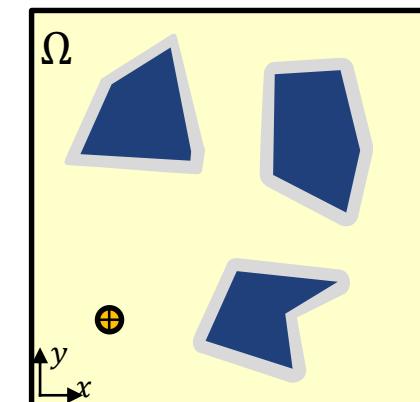
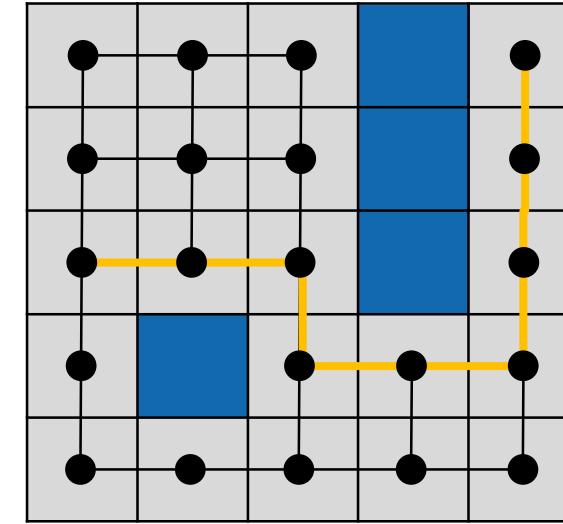
# Potential Field Methods

- Relatively simple to implement
- Simplest versions can have issues with stationary points or local minima
- Modifying the potential functions (iterative solution) can allow these conditions to be avoided (harmonic potentials, homework!)
- Dealing with higher-order state spaces (arms etc.) can be difficult to grasp/visualise, but potential functions are applicable



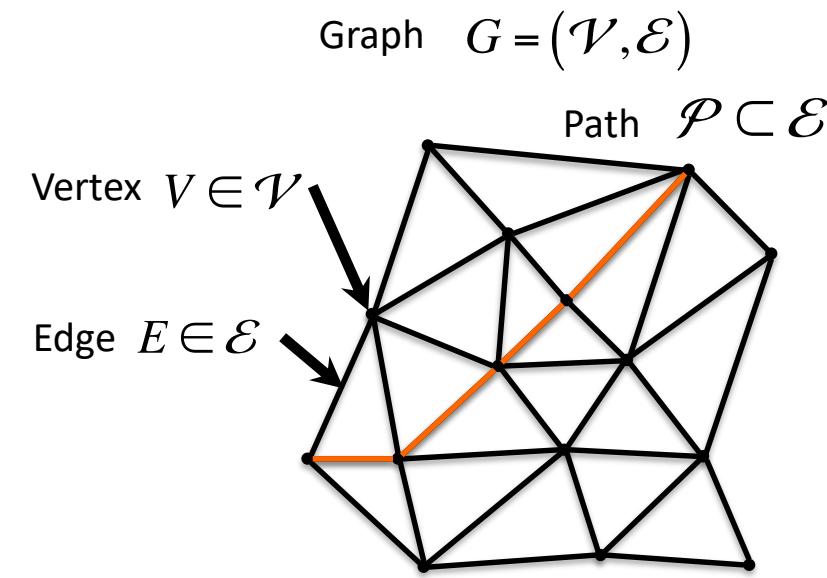
# Summary

- Motion planning
  - Representation – how to define the robot's understanding of the world, and ensure that it is sufficient to complete the task
  - Work space – the world without the robot
  - Configuration space – the robot's configuration (joint angles etc.) in the world



# Summary

- Graph search methods
  - Graphs are constructions of vertices and connecting edges
  - Graph search techniques are used to find low-cost paths through graphs
  - Breadth-first and depth-first search – complete searches from start (unweighted graphs)
  - Dijkstra – search outwards in order of cost from start (weighted graphs)
  - A\* – focused search that prioritises searching towards the goal using an admissible heuristic



# Summary

- Hierarchical planning
  - Global planner over the whole search space
  - Local planner to respond to changes in environment, avoid collisions, stay on global path
- Potential fields
  - Design a function such that descending the gradient leads to a collision-free path to the goal
- Additional references
  - Course text book and online lectures
  - Howie Choset (CMU) motion planning lecture notes:
    - [https://www.cs.cmu.edu/~motionplanning/lecture/Chap4-Potential-Field\\_howie.pdf](https://www.cs.cmu.edu/~motionplanning/lecture/Chap4-Potential-Field_howie.pdf)
  - Steven LaValle's Planning Algorithm Textbook
    - <http://planning.cs.uiuc.edu/>

Please fill out the  
class evaluation!

Python or Matlab?

<http://etc.ch/qbtw>



Thank you for listening!

**QUESTIONS?**