

# EECE 5550: Mobile Robotics



## Lecture 16: Motion planning


# Plan of the day 😊

**Last time:** Planning as search

- Definition of planning problems & solutions
- Planning as graph search
- Graph search algorithms

**Today:** Application to robot motion planning (yay 😊!)

- Robot workspaces & configuration spaces
- Grid-based motion planners
- Sampling-based planners for **high-dimensional** spaces
  - Probabilistic road maps (**PRMs**)
  - Rapidly-exploring random trees (**RRTs**)



# References



Classic (and very beautiful!) papers:

- “Probabilistic Roadmaps for Path Planning in High-Dimensional Configuration Spaces”
- “Randomized Kinodynamic Planning”

Lecture “Planning II” from ETH Zurich’s Autonomous Mobile Robots course



# Workspace and configuration spaces

**Workspace:** The *environment* in which the robot is operating.  
Often modeled at the level of *free* and *occupied* space


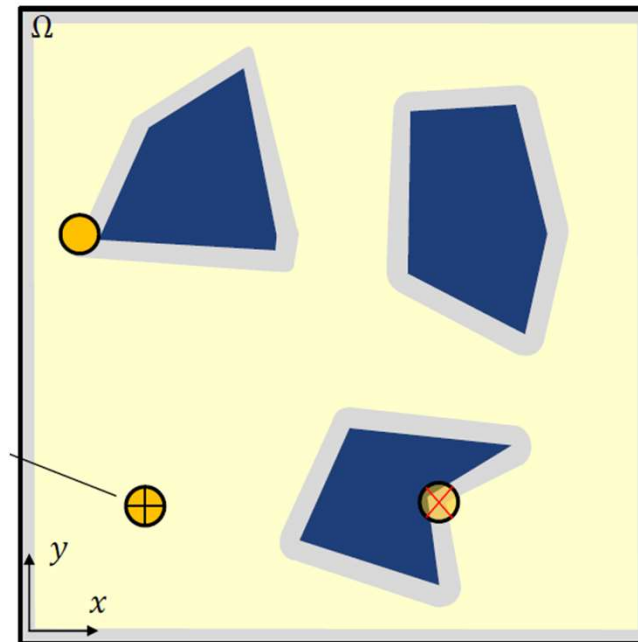
**Configuration space:** The set of feasible *robot states*

- Partially determined by workspace (no collisions permitted)
- For “interesting” (i.e. non-point) robots, also includes:
  - Orientations
  - Actuator limits ...

# Workspace and configuration spaces




Configuration space for a point robot




Configuration space for a disc robot

# Workspace and configuration space: two-link arm



Workspace




Configuration space

# Robot motion planning

## Given:

- **Workspace**  $W$ , partitioned into *free* and *occupied* subsets
- Robot **configuration space**  $C$ , partitioned into corresponding *free* and *occupied* subsets
- **Initial configuration**  $q_I \in C_{free}$
- **Goal configuration**  $q_G \in C_{free}$

**Find:** A *path*  $\tau: [0,1] \rightarrow C_{free}$  such that  $\tau(0) = q_I$  and  $\tau(1) = q_G$




# Robot motion planning: Challenges

**Basic challenge:** Workspaces and robot configuration spaces are generally *continuous*

⇒ **Very** hard to model arbitrary continuous shapes (Fun fact: these form *infinite-dimensional* spaces.)

**One natural approach:** *Discretize* the configuration space to get a *finite approximation*, and then solve the problem via *search*.





# Simple case: route planning for a planar robot


## Recall:

- **Occupancy grids** already provide a discrete model of free and occupied space
- For a *point* robot,  $C_{free}$  is just the set of unoccupied cells
- For *disc* robot, can also “grow” occupied cells to account for the body

**Now:** Construct a graph  $G = (V, E)$  where:

- Vertices are free cells
- Edges model connectivity between free cells

**Then:** Motion planning is just graph search!



# Grid-based representations


More generally, one can “voxelize” any *generic* configuration space  $C_{free}$ , and then apply the same graph search strategy

## Pros:

- Very simple idea
- Easy to implement
- **Resolution completeness**: If there is a feasible plan, this approach is *guaranteed* to find it as grid resolution  $r \rightarrow 0$ .

## Con:

- Voxelization can introduce weird artifacts
- Not always clear how to choose resolution  $r$ : how small is “small enough”??
- **Curse of dimensionality**: Number of cells  $N$  in the voxel grid grows as  $O(r^{-d})$ !





# Sampling-based planners

**Main idea:** Rather than trying to capture *every* point in the configuration space  $C$  via voxelization, let's *randomly sample* a *representative set*  $S$  of points in  $C_{free}$

**Then:**

- $S \subset C_{free}$  provides a *sample-based inner approximation* of  $C_{free}$
- **NB:** It's often easy to plan feasible motions between two *nearby* points  $x, y \in S$  (E.g.: a straight-line path often suffices ...)
- If we draw an edge  $e$  between two points  $x, y \in S$  whenever we can *locally* plan a feasible path from one to the other, we get a graph  $G = (V, E)$  that models *reachability*


# Grid- vs. sample-based planners



# Sampling-based planning: General framework

Starting with an empty graph  $G$ , repeat:


1. Sample a random point  $x \in C_{free}$  and add to  $G$   
[Q: How to sample  $x$ ?]
2. For each vertex  $y \in G$  s.t.  $d(x, y) < \epsilon$ , try to  
**plan a path from  $x$  to  $y$**   
  
**NB:** This requires:
  - A suitable notion of **distance** for  $C$
  - A **fast local planner**
3. If a feasible path is found, add edge  $(x, y)$  to  $G$ .




# The Probabilistic Roadmap (PRM) Algorithm

Two-phase algorithm for sampling-based planning:


1. **Construction:** Build a **roadmap** (graph) by randomly sampling points over the **entire** configuration space and planning local paths
2. **Query:** At run-time, given initial point  $x \in G$  and goal  $y \in G$ :
  - Plan **local paths** from  $x$  and  $y$  to (nearby) vertices in the **same connected component** of  $G$
  - **Graph search!**



# Example: PRM planning with a 5 DOF robotic arm



## Example: Solving the piano mover's problem with PRMs





# Probabilistic Roadmaps: Key Properties

**Recap:** Builds a roadmap (graph)  $G$  over the *entire* space  $C_{free}$  by joining a *sparse sample set* using *local planning*


- **Much** more scalable than grid methods
  - We *sparsely (inner) approximate* true space  $C_{free}$  using sample-set  $S$
  - Works well for high degree-of-freedom robots (e.g. robot arms, etc.)
- **Probabilistically complete:** If a feasible plan exists, the probability of finding it approaches 1 as the number of samples grows.
- Convenient “anytime” flavor: We can stop building the graph whenever we want, and we still get something useful
- Reusable! Enables fast *multi-query* planning

**BUT:** Costs a lot up-front (we build the roadmap  $G$  over the *entire* space  $C_{free}$ )

**Q:** Can we get something faster in the *single-query* case?

# Rapidly-exploring Random Trees (RRTs)

**Main idea:** Instead of building a *graph* by sampling vertices *uniformly* over  $C_{free}$ , we **build a tree outwards from the initial state  $x$  towards the goal  $y$ .**



# RRT algorithm

---

```


BUILD_RRT( $x_{init}$ )
1    $\mathcal{T}.\text{init}(x_{init});$ 
2   for  $k = 1$  to  $K$  do
3        $x_{rand} \leftarrow \text{RANDOM\_STATE}();$ 
4        $\text{EXTEND}(\mathcal{T}, x_{rand});$ 
5   Return  $\mathcal{T}$ 
    
```

---

```

EXTEND( $\mathcal{T}, x$ )
1    $x_{near} \leftarrow \text{NEAREST\_NEIGHBOR}(x, \mathcal{T});$ 
2   if  $\text{NEW\_STATE}(x, x_{near}, x_{new}, u_{new})$  then
3        $\mathcal{T}.\text{add\_vertex}(x_{new});$ 
4        $\mathcal{T}.\text{add\_edge}(x_{near}, x_{new}, u_{new});$ 
5       if  $x_{new} = x$  then
6           Return Reached;
7       else
8           Return Advanced;
9   Return Trapped;
    
```

---



NEW\_STATE:  $f_{new}$  is gotten from  $x_{near}$  and  $u_{new}$  by forward simulation using system dynamics:


$$\dot{x} = f(x, u)$$

Fig. 5. Basic rapidly exploring random tree construction algorithm.

**Key point:** Unlike PRM, in RRT we *do not* require that the control  $u_{new}$  drives  $x_{near}$  to  $x$ ; only that it *makes progress towards*  $x$

# What makes RRTs “rapidly exploring”?

**Voronoi bias:** Random sampling tends to place new vertices in larger Voronoi cells ( $\Rightarrow$  unexplored regions)



# RRT: Key Properties

- Probabilistically complete
- Single query & directed
  - Edges in RRTs are **directed**: travel is **from** the initial location **towards** the goal ( $x \rightarrow y \neq y \rightarrow x$ )
  - RRTs **retain control information**  $u$  in their edges
  - **Key payoff**: Unlike PRMs, RRTs can easily handle **dynamic constraints**:

$$\dot{x} = f(x, u)$$

In fact, RRTs were explicitly developed for **kinodynamic planning** (planning w/ kinematic and dynamic constraints)

⇒ Often applied to planning in **phase space** (position *and* velocity)

# Planning with phase space as configuration space





Fig. 2. Slices of  $\mathcal{X}$  for a point mass robot in two dimensions with increasingly higher initial speeds. White areas represent  $\mathcal{X}_{free}$ , black areas are  $\mathcal{X}_{obst}$ , and gray areas approximate  $\mathcal{X}_{ric}$ .

## Example: RRT Dubin's path planner



# RRT trajectory planning with racecar dynamics





# Variations on a Theme: Biased RRT sampling

---

**Input:**  $q_{\text{start}}$ ,  $q_{\text{goal}}$ , number  $n$  of nodes, stepsize  $\alpha$ ,  $\beta$

**Output:** tree  $T = (V, E)$

- 1: initialize  $V = \{q_{\text{start}}\}$ ,  $E = \emptyset$
  - 2: **for**  $i = 0 : n$  **do**
  - 3:     **if**  $\text{rand}(0, 1) < \beta$  **then**  $q_{\text{target}} \leftarrow q_{\text{goal}}$
  - 4:     **else**  $q_{\text{target}} \leftarrow$  random sample from  $Q$
  - 5:      $q_{\text{near}} \leftarrow$  nearest neighbor of  $q_{\text{target}}$  in  $V$
  - 6:      $q_{\text{new}} \leftarrow q_{\text{near}} + \frac{\alpha}{|q_{\text{target}} - q_{\text{near}}|} (q_{\text{target}} - q_{\text{near}})$
  - 7:     **if**  $q_{\text{new}} \in Q_{\text{free}}$  **then**  $V \leftarrow V \cup \{q_{\text{new}}\}$ ,  $E \leftarrow E \cup \{(q_{\text{near}}, q_{\text{new}})\}$
  - 8: **end for**
- 

Biasing tree expansion towards the goal: Sample the goal state  $q_{\text{goal}}$  itself with probability  $\beta > 0$

# Variations on a Theme: Bi-directional RRT

---

```
RRT_BIDIRECTIONAL( $x_{init}, x_{goal}$ );  
1    $\mathcal{T}_a.\text{init}(x_{init}); \mathcal{T}_b.\text{init}(x_{goal});$   
2   for  $k = 1$  to  $K$  do  
3        $x_{rand} \leftarrow \text{RANDOM\_STATE}();$   
4       if not ( $\text{EXTEND}(\mathcal{T}_a, x_{rand}) = \text{Trapped}$ ) then  
5           if ( $\text{EXTEND}(\mathcal{T}_b, x_{new}) = \text{Reached}$ ) then  
6               Return  $\text{PATH}(\mathcal{T}_a, \mathcal{T}_b);$   
7        $\text{SWAP}(\mathcal{T}_a, \mathcal{T}_b);$   
8   Return Failure
```

---

Fig. 7. A bidirectional rapidly exploring random trees-based planner.

Bi-directional RRT grows two trees towards each other: one from the initial state, and one from the goal

# Example: Bi-directional RRT path-planning

